



# **Red Hat Enterprise Linux 7 High Availability Add-On Reference**

---

Reference Document for the High Availability Add-On for Red Hat Enterprise Linux 7

Steven Levine



# Red Hat Enterprise Linux 7 High Availability Add-On Reference

---

## Reference Document for the High Availability Add-On for Red Hat Enterprise Linux 7

Steven Levine  
Red Hat Customer Content Services  
[slevine@redhat.com](mailto:slevine@redhat.com)

## **Legal Notice**

Copyright © 2016 Red Hat, Inc. and others.

This document is licensed by Red Hat under the [Creative Commons Attribution-ShareAlike 3.0 Unported License](#). If you distribute this document, or a modified version of it, you must provide attribution to Red Hat, Inc. and provide a link to the original. If the document is modified, all Red Hat trademarks must be removed.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux ® is the registered trademark of Linus Torvalds in the United States and other countries.

Java ® is a registered trademark of Oracle and/or its affiliates.

XFS ® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL ® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js ® is an official trademark of Joyent. Red Hat Software Collections is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack ® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

## **Abstract**

Red Hat High Availability Add-On Reference provides reference information about installing, configuring, and managing the Red Hat High Availability Add-On for Red Hat Enterprise Linux 7.

## Table of Contents

<b>Chapter 1. Red Hat High Availability Add-On Configuration and Management Reference.....</b>	
<b>Overview</b>	<b>4</b>
1.1. New and Changed Features	4
1.2. Installing Pacemaker configuration tools	6
1.3. Configuring the iptables Firewall to Allow Cluster Components	6
1.4. The Cluster and Pacemaker Configuration Files	7
1.5. Cluster Configuration Considerations	7
<b>Chapter 2. The pcsd Web UI .....</b>	<b>8</b>
2.1. pcsd Web UI Setup	8
2.2. Creating a Cluster with the pcsd Web UI	9
2.3. Configuring Cluster Components	11
<b>Chapter 3. The pcs Command Line Interface .....</b>	<b>14</b>
3.1. The pcs Commands	14
3.2. pcs Usage Help Display	14
3.3. Viewing the Raw Cluster Configuration	15
3.4. Saving a Configuration Change to a File	15
3.5. Displaying Status	15
3.6. Displaying the Full Cluster Configuration	16
3.7. Displaying The Current pcs Version	16
3.8. Backing Up and Restoring a Cluster Configuration	16
<b>Chapter 4. Cluster Creation and Administration .....</b>	<b>17</b>
4.1. Cluster Creation	17
4.2. Configuring Timeout Values for a Cluster	18
4.3. Configuring Redundant Ring Protocol (RRP)	19
4.4. Managing Cluster Nodes	19
4.5. Setting User Permissions	21
4.6. Removing the Cluster Configuration	23
4.7. Displaying Cluster Status	23
<b>Chapter 5. Fencing: Configuring STONITH .....</b>	<b>25</b>
5.1. Available STONITH (Fencing) Agents	25
5.2. General Properties of Fencing Devices	25
5.3. Displaying Device-Specific Fencing Options	26
5.4. Creating a Fencing Device	27
5.5. Configuring Storage-Based Fence Devices with unfencing	27
5.6. Displaying Fencing Devices	27
5.7. Modifying and Deleting Fencing Devices	28
5.8. Managing Nodes with Fence Devices	28
5.9. Additional Fencing Configuration Options	28
5.10. Configuring Fencing Levels	31
5.11. Configuring Fencing for Redundant Power Supplies	32
<b>Chapter 6. Configuring Cluster Resources .....</b>	<b>33</b>
6.1. Resource Creation	33
6.2. Resource Properties	34
6.3. Resource-Specific Parameters	34
6.4. Resource Meta Options	35
6.5. Resource Groups	37
6.6. Resource Operations	39
6.7. Displaying Configured Resources	41

6.8. Modifying Resource Parameters	41
6.9. Multiple Monitoring Operations	42
6.10. Enabling and Disabling Cluster Resources	42
6.11. Cluster Resources Cleanup	43
<b>Chapter 7. Resource Constraints</b>	<b>44</b>
7.1. Location Constraints	44
7.2. Order Constraints	46
7.3. Colocation of Resources	49
7.4. Displaying Constraints	51
<b>Chapter 8. Managing Cluster Resources</b>	<b>52</b>
8.1. Manually Moving Resources Around the Cluster	52
8.2. Moving Resources Due to Failure	53
8.3. Moving Resources Due to Connectivity Changes	54
8.4. Enabling, Disabling, and Banning Cluster Resources	55
8.5. Disabling a Monitor Operations	56
8.6. Managed Resources	56
<b>Chapter 9. Advanced Resource types</b>	<b>58</b>
9.1. Resource Clones	58
9.2. MultiState Resources: Resources That Have Multiple Modes	60
9.3. The pacemaker_remote Service	62
<b>Chapter 10. Cluster Quorum</b>	<b>69</b>
10.1. Configuring Quorum Options	69
10.2. Quorum Administration Commands (Red Hat Enterprise Linux 7.3 and Later)	70
10.3. Modifying Quorum Options (Red Hat Enterprise Linux 7.3 and later)	70
10.4. The quorum unblock Command	71
10.5. Quorum Devices (Technical Preview)	71
<b>Chapter 11. Pacemaker Rules</b>	<b>78</b>
11.1. Node Attribute Expressions	78
11.2. Time/Date Based Expressions	79
11.3. Date Specifications	79
11.4. Durations	80
11.5. Configuring Rules with pcs	80
11.6. Sample Time Based Expressions	80
11.7. Using Rules to Determine Resource Location	81
<b>Chapter 12. Pacemaker Cluster Properties</b>	<b>82</b>
12.1. Summary of Cluster Properties and Options	82
12.2. Setting and Removing Cluster Properties	84
12.3. Querying Cluster Property Settings	84
<b>Chapter 13. Triggering Scripts for Cluster Events</b>	<b>86</b>
13.1. Pacemaker Alert Agents (Red Hat Enterprise Linux 7.3 and later)	86
13.2. Event Notification with Monitoring Resources	92
<b>Chapter 14. Configuring Multi-Site Clusters with Pacemaker (Technical Preview)</b>	<b>95</b>
<b>Appendix A. Cluster Creation in Red Hat Enterprise Linux 6 and Red Hat Enterprise Linux 7</b>	<b>98</b>
A.1. Cluster Creation with rgmanager and with Pacemaker	98
A.2. Cluster Creation with Pacemaker in Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7	102

<b>Appendix B. Revision History .....</b>	<b>104</b>
<b>Index .....</b>	<b>104</b>

# Chapter 1. Red Hat High Availability Add-On Configuration and Management Reference Overview

This document provides descriptions of the options and features that the Red Hat High Availability Add-On using Pacemaker supports. For a step by step basic configuration example, refer to *Red Hat High Availability Add-On Administration*.

You can configure a Red Hat High Availability Add-On cluster with the **pcs** configuration interface or with the **pcsd** GUI interface.

## 1.1. New and Changed Features

This section lists features of the Red Hat High Availability Add-On that are new since the initial release of Red Hat Enterprise Linux 7.

### 1.1.1. New and Changed Features for Red Hat Enterprise Linux 7.1

Red Hat Enterprise Linux 7.1 includes the following documentation and feature updates and changes.

- » The **pcs resource cleanup** command can now reset the resource status and **failcount** for all resources, as documented in [Section 6.11, “Cluster Resources Cleanup”](#).
- » You can specify a **lifetime** parameter for the **pcs resource move** command, as documented in [Section 8.1, “Manually Moving Resources Around the Cluster”](#).
- » As of Red Hat Enterprise Linux 7.1, you can use the **pcs acl** command to set permissions for local users to allow read-only or read-write access to the cluster configuration by using access control lists (ACLs). For information on ACLs, see [Section 4.5, “Setting User Permissions”](#).
- » [Section 7.2.3, “Ordered Resource Sets”](#) and [Section 7.3, “Colocation of Resources”](#) have been extensively updated and clarified.
- » [Section 6.1, “Resource Creation”](#) documents the **disabled** parameter of the **pcs resource create** command, to indicate that the resource being created is not started automatically.
- » [Section 10.1, “Configuring Quorum Options”](#) documents the new **cluster quorum unlock** feature, which prevents the cluster from waiting for all nodes when establishing quorum.
- » [Section 6.1, “Resource Creation”](#) documents the **before** and **after** parameters of the **pcs resource create** command, which can be used to configure resource group ordering.
- » As of the Red Hat Enterprise Linux 7.1 release, you can backup the cluster configuration in a tarball and restore the cluster configuration files on all nodes from backup with the **backup** and **restore** options of the **pcs config** command. For information on this feature, see [Section 3.8, “Backing Up and Restoring a Cluster Configuration”](#).
- » Small clarifications have been made throughout this document.

### 1.1.2. New and Changed Features for Red Hat Enterprise Linux 7.2

Red Hat Enterprise Linux 7.2 includes the following documentation and feature updates and changes.

- ▶ You can now use the **pcs resource relocate run** command to move a resource to its preferred node, as determined by current cluster status, constraints, location of resources and other settings. For information on this command, see [Section 8.1.2, “Moving a Resource to its Preferred Node”](#).
- ▶ [Section 13.2, “Event Notification with Monitoring Resources”](#) has been modified and expanded to better document how to configure the **ClusterMon** resource to execute an external program to determine what to do with cluster notifications.
- ▶ When configuring fencing for redundant power supplies, you now are only required to define each device once and to specify that both devices are required to fence the node. For information on configuring fencing for redundant power supplies, see [Section 5.11, “Configuring Fencing for Redundant Power Supplies”](#).
- ▶ This document now provides a procedure for adding a node to an existing cluster in [Section 4.4.3, “Adding Cluster Nodes”](#).
- ▶ The new **resource-discovery** location constraint option allows you to indicate whether Pacemaker should perform resource discovery on a node for a specified resource, as documented in [Table 7.1, “Location Constraint Options”](#).
- ▶ Small clarifications and corrections have been made throughout this document.

### 1.1.3. New and Changed Features for Red Hat Enterprise Linux 7.3

Red Hat Enterprise Linux 7.3 includes the following documentation and feature updates and changes.

- ▶ [Section 9.3, “The pacemaker\\_remote Service”](#), has been wholly rewritten for this version of the document.
- ▶ You can configure Pacemaker alerts by means of alert agents, which are external programs that the cluster calls in the same manner as the cluster calls resource agents to handle resource configuration and operation. Pacemaker alert agents are described in [Section 13.1, “Pacemaker Alert Agents \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- ▶ New quorum administration commands are supported with this release which allow you to display the quorum status and to change the **expected\_votes** parameter. These commands are described in [Section 10.2, “Quorum Administration Commands \(Red Hat Enterprise Linux 7.3 and Later\)”](#).
- ▶ You can now modify general quorum options for your cluster with the **pcs quorum update** command, as described in [Section 10.3, “Modifying Quorum Options \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- ▶ You can configure a separate quorum device which acts as a third-party arbitration device for the cluster. The primary use of this feature is to allow a cluster to sustain more node failures than standard quorum rules allow. This feature is provided for technical preview only. For information on quorum devices, see [Section 10.5, “Quorum Devices \(Technical Preview\)”](#).
- ▶ Red Hat Enterprise Linux release 7.3 provides the ability to configure high availability clusters that span multiple sites through the use of a Booth cluster ticket manager. This feature is provided for technical preview only. For information on the Booth cluster ticket manager, see [Chapter 14, “Configuring Multi-Site Clusters with Pacemaker \(Technical Preview\)”](#).

- When configuring a KVM guest node running a the **pacemaker\_remote** service, you can include guest nodes in groups, which allows you to group a storage device, file system, and VM. For information on configuring KVM guest nodes, see [Section 9.3.5, “Configuration Overview: KVM Guest Node”](#).

Additionally, small clarifications and corrections have been made throughout this document.

## 1.2. Installing Pacemaker configuration tools

You can use the following **yum install** command to install the Red Hat High Availability Add-On software packages along with all available fence agents from the High Availability channel.

```
# yum install pcs pacemaker fence-agents-all
```

Alternately, you can install the Red Hat High Availability Add-On software packages along with only the fence agent that you require with the following command.

```
# yum install pcs pacemaker fence-agents-model
```

The following command displays a listing of the available fence agents.

```
# rpm -q -a | grep fence
fence-agents-rhev-4.0.2-3.el7.x86_64
fence-agents-ilo-mp-4.0.2-3.el7.x86_64
fence-agents-ipmilan-4.0.2-3.el7.x86_64
...
...
```

The **lvm2-cluster** and **gfs2-utils** packages are part of ResilientStorage channel. You can install them, as needed, with the following command.

```
# yum install lvm2-cluster gfs2-utils
```



### Warning

After you install the Red Hat High Availability Add-On packages, you should ensure that your software update preferences are set so that nothing is installed automatically. Installation on a running cluster can cause unexpected behaviors.

## 1.3. Configuring the iptables Firewall to Allow Cluster Components

The Red Hat High Availability Add-On requires that the following ports be enabled for incoming traffic:

- For TCP: Ports 2224, 3121, 21064
- For UDP: Ports 5405
- For DLM (if using the DLM lock manager with clvm/GFS2): Port 21064

You can enable these ports by means of the **firewalld** daemon by executing the following commands.

```
# firewall-cmd --permanent --add-service=high-availability  
# firewall-cmd --add-service=high-availability
```

## 1.4. The Cluster and Pacemaker Configuration Files

The configuration files for the Red Hat High Availability add-on are **corosync.conf** and **cib.xml**. Do not edit these files directly; use the **pcs** or **pcsd** interface instead.

The **corosync.conf** file provides the cluster parameters used by **corosync**, the cluster manager that Pacemaker is built on.

The **cib.xml** file is an XML file that represents both the cluster's configuration and current state of all resources in the cluster. This file is used by Pacemaker's Cluster Information Base (CIB). The contents of the CIB are automatically kept in sync across the entire cluster

## 1.5. Cluster Configuration Considerations

When configuring a Red Hat High Availability Add-On cluster, you must take the following considerations into account:

- » Red Hat does not support cluster deployments greater than 16 full cluster nodes. It is possible, however, to scale beyond that limit with remote nodes running the **pacemaker\_remote** service. For information on the **pacemaker\_remote** service, see [Section 9.3, “The pacemaker\\_remote Service”](#).
- » The use of Dynamic Host Configuration Protocol (DHCP) for obtaining an IP address on a network interface that is utilized by the **corosync** daemons is not supported. The DHCP client can periodically remove and re-add an IP address to its assigned interface during address renewal. This will result in **corosync** detecting a connection failure, which will result in fencing activity from any other nodes in the cluster using **corosync** for heartbeat connectivity.

## Chapter 2. The pcsd Web UI

This chapter provides an overview of configuring a Red Hat High Availability cluster with the **pcsd** Web UI.

### 2.1. pcsd Web UI Setup

To set up your system to use the **pcsd** Web UI to configure a cluster, use the following procedure.

1. Install the Pacemaker configuration tools, as described in [Section 1.2, “Installing Pacemaker configuration tools”](#).
2. On each node that will be part of the cluster, use the **passwd** command to set the password for user **hacluster**, using the same password on each node.
3. Start and enable the **pcsd** daemon on each node:

```
# systemctl start pcsd.service  
# systemctl enable pcsd.service
```

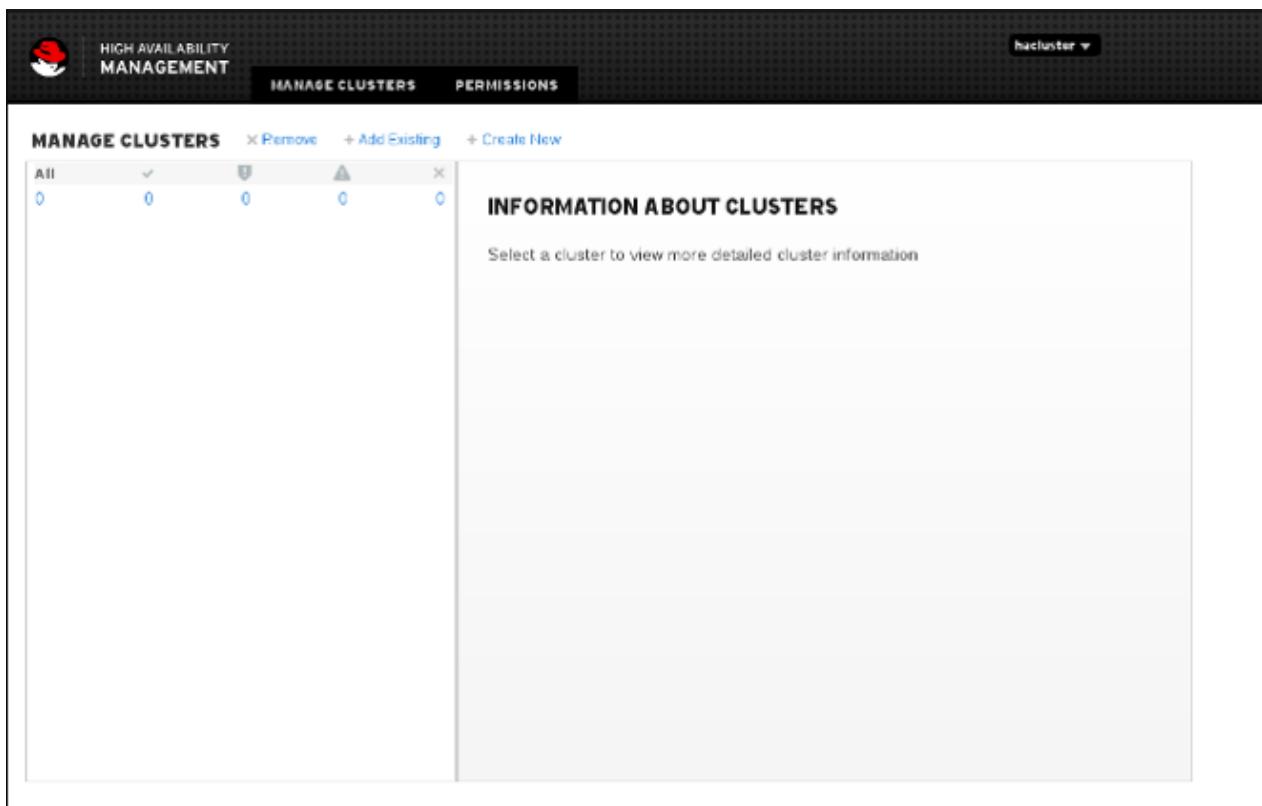
4. On one node of the cluster, authenticate the nodes that will constitute the cluster with the following command. After executing this command, you will be prompted for a **Username** and a **Password**. Specify **hacluster** as the **Username**.

```
# pcs cluster auth node1 node2 ... nodeN
```

5. On any system, open a browser to the following URL, specifying one of the nodes you have authorized (note that this uses the **https** protocol). This brings up the **pcsd** Web UI login screen.

```
https://nodename:2224
```

6. Log in as user **hacluster**. This brings up the **Manage Clusters** page as shown in [Figure 2.1, “Manage Clusters page”](#).



**Figure 2.1.** Manage Clusters page

## 2.2. Creating a Cluster with the pcsd Web UI

From the **Manage Clusters** page, you can create a new cluster, add an existing cluster to the Web UI, or remove a cluster from the Web UI.

- To create a cluster, click on **Create New** and enter the name of the cluster to create and the nodes that constitute the cluster. You can also configure advanced cluster options from this screen, including the transport mechanism for cluster communication, as described in [Section 2.2.1, “Advanced Cluster Configuration Options”](#). After entering the cluster information, click **Create Cluster**.
- To add an existing cluster to the Web UI, click on **Add Existing** and enter the host name or IP address of a node in the cluster that you would like to manage with the Web UI.

Once you have created or added a cluster, the cluster name is displayed on the **Manage Cluster** page. Selecting the cluster displays information about the cluster.

### Note

When using the **pcsd** Web UI to configure a cluster, you can move your mouse over the text describing many of the options to see longer descriptions of those options as a **tooltip** display.

### 2.2.1. Advanced Cluster Configuration Options

When creating a cluster, you can click on **Advanced Options** to configure additional cluster options, as shown in [Figure 2.2, “Create Clusters page”](#). For information about the options displayed, move your mouse over the text for that option.

Note that you can configure a cluster with Redundant Ring Protocol by specifying the interfaces for each node. The Redundant Ring Protocol settings display will change if you select **UDP** rather than the default value of **UDPU** as the transport mechanism for the cluster.

**Create Cluster** ×

Enter the hostnames of the nodes you would like to use to create a cluster:

Cluster Name:

Node 1:

Node 2:

Node 3:

[More nodes...](#)

▼ Advanced Options:

Transport: **UDPU (default)** ▼

Wait for All:

Auto Tie Breaker:

Last Man Standing:

Last Man Standing Window:  ms

Use IPv6:

Token Timeout:  ms

Token Timeout Coefficient:  ms

Join Timeout:  ms

Consensus Timeout:  ms

Missed Messages Count:

Failures Count:

Redundant Ring Protocol settings for UDPU transport:

Node 1 (Ring 1):

Node 2 (Ring 1):

Node 3 (Ring 1):

**Create Cluster** **Cancel**

**Figure 2.2. Create Clusters page**

## 2.2.2. Setting Cluster Permissions

From the **Manage Clusters** page, you can click on **Permissions** to set permissions for individual users or for groups. Note that user **hacluster** always has full permissions.

### Note

The permissions you set on this screen are for managing the cluster with the Web UI and are not the same as the permissions you set by using access control lists (ACL), which are described in [Section 2.3.4, “Configuring ACLs”](#).

From this screen, you can grant the following permissions:

- » Read permissions, to view the cluster settings.
- » Write permissions, to modify cluster settings (except for permissions and ACLs)
- » Grant permissions, to modify cluster permissions and ACLs
- » Full permissions, for unrestricted access to a cluster, including adding and removing nodes, with access to keys and certificates

By default, members of **haclient** group have full access to clusters.

## 2.3. Configuring Cluster Components

To configure the components and attributes of a cluster, click on the name of the cluster displayed on the **Manage Clusters** screen. This brings up the **Nodes** page, as described in [Section 2.3.1, “Cluster Nodes”](#). This page displays a menu along the top of the page, as shown in [Figure 2.3, “Cluster Components Menu”](#), with the following entries:

- » **Nodes**, as described in [Section 2.3.1, “Cluster Nodes”](#)
- » **Resources**, as described in [Section 2.3.2, “Cluster Resources”](#)
- » **Fence Devices**, as described in [Section 2.3.3, “Fence Devices”](#)
- » **ACLs**, as described in [Section 2.3.4, “Configuring ACLs”](#)
- » **Manage Clusters**, as described in [Section 2.3.5, “Cluster Properties”](#)

**Figure 2.3. Cluster Components Menu**

### 2.3.1. Cluster Nodes

Selecting the **Nodes** option from the menu along the top of the cluster management page displays the currently configured nodes and the status of the currently selected node, including which resources are running on the node and the resource location preferences. This is the default page that displays when you select a cluster from the **Manage Clusters** screen.

You can add or remove nodes from this page, and you can start, stop, restart, or put a node in standby mode. For information on standby mode, see [Section 4.4.5, “Standby Mode”](#).

You can also configure fence devices directly from this page, as described in [Section 2.3.3, “Fence Devices”](#), by selecting **Configure Fencing**.

### 2.3.2. Cluster Resources

Selecting the **Resources** option from the menu along the top of the cluster management page displays the currently configured resources for the cluster, organized according to resource groups. Selecting a group or a resource displays the attributes of that group or resource.

From this screen, you can add or remove resources, you can edit the configuration of existing resources, and you can create a resource group.

To add a new resource to the cluster, click **Add**. This brings up the **Add Resource** screen. When you select a resource type from the dropdown **Type** menu, the arguments you must specify for that resource appear in the menu. You can click **Optional Arguments** to display additional arguments you can specify for the resource you are defining. After entering the parameters for the resource you are creating, click **Create Resource**.

When configuring the arguments for a resource, a brief description of the argument appears in the menu. If you move the cursor to the field, a longer help description of that argument is displayed.

You can define a resource as a cloned resource, or as a master/slave resource. For information on these resource types, see [Chapter 9, Advanced Resource types](#).

Once you have created at least one resource, you can create a resource group. For information on resource groups, see [Section 6.5, “Resource Groups”](#).

To create a resource group, select a resource that will be part of the group from the **Resources** screen, then click **Create Group**. This displays the **Create Group** screen. Enter a group name and click **Create Group**. This returns you to the **Resources** screen, which now displays the group name for the resource. After you have created a resource group, you can indicate that group name as a resource parameter when you create or modify additional resources.

### 2.3.3. Fence Devices

Selecting the **Fence Devices** option from the menu along the top of the cluster management page displays **Fence Devices** screen, showing the currently configured fence devices.

To add a new fence device to the cluster, click **Add**. This brings up the **Add Fence Device** screen. When you select a fence device type from the drop-down **Type** menu, the arguments you must specify for that fence device appear in the menu. You can click on **Optional Arguments** to display additional arguments you can specify for the fence device you are defining. After entering the parameters for the new fence device, click **Create Fence Instance**.

For information on configuring fence devices with Pacemaker, see [Chapter 5, Fencing: Configuring STONITH](#).

### 2.3.4. Configuring ACLs

Selecting the **ACLS** option from the menu along the top of the cluster management page displays a screen from which you can set permissions for local users, allowing read-only or read-write access to the cluster configuration by using access control lists (ACLs).

To assign ACL permissions, you create a role and define the access permissions for that role. Each role can have an unlimited number of permissions (read/write/deny) applied to either an XPath query or the ID of a specific element. After defining the role, you can assign it to an existing user or group.

### 2.3.5. Cluster Properties

Selecting the **Cluster Properties** option from the menu along the top of the cluster management page displays the cluster properties and allows you to modify these properties from their default values. For information on the Pacemaker cluster properties, see [Chapter 12, Pacemaker Cluster Properties](#).

## Chapter 3. The pcs Command Line Interface

The **pcs** command line interface controls and configures **corosync** and Pacemaker by providing an interface to the **corosync.conf** and **cib.xml** files.

The general format of the **pcs** command is as follows.

```
pcs [-f file] [-h] [commands]...
```

### 3.1. The pcs Commands

The **pcs** commands are as follows.

- » **cluster**

Configure cluster options and nodes. For information on the **pcs cluster** command, see [Chapter 4, Cluster Creation and Administration](#).

- » **resource**

Create and manage cluster resources. For information on the **pcs cluster** command, see [Chapter 6, Configuring Cluster Resources](#), [Chapter 8, Managing Cluster Resources](#), and [Chapter 9, Advanced Resource types](#).

- » **stonith**

Configure fence devices for use with Pacemaker. For information on the **pcs stonith** command, see [Chapter 5, Fencing: Configuring STONITH](#).

- » **constraint**

Manage resource constraints. For information on the **pcs constraint** command, see [Chapter 7, Resource Constraints](#).

- » **property**

Set Pacemaker properties. For information on setting properties with the **pcs property** command, see [Chapter 12, Pacemaker Cluster Properties](#).

- » **status**

View current cluster and resource status. For information on the **pcs status** command, see [Section 3.5, “Displaying Status”](#).

- » **config**

Display complete cluster configuration in user readable form. For information on the **pcs config** command, see [Section 3.6, “Displaying the Full Cluster Configuration”](#).

### 3.2. pcs Usage Help Display

You can use the **-h** option of **pcs** to display the parameters of a **pcs** command and a description of those parameters. For example, the following command displays the parameters of the **pcs resource** command. Only a portion of the output is shown.

```
# pcs resource -h
```

Usage: **pcs resource [commands]...**

Manage pacemaker resources

Commands:

**show [resource id] [-all]**

Show all currently configured resources or if a resource is specified show the options for the configured resource. If --all is specified resource options will be displayed

**start <resource id>**

Start resource specified by resource\_id

...

### 3.3. Viewing the Raw Cluster Configuration

Although you should not edit the cluster configuration file directly, you can view the raw cluster configuration with the **pcs cluster cib** command.

You can save the raw cluster configuration to a specified file with the **pcs cluster cib filename** as described in [Section 3.4, “Saving a Configuration Change to a File”](#).

### 3.4. Saving a Configuration Change to a File

When using the **pcs** command, you can use the **-f** option to save a configuration change to a file without affecting the active CIB.

If you have previously configured a cluster and there is already an active CIB, you use the following command to save the raw xml a file.

**pcs cluster cib filename**

For example, the following command saves the raw xml from the CIB into a file name **testfile**.

**pcs cluster cib testfile**

The following command creates a resource in the file **testfile1** but does not add that resource to the currently running cluster configuration.

```
# pcs -f testfile1 resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.120 cidr_netmask=24 op monitor interval=30s
```

You can push the current content of **testfile** to the CIB with the following command.

**pcs cluster cib-push filename**

### 3.5. Displaying Status

You can display the status of the cluster and the cluster resources with the following command.

**pcs status commands**

If you do not specify a *commands* parameter, this command displays all information about the cluster and the resources. You display the status of only particular cluster components by specifying **resources**, **groups**, **cluster**, **nodes**, or **pcsd**.

## 3.6. Displaying the Full Cluster Configuration

Use the following command to display the full current cluster configuration.

```
pcs config
```

## 3.7. Displaying The Current pcs Version

The following command displays the current version of **pcs** that is running.

```
pcs --version
```

## 3.8. Backing Up and Restoring a Cluster Configuration

As of the Red Hat Enterprise Linux 7.1 release, you can back up the cluster configuration in a tarball with the following command. If you do not specify a file name, the standard output will be used.

```
pcs config backup filename
```

Use the following command to restore the cluster configuration files on all nodes from the backup. If you do not specify a file name, the standard input will be used. Specifying the **--local** option restores only the files on the current node.

```
pcs config restore [--local] [filename]
```

# Chapter 4. Cluster Creation and Administration

This chapter describes how to perform basic cluster administration with Pacemaker, including creating the cluster, managing the cluster components, and displaying cluster status.

## 4.1. Cluster Creation

To create a running cluster, perform the following steps:

1. Start the **pcsd** on each node in the cluster.
2. Authenticate the nodes that will constitute the cluster.
3. Configure and sync the cluster nodes.
4. Start cluster services on the cluster nodes.

The following sections described the commands that you use to perform these steps.

### 4.1.1. Starting the **pcsd** daemon

The following commands start the **pcsd** service and enable **pcsd** at system start. These commands should be run on each node in the cluster.

```
# systemctl start pcsd.service
# systemctl enable pcsd.service
```

### 4.1.2. Authenticating the Cluster Nodes

The following command authenticates **pcs** to the **pcs** daemon on the nodes in the cluster.

- ▶ The user name for the **pcs** administrator must be **hacluster** on every node. It is recommended that the password for user **hacluster** be the same on each node.
- ▶ If you do not specify **username** or **password**, the system will prompt you for those parameters for each node when you execute the command.
- ▶ If you do not specify any nodes, this command will authenticate **pcs** on the nodes that are specified with a **pcs cluster setup** command, if you have previously executed that command.

```
pcs cluster auth [node] [...] [-u username] [-p password]
```

For example, the following command authenticates user **hacluster** on **z1.example.com** for both of the nodes in the cluster that consist of **z1.example.com** and **z2.example.com**. This command prompts for the password for user **hacluster** on the cluster nodes.

```
root@z1 ~]# pcs cluster auth z1.example.com z2.example.com
Username: hacluster
Password:
z1.example.com: Authorized
z2.example.com: Authorized
```

Authorization tokens are stored in the file **~/.pcs/tokens** (or **/var/lib/pcsd/tokens**).

### 4.1.3. Configuring and Starting the Cluster Nodes

The following command configures the cluster configuration file and syncs the configuration to the specified nodes.

- If you specify the **--start** option, the command will also start the cluster services on the specified nodes. If necessary, you can also start the cluster services with a separate **pcs cluster start** command.

When you create a cluster with the **pcs cluster setup --start** command or when you start cluster services with the **pcs cluster start** command, there may be a slight delay before the cluster is up and running. Before performing any subsequent actions on the cluster and its configuration, it is recommended that you use the **pcs cluster status** command to be sure that the cluster is up and running.

- If you specify the **--local** option, the command will perform changes on the local node only.

```
pcs cluster setup [--start] [--local] --name cluster_name node1 [node2] [...]
```

The following command starts cluster services on the specified node or nodes.

- If you specify the **--all** option, the command starts cluster services on all nodes.
- If you do not specify any nodes, cluster services are started on the local node only.

```
pcs cluster start [--all] [node] [...]
```

## 4.2. Configuring Timeout Values for a Cluster

When you create a cluster with the **pcs cluster setup** command, timeout values for the cluster are set to default values that should be suitable for most cluster configurations. If your system requires different timeout values, however, you can modify these values with the **pcs cluster setup** options summarized in [Table 4.1, “Timeout Options”](#).

**Table 4.1. Timeout Options**

Option	Description
<b>--token <i>timeout</i></b>	Sets time in milliseconds until a token loss is declared after not receiving a token (default 1000 ms)
<b>--join <i>timeout</i></b>	sets time in milliseconds to wait for join messages (default 50 ms)
<b>--consensus <i>timeout</i></b>	sets time in milliseconds to wait for consensus to be achieved before starting a new round of membership configuration (default 1200 ms)
<b>--miss_count_const <i>count</i></b>	sets the maximum number of times on receipt of a token a message is checked for retransmission before a retransmission occurs (default 5 messages)
<b>--fail_recv_const <i>failures</i></b>	specifies how many rotations of the token without receiving any messages when messages should be received may occur before a new configuration is formed (default 2500 failures)

For example, the following command creates the cluster **new\_cluster** and sets the token timeout value to 10000 milliseconds (10 seconds) and the join timeout value to 100 milliseconds.

```
# pcs cluster setup --name new_cluster nodeA nodeB --token 10000 --join 100
```

## 4.3. Configuring Redundant Ring Protocol (RRP)

When you create a cluster with the **pcs cluster setup** command, you can configure a cluster with Redundant Ring Protocol by specifying both interfaces for each node. When using the default udpu transport, when you specify the cluster nodes you specify the ring 0 address followed by a ',', then the ring 1 address.

For example, the following command configures a cluster named **my\_rrp\_clusterM** with two nodes, node A and node B. Node A has two interfaces, **nodeA-0** and **nodeA-1**. Node B has two interfaces, **nodeB-0** and **nodeB-1**. To configure these nodes as a cluster using RRP, execute the following command.

```
# pcs cluster setup --name my_rrp_cluster nodeA-0,nodeA-1 nodeB-0,nodeB-1
```

For information on configuring RRP in a cluster that uses **udp** transport, see the help screen for the **pcs cluster setup** command.

## 4.4. Managing Cluster Nodes

The following sections describe the commands you use to manage cluster nodes, including commands to start and stop cluster services and to add and remove cluster nodes.

### 4.4.1. Stopping Cluster Services

The following command stops cluster services on the specified node or nodes. As with the **pcs cluster start**, the **--all** option stops cluster services on all nodes and if you do not specify any nodes, cluster services are stopped on the local node only.

```
pcs cluster stop [--all] [node] [...]
```

You can force a stop of cluster services on the local node with the following command, which performs a **kill -9** command.

```
pcs cluster kill
```

### 4.4.2. Enabling and Disabling Cluster Services

Use the following command to configure the cluster services to run on startup on the specified node or nodes.

- » If you specify the **--all** option, the command enables cluster services on all nodes.
- » If you do not specify any nodes, cluster services are enabled on the local node only.

```
pcs cluster enable [--all] [node] [...]
```

Use the following command to configure the cluster services not to run on startup on the specified node or nodes.

- » If you specify the **--all** option, the command disables cluster services on all nodes.
- » If you do not specify any nodes, cluster services are disabled on the local node only.

```
pcs cluster disable [--all] [node] [...]
```

#### 4.4.3. Adding Cluster Nodes

**Note**

It is highly recommended that you add nodes to existing clusters only during a production maintenance window. This allows you to perform appropriate resource and deployment testing for the new node and its fencing configuration.

Use the following procedure to add a new node to an existing cluster. In this example, the existing cluster nodes are **clusternode-01.example.com**, **clusternode-02.example.com**, and **clusternode-03.example.com**. The new node is **newnode.example.com**.

On the new node to add to the cluster, perform the following tasks.

1. Install the cluster packages. If the cluster uses SBD, the Booth ticket manager, or a quorum device, you must manually install the respective packages (**sbd**, **booth-site**, **corosync-device**) on the new node as well.

```
[root@newnode ~]# yum install -y pcs fence-agents-all
```

2. If you are running the **firewalld** daemon, execute the following commands to enable the ports that are required by the Red Hat High Availability Add-On.

```
# firewall-cmd --permanent --add-service=high-availability
# firewall-cmd --add-service=high-availability
```

3. Set a password for the user ID **hacluster**. It is recommended that you use the same password for each node in the cluster.

```
[root@newnode ~]# passwd hacluster
Changing password for user hacluster.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
```

4. Execute the following commands to start the **pcsd** service and to enable **pcsd** at system start.

```
# systemctl start pcasd.service
# systemctl enable pcasd.service
```

On a node in the existing cluster, perform the following tasks.

1. Authenticate user **hacluster** on the new cluster node.

```
[root@clusternode-01 ~]# pcs cluster auth newnode.example.com
Username: hacluster
Password:
newnode.example.com: Authorized
```

2. Add the new node to the existing cluster. This command also syncs the cluster configuration file **corosync.conf** to all nodes in the cluster, including the new node you are adding.

```
[root@clusternode-01 ~]# pcs cluster node add newnode.example.com
```

On the new node to add to the cluster, perform the following tasks.

1. Start and enable cluster services on the new node.

```
[root@newnode ~]# pcs cluster start
Starting Cluster...
[root@newnode ~]# pcs cluster enable
```

2. Ensure that you configure and test a fencing device for the new cluster node. For information on configuring fencing devices, see [Chapter 5, Fencing: Configuring STONITH](#).

#### 4.4.4. Removing Cluster Nodes

The following command shuts down the specified node and removes it from the cluster configuration file, **corosync.conf**, on all of the other nodes in the cluster. For information on removing all information about the cluster from the cluster nodes entirely, thereby destroying the cluster permanently, refer to [Section 4.6, “Removing the Cluster Configuration”](#).

```
pcs cluster node remove node
```

#### 4.4.5. Standby Mode

The following command puts the specified node into standby mode. The specified node is no longer able to host resources. Any resources currently active on the node will be moved to another node. If you specify the **--all**, this command puts all nodes into standby mode.

You can use this command when updating a resource's packages. You can also use this command when testing a configuration, to simulate recovery without actually shutting down a node.

```
pcs cluster standby node | --all
```

The following command removes the specified node from standby mode. After running this command, the specified node is then able to host resources. If you specify the **--all**, this command removes all nodes from standby mode.

```
pcs cluster unstandby node | --all
```

Note that when you execute the **pcs cluster standby** command, this adds constraints to the resources to prevent them from running on the indicated node. When you execute the **pcs cluster unstandby** command, this removes the constraints. This does not necessarily move the resources back to the indicated node; where the resources can run at that point depends on how you have configured your resources initially. For information on resource constraints, refer to [Chapter 7, Resource Constraints](#).

### 4.5. Setting User Permissions

By default, the root user and any user who is a member of the group **haclient** has full read/write access to the cluster configuration. As of Red Hat Enterprise Linux 7.1, you can use the **pcs acl**

command to set permissions for local users to allow read-only or read-write access to the cluster configuration by using access control lists (ACLs).

Setting permissions for local users is a two step process:

1. Execute the **pcs acl role create...** command to create a *role* which defines the permissions for that role.
2. Assign the role you created to a user with the **pcs acl user create** command.

The following example procedure provides read-only access for a cluster configuration to a local user named **rouser**.

1. This procedure requires that the user **rouser** exists on the local system and that the user **rouser** is a member of the group **haclient**.

```
# adduser rouser
# usermod -a -G haclient rouser
```

2. Enable Pacemaker ACLs with the **enable-acl** cluster property.

```
# pcs property set enable-acl=true --force
```

3. Create a role named **read-only** with read-only permissions for the cib.

```
# pcs acl role create read-only description="Read access to cluster" read
xpath /cib
```

4. Create the user **rouser** in the pcs ACL system and assign that user the **read-only** role.

```
# pcs acl user create rouser read-only
```

5. View the current ACLs.

```
# pcs acl
User: rouser
Roles: read-only
Role: read-only
Description: Read access to cluster
Permission: read xpath /cib (read-only-read)
```

The following example procedure provides write access for a cluster configuration to a local user named **wuser**.

1. This procedure requires that the user **wuser** exists on the local system and that the user **wuser** is a member of the group **haclient**.

```
# adduser wuser
# usermod -a -G haclient wuser
```

2. Enable Pacemaker ACLs with the **enable-acl** cluster property.

```
# pcs property set enable-acl=true --force
```

3. Create a role named **write-access** with write permissions for the cib.

```
# pcs acl role create write-access description="Full access" write xpath /cib
```

4. Create the user **wuser** in the pcs ACL system and assign that user the **write-access** role.

```
# pcs acl user create wuser write-access
```

5. View the current ACLs.

```
# pcs acl
User: rouser
  Roles: read-only
User: wuser
  Roles: write-access
Role: read-only
  Description: Read access to cluster
  Permission: read xpath /cib (read-only-read)
Role: write-access
  Description: Full Access
  Permission: write xpath /cib (write-access-write)
```

For further information about cluster ACLs, see the help screen for the **pcs acl** command.

## 4.6. Removing the Cluster Configuration

To remove all cluster configuration files and stop all cluster services, thus permanently destroying a cluster, use the following command.



### Warning

This command permanently removes any cluster configuration that has been created. It is recommended that you run **pcs cluster stop** before destroying the cluster.

```
pcs cluster destroy
```

## 4.7. Displaying Cluster Status

The following command displays the current status of the cluster and the cluster resources.

```
pcs status
```

You can display a subset of information about the current status of the cluster with the following commands.

The following command displays the status of the cluster, but not the cluster resources.

```
pcs cluster status
```

The following command displays the status of the cluster resources.

pcs status resources

# Chapter 5. Fencing: Configuring STONITH

STONITH is an acronym for "Shoot The Other Node In The Head" and it protects your data from being corrupted by rogue nodes or concurrent access.

Just because a node is unresponsive, this does not mean it is not accessing your data. The only way to be 100% sure that your data is safe, is to fence the node using STONITH so we can be certain that the node is truly offline, before allowing the data to be accessed from another node.

STONITH also has a role to play in the event that a clustered service cannot be stopped. In this case, the cluster uses STONITH to force the whole node offline, thereby making it safe to start the service elsewhere.

## 5.1. Available STONITH (Fencing) Agents

Use the following command to view of list of all available STONITH agents. You specify a filter, then this command displays only the STONITH agents that match the filter.

```
pcs stonith list [filter]
```

## 5.2. General Properties of Fencing Devices

### Note

To disable a fencing device/resource, you can set the **target-role** as you would for a normal resource.

### Note

To prevent a specific node from using a fencing device, you can configure location constraints for the fencing resource.

[Table 5.1, “General Properties of Fencing Devices”](#) describes the general properties you can set for fencing devices. Refer to [Section 5.3, “Displaying Device-Specific Fencing Options”](#) for information on fencing properties you can set for specific fencing devices.

### Note

For information on more advanced fencing configuration properties, refer to [Section 5.9, “Additional Fencing Configuration Options”](#).

**Table 5.1. General Properties of Fencing Devices**

Field	Type	Default	Description
-------	------	---------	-------------

Field	Type	Default	Description
<b>priority</b>	integer	0	The priority of the stonith resource. Devices are tried in order of highest priority to lowest.
<b>pcmk_host_map</b>	string		A mapping of host names to ports numbers for devices that do not support host names. For example: <b>node1:1;node2:2,3</b> tells the cluster to use port 1 for node1 and ports 2 and 3 for node2
<b>pcmk_host_list</b>	string		A list of machines controlled by this device (Optional unless <b>pcmk_host_check=static-list</b> ).
<b>pcmk_host_check</b>	string	dynamic-list	How to determine which machines are controlled by the device. Allowed values: <b>dynamic-list</b> (query the device), <b>static-list</b> (check the <b>pcmk_host_list</b> attribute), none (assume every device can fence every machine)

## 5.3. Displaying Device-Specific Fencing Options

Use the following command to view the options for the specified STONITH agent.

```
pcs stonith describe stonith_agent
```

For example, the following command displays the options for the fence agent for APC over telnet/SSH.

```
# pcs stonith describe fence_apc
Stonith options for: fence_apc
ipaddr (required): IP Address or Hostname
login (required): Login Name
passwd: Login password or passphrase
passwd_script: Script to retrieve password
cmd_prompt: Force command prompt
secure: SSH connection
port (required): Physical plug number or name of virtual machine
identity_file: Identity file for ssh
switch: Physical switch number on device
inet4_only: Forces agent to use IPv4 addresses only
inet6_only: Forces agent to use IPv6 addresses only
ipport: TCP port to use for connection with device
action (required): Fencing Action
verbose: Verbose mode
debug: Write debug information to given file
version: Display version information and exit
help: Display help and exit
separator: Separator for CSV created by operation list
power_timeout: Test X seconds for status change after ON/OFF
shell_timeout: Wait X seconds for cmd prompt after issuing command
login_timeout: Wait X seconds for cmd prompt after login
power_wait: Wait X seconds after issuing ON/OFF
delay: Wait X seconds before fencing is started
```

`retry_on`: Count of attempts to retry power on

## 5.4. Creating a Fencing Device

The following command creates a stonith device.

```
pcs stonith create stonith_id stonith_device_type [stonith_device_options]
```

```
# pcs stonith create MyStonith fence_virt pcmk_host_list=f1 op monitor
interval=30s
```

If you use a single fence device for several nodes, using a different port of each node, you do not need to create a device separately for each node. Instead you can use the **pcmk\_host\_map** option to define which port goes to which node. For example, the following command creates a single fencing device called **myapc-west-13** that uses an APC power switch called **west-apc** and uses port 15 for node **west-13**.

```
# pcs stonith create myapc-west-13 fence_apc pcmk_host_list="west-13"
ipaddr="west-apc" login="apc" passwd="apc" port="15"
```

The following example, however, uses the APC power switch named **west-apc** to fence nodes **west-13** using port 15, **west-14** using port 17, **west-15** using port 18, and **west-16** using port 19.

```
# pcs stonith create myapc fence_apc pcmk_host_list="west-13,west-14,west-
15,west-16" pcmk_host_map="west-13:15;west-14:17;west-15:18;west-16:19"
ipaddr="west-apc" login="apc" passwd="apc"
```

## 5.5. Configuring Storage-Based Fence Devices with unfencing

When creating a SAN/storage fence device (that is, one that uses a non-power-based fencing agent), you must set the meta option **provides=unfencing** when creating the **stonith** device. This ensures that a fenced node is unfenced before the node is rebooted and the cluster services are started on the node.

Setting the **provides=unfencing** meta option is not necessary when configuring a power-based fence device, since the device itself is providing power to the node in order for it to boot (and attempt to rejoin the cluster). The act of booting in this case implies that unfencing occurred.

The following command configures a stonith device named **my-scsi-shooter** that uses the **fence\_scsi** fence agent, enabling unfencing for the device.

```
pcs stonith create my-scsi-shooter fence_scsi devices=/dev/sda meta provides=unfencing
```

## 5.6. Displaying Fencing Devices

The following command shows all currently configured fencing devices. If a *stonith\_id* is specified, the command shows the options for that configured stonith device only. If the **--full** option is specified, all configured stonith options are displayed.

```
pcs stonith show [stonith_id] [--full]
```

## 5.7. Modifying and Deleting Fencing Devices

Use the following command to modify or add options to a currently configured fencing device.

```
pcs stonith update stonith_id [stonith_device_options]
```

Use the following command to remove a fencing device from the current configuration.

```
pcs stonith delete stonith_id
```

## 5.8. Managing Nodes with Fence Devices

You can fence a node manually with the following command. If you specify **--off** this will use the **off** API call to stonith which will turn the node off instead of rebooting it.

```
pcs stonith fence node [--off]
```

You can confirm whether a specified node is currently powered off with the following command.



### Note

If the node you specify is still running the cluster software or services normally controlled by the cluster, data corruption/cluster failure will occur.

```
pcs stonith confirm node
```

## 5.9. Additional Fencing Configuration Options

[Table 5.2, “Advanced Properties of Fencing Devices”](#). summarizes additional properties you can set for fencing devices. Note that these properties are for advanced use only.

**Table 5.2. Advanced Properties of Fencing Devices**

Field	Type	Default	Description
<b>pcmk_host_argument</b>	string	port	An alternate parameter to supply instead of port. Some devices do not support the standard port parameter or may provide additional ones. Use this to specify an alternate, device-specific, parameter that should indicate the machine to be fenced. A value of <b>none</b> can be used to tell the cluster not to supply any additional parameters.
<b>pcmk_reboot_action</b>	string	reboot	An alternate command to run instead of <b>reboot</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the reboot action.

Field	Type	Default	Description
<code>pcmk_reboot_timeout</code>	time	60s	Specify an alternate timeout to use for reboot actions instead of <code>stonith-timeout</code> . Some devices need much more/less time to complete than normal. Use this to specify an alternate, device-specific, timeout for reboot actions.
<code>pcmk_reboot_retries</code>	integer	2	The maximum number of times to retry the <code>reboot</code> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries reboot actions before giving up.
<code>pcmk_off_action</code>	string	off	An alternate command to run instead of <code>off</code> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the off action.
<code>pcmk_off_timeout</code>	time	60s	Specify an alternate timeout to use for off actions instead of <code>stonith-timeout</code> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for off actions.
<code>pcmk_off_retries</code>	integer	2	The maximum number of times to retry the off command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries off actions before giving up.
<code>pcmk_list_action</code>	string	list	An alternate command to run instead of <code>list</code> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the list action.
<code>pcmk_list_timeout</code>	time	60s	Specify an alternate timeout to use for list actions instead of <code>stonith-timeout</code> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for list actions.

Field	Type	Default	Description
<b>pcmk_list_retries</b>	integer	2	The maximum number of times to retry the <b>list</b> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries list actions before giving up.
<b>pcmk_monitor_action</b>	string	monitor	An alternate command to run instead of <b>monitor</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the monitor action.
<b>pcmk_monitor_timeout</b>	time	60s	Specify an alternate timeout to use for monitor actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for monitor actions.
<b>pcmk_monitor_retries</b>	integer	2	The maximum number of times to retry the <b>monitor</b> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries monitor actions before giving up.
<b>pcmk_status_action</b>	string	status	An alternate command to run instead of <b>status</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the status action.
<b>pcmk_status_timeout</b>	time	60s	Specify an alternate timeout to use for status actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for status actions.

Field	Type	Default	Description
<code>pcmk_status_retries</code>	integer	2	The maximum number of times to retry the status command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries status actions before giving up.

## 5.10. Configuring Fencing Levels

Pacemaker supports fencing nodes with multiple devices through a feature called fencing topologies. To implement topologies, create the individual devices as you normally would and then define one or more fencing levels in the fencing topology section in the configuration.

- » Each level is attempted in ascending numeric order, starting at 1.
- » If a device fails, processing terminates for the current level. No further devices in that level are exercised and the next level is attempted instead.
- » If all devices are successfully fenced, then that level has succeeded and no other levels are tried.
- » The operation is finished when a level has passed (success), or all levels have been attempted (failed).

Use the following command to add a fencing level to a node. The devices are given as a comma-separated list of stonith ids, which are attempted for the node at that level.

```
pcs stonith level add level node devices
```

The following command lists all of the fencing levels that are currently configured.

```
pcs stonith level
```

In the following example, there are two fence devices configured for node **rh7-2**: an ilo fence device called **my\_il0** and an apc fence device called **my\_apc**. These commands sets up fence levels so that if the device **my\_il0** fails and is unable to fence the node, then Pacemaker will attempt to use the device **my\_apc**. This example also shows the output of the **pcs stonith level** command after the levels are configured.

```
# pcs stonith level add 1 rh7-2 my_il0
# pcs stonith level add 2 rh7-2 my_apc
# pcs stonith level
Node: rh7-2
  Level 1 - my_il0
  Level 2 - my_apc
```

The following command removes the fence level for the specified node and devices. If no nodes or devices are specified then the fence level you specify is removed from all nodes.

```
pcs stonith level remove level [node_id] [stonith_id] ... [stonith_id]
```

The following command clears the fence levels on the specified node or stonith id. If you do not specify a node or stonith id, all fence levels are cleared.

```
pcs stonith level clear [node|stonith_id(s)]
```

If you specify more than one stonith id, they must be separated by a comma and no spaces, as in the following example.

```
# pcs stonith level clear dev_a,dev_b
```

The following command verifies that all fence devices and nodes specified in fence levels exist.

```
pcs stonith level verify
```

## 5.11. Configuring Fencing for Redundant Power Supplies

When configuring fencing for redundant power supplies, the cluster must ensure that when attempting to reboot a host, both power supplies are turned off before either power supply is turned back on.

If the node never completely loses power, the node may not release its resources. This opens up the possibility of nodes accessing these resources simultaneously and corrupting them.

Prior to Red Hat Enterprise Linux 7.2, you needed to explicitly configure different versions of the devices which used either the 'on' or 'off' actions. Since Red Hat Enterprise Linux 7.2, it is now only required to define each device once and to specify that both are required to fence the node, as in the following example.

```
# pcs stonith create apc1 fence_apc_snmp ipaddr=apc1.example.com login=user  
passwd='7a4D#1j!pz864'  
pcmk_host_map="node1.example.com:1;node2.example.com:2"  
  
# pcs stonith create apc2 fence_apc_snmp ipaddr=apc2.example.com login=user  
passwd='7a4D#1j!pz864'  
pcmk_host_map="node1.example.com:1;node2.example.com:2"  
  
# pcs stonith level add 1 node1.example.com apc1,apc2  
# pcs stonith level add 1 node2.example.com apc1,apc2
```

# Chapter 6. Configuring Cluster Resources

This chapter provides information on configuring resources in a cluster.

## 6.1. Resource Creation

Use the following command to create a cluster resource.

```
pcs resource create resource_id standard:provider:type|type [resource options]
[op operation_action operation_options [operation_action operation_options]...]
[meta meta_options...] [-clone clone_options | 
--master master_options | --group group_name
[--before resource_id | --after resource_id]] [--disabled]
```

When you specify the **--group** option, the resource is added to the resource group named. If the group does not exist, this creates the group and adds this resource to the group. For information on resource groups, refer to [Section 6.5, “Resource Groups”](#).

The **--before** and **--after** options specify the position of the added resource relative to a resource that already exists in a resource group.

Specifying the **--disabled** option indicates that the resource is not started automatically.

The following command creates a resource with the name **VirtualIP** of standard **ocf**, provider **heartbeat**, and type **IPaddr2**. The floating address of this resource is 192.168.0.120, the system will check whether the resource is running every 30 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.120
cidr_netmask=24 op monitor interval=30s
```

Alternately, you can omit the *standard* and *provider* fields and use the following command. This will default to a standard of **ocf** and a provider of **heartbeat**.

```
# pcs resource create VirtualIP IPaddr2 ip=192.168.0.120 cidr_netmask=24 op
monitor interval=30s
```

Use the following command to delete a configured resource.

```
pcs resource delete resource_id
```

For example, the following command deletes an existing resource with a resource ID of **VirtualIP**

```
# pcs resource delete VirtualIP
```

- » For information on the *resource\_id*, *standard*, *provider*, and *type* fields of the **pcs resource create** command, refer to [Section 6.2, “Resource Properties”](#).
- » For information on defining resource parameters for individual resources, refer to [Section 6.3, “Resource-Specific Parameters”](#).
- » For information on defining resource meta options, which are used by the cluster to decide how a resource should behave, refer to [Section 6.4, “Resource Meta Options”](#).

- For information on defining the operations to perform on a resource, refer to [Section 6.6, “Resource Operations”](#).
- Specifying the **--clone** creates a clone resource. Specifying the **--master** creates a master/slave resource. For information on resource clones and resources with multiple modes, refer to [Chapter 9, Advanced Resource types](#).

## 6.2. Resource Properties

The properties that you define for a resource tell the cluster which script to use for the resource, where to find that script and what standards it conforms to. [Table 6.1, “Resource Properties”](#) describes these properties.

**Table 6.1. Resource Properties**

Field	Description
resource_id	Your name for the resource
standard	The standard the script conforms to. Allowed values: <b>ocf</b> , <b>service</b> , <b>upstart</b> , <b>systemd</b> , <b>lsb</b> , <b>stonith</b>
type	The name of the Resource Agent you wish to use, for example <b>IPaddr</b> or <b>Filesystem</b>
provider	The OCF spec allows multiple vendors to supply the same resource agent. Most of the agents shipped by Red Hat use <b>heartbeat</b> as the provider.

[Table 6.2, “Commands to Display Resource Properties”](#). summarizes the commands that display the available resource properties.

**Table 6.2. Commands to Display Resource Properties**

pcs Display Command	Output
<b>pcs resource list</b>	Displays a list of all available resources.
<b>pcs resource standards</b>	Displays a list of available resources agent standards.
<b>pcs resource providers</b>	Displays a list of available resources agent providers.
<b>pcs resource list string</b>	Displays a list of available resources filtered by the specified string. You can use this command to display resources filtered by the name of a standard, a provider, or a type.

## 6.3. Resource-Specific Parameters

For any individual resource, you can use the following command to display the parameters you can set for that resource.

```
# pcs resource describe standard:provider:type|type
```

For example, the following command displays the parameters you can set for a resource of type **LVM**.

```
# pcs resource describe LVM
Resource options for: LVM
volgrpname (required): The name of volume group.
exclusive: If set, the volume group will be activated exclusively.
```

`partial_activation`: If set, the volume group will be activated even only partial of the physical volumes available. It helps to set to true, when you are using mirroring logical volumes.

## 6.4. Resource Meta Options

In addition to the resource-specific parameters, you can configure additional resource options for any resource. These options are used by the cluster to decide how your resource should behave. [Table 6.3, “Resource Meta Options”](#) describes this options.

**Table 6.3. Resource Meta Options**

Field	Default	Description
<code>priority</code>	0	If not all resources can be active, the cluster will stop lower priority resources in order to keep higher priority ones active.
<code>target-role</code>	<code>Started</code>	What state should the cluster attempt to keep this resource in? Allowed values:  * <i>Stopped</i> - Force the resource to be stopped  * <i>Started</i> - Allow the resource to be started (In the case of multistate resources, they will not promoted to master)  * <i>Master</i> - Allow the resource to be started and, if appropriate, promoted
<code>is-managed</code>	<code>true</code>	Is the cluster allowed to start and stop the resource? Allowed values: <code>true</code> , <code>false</code>
<code>resource-stickiness</code>	0	Value to indicate how much the resource prefers to stay where it is.

Field	Default	Description
<b>requires</b>	Calculated	<p>Indicates under what conditions the resource can be started.</p> <p>Defaults to <b>fencing</b> except under the conditions noted below. Possible values:</p> <ul style="list-style-type: none"> <li>* <b>nothing</b> - The cluster can always start the resource.</li> <li>* <b>quorum</b> - The cluster can only start this resource if a majority of the configured nodes are active. This is the default value if <b>stonith-enabled</b> is <b>false</b> or the resource's <b>standard</b> is <b>stonith</b>.</li> <li>* <b>fencing</b> - The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been powered off.</li> <li>* <b>unfencing</b> - The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been powered off <i>and</i> only on nodes that have been <i>unfenced</i>. This is the default value if the <b>provides=unfencing stonith</b> meta option has been set for a fencing device. For information on the <b>provides=unfencing stonith</b> meta option, see <a href="#">Section 5.5, “Configuring Storage-Based Fence Devices with unfencing”</a>.</li> </ul>
<b>migration-threshold</b>	<b>INFINITY</b> (disabled)	How many failures may occur for this resource on a node, before this node is marked ineligible to host this resource. For information on configuring the <b>migration-threshold</b> option, refer to <a href="#">Section 8.2, “Moving Resources Due to Failure”</a> .
<b>failure-timeout</b>	<b>0</b> (disabled)	Used in conjunction with the <b>migration-threshold</b> option, indicates how many seconds to wait before acting as if the failure had not occurred, and potentially allowing the resource back to the node on which it failed. For information on configuring the <b>failure-timeout</b> option, refer to <a href="#">Section 8.2, “Moving Resources Due to Failure”</a> .
<b>multiple-active</b>	<b>stop_start</b>	<p>What should the cluster do if it ever finds the resource active on more than one node. Allowed values:</p> <ul style="list-style-type: none"> <li>* <b>block</b> - mark the resource as unmanaged</li> <li>* <b>stop_only</b> - stop all active instances and leave them that way</li> <li>* <b>stop_start</b> - stop all active instances and start the resource in one location only</li> </ul>

To change the default value of a resource option, use the following command.

```
pcs resource defaults options
```

For example, the following command resets the default value of **resource-stickiness** to 100.

```
# pcs resource defaults resource-stickiness=100
```

Omitting the *options* parameter from the **pcs resource defaults** displays a list of currently configured default values for resource options. The following example shows the output of this command after you have reset the default value of **resource-stickiness** to 100.

```
# pcs resource defaults
resource-stickiness:100
```

Whether you have reset the default value of a resource meta option or not, you can set a resource option for a particular resource to a value other than the default when you create the resource. The following shows the format of the **pcs resource create** command you use when specifying a value for a resource meta option.

```
pcs resource create resource_id standard:provider:type|type [resource options] [meta meta_options...]
```

For example, the following command creates a resource with a **resource-stickiness** value of 50.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.120
cidr_netmask=24 meta resource-stickiness=50
```

You can also set the value of a resource meta option for an existing resource, group, cloned resource, or master resource with the following command.

```
pcs resource meta resource_id | group_id | clone_id | master_id meta_options
```

In the following example, there is an existing resource named **dummy\_resource**. This command sets the **failure-timeout** meta option to 20 seconds, so that the resource can attempt to restart on the same node in 20 seconds.

```
# pcs resource meta dummy_resource failure-timeout=20s
```

After executing this command, you can display the values for the resource to verify that **failure-timeout=20s** is set.

```
# pcs resource show dummy_resource
Resource: dummy_resource (class=ocf provider=heartbeat type=Dummy)
Meta Attrs: failure-timeout=20s
Operations: start interval=0s timeout=20 (dummy_resource-start-timeout-20)
           stop interval=0s timeout=20 (dummy_resource-stop-timeout-20)
           monitor interval=10 timeout=20 (dummy_resource-monitor-interval-10)
```

For information on resource clone meta options, see [Section 9.1, “Resource Clones”](#). For information on resource master meta options, see [Section 9.2, “MultiState Resources: Resources That Have Multiple Modes”](#).

## 6.5. Resource Groups

One of the most common elements of a cluster is a set of resources that need to be located together, start sequentially, and stop in the reverse order. To simplify this configuration, Pacemaker supports the concept of groups.

You create a resource group with the following command, specifying the resources to include in the group. If the group does not exist, this command creates the group. If the group exists, this command adds additional resources to the group. The resources will start in the order you specify them with this command, and will stop in the reverse order of their starting order.

```
pcs resource group add group_name resource_id [resource_id] ... [resource_id]
[--before resource_id] [--after resource_id]
```

You can use the **--before** and **--after** options of this command to specify the position of the added resources relative to a resource that already exists in the group.

You can also add a new resource to an existing group when you create the resource, using the following command. The resource you create is added to the group named *group\_name*.

```
pcs resource create resource_id standard:provider:type|type [resource_options] [op operation_action
operation_options] --group group_name
```

You remove a resource from a group with the following command. If there are no resources in the group, this command removes the group itself.

```
pcs resource group remove group_name resource_id...
```

The following command lists all currently configured resource groups.

```
pcs resource group list
```

The following example creates a resource group named **shortcut** that contains the existing resources **IPAddr** and **Email**.

```
# pcs resource group add shortcut IPAddr Email
```

There is no limit to the number of resources a group can contain. The fundamental properties of a group are as follows.

- » Resources are started in the order in which you specify them (in this example, **IPAddr** first, then **Email**).
- » Resources are stopped in the reverse order in which you specify them. (**Email** first, then **IPAddr**).

If a resource in the group cannot run anywhere, then no resource specified after that resource is allowed to run.

- » If **IPAddr** cannot run anywhere, neither can **Email**.
- » If **Email** cannot run anywhere, however, this does not affect **IPAddr** in any way.

Obviously as the group grows bigger, the reduced configuration effort of creating resource groups can become significant.

## 6.5.1. Group Options

A resource group inherits the following options from the resources that it contains: **priority**, **target-role**, **is-managed**. For information on resource options, refer to [Table 6.3, “Resource Meta Options”](#).

## 6.5.2. Group Stickiness

Stickiness, the measure of how much a resource wants to stay where it is, is additive in groups. Every active resource of the group will contribute its stickiness value to the group’s total. So if the default **resource-stickiness** is 100, and a group has seven members, five of which are active, then the group as a whole will prefer its current location with a score of 500.

## 6.6. Resource Operations

To ensure that resources remain healthy, you can add a monitoring operation to a resource’s definition. If you do not specify a monitoring operation for a resource, by default the **pcs** command will create a monitoring operation, with an interval that is determined by the resource agent. If the resource agent does not provide a default monitoring interval, the **pcs** command will create a monitoring operation with an interval of 60 seconds.

[Table 6.4, “Properties of an Operation”](#) summarizes the properties of a resource monitoring operation.

**Table 6.4. Properties of an Operation**

Field	Description
<b>id</b>	Unique name for the action. The system assigns this when you configure an operation.
<b>name</b>	The action to perform. Common values: <b>monitor</b> , <b>start</b> , <b>stop</b>
<b>interval</b>	How frequently (in seconds) to perform the operation. Default value: <b>0</b> , meaning never.
<b>timeout</b>	How long to wait before declaring the action has failed. If you find that your system includes a resource that takes a long time to start or stop or perform a non-recurring monitor action at startup, and requires more time than the system allows before declaring that the start action has failed, you can increase this value from the default of 20 or the value of <b>timeout</b> in “op defaults”.
<b>on-fail</b>	<p>The action to take if this action ever fails. Allowed values:</p> <ul style="list-style-type: none"> <li>* <b>ignore</b> - Pretend the resource did not fail</li> <li>* <b>block</b> - Do not perform any further operations on the resource</li> <li>* <b>stop</b> - Stop the resource and do not start it elsewhere</li> <li>* <b>restart</b> - Stop the resource and start it again (possibly on a different node)</li> <li>* <b>fence</b> - STONITH the node on which the resource failed</li> <li>* <b>standby</b> - Move <i>all</i> resources away from the node on which the resource failed</li> </ul> <p>The default for the <b>stop</b> operation is <b>fence</b> when STONITH is enabled and <b>block</b> otherwise. All other operations default to <b>restart</b>.</p>
<b>enabled</b>	If <b>false</b> , the operation is treated as if it does not exist. Allowed values: <b>true</b> , <b>false</b>

You can configure monitoring operations when you create a resource, using the following command.

```
pcs resource create resource_id standard:provider:type|type [resource_options] [op operation_action operation_options [operation_type operation_options]...]
```

For example, the following command creates an **IPaddr2** resource with a monitoring operation. The new resource is called **VirtualIP** with an IP address of 192.168.0.99 and a netmask of 24 on **eth2**. A monitoring operation will be performed every 30 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99
cidr_netmask=24 nic=eth2 op monitor interval=30s
```

Alternately, you can add a monitoring operation to an existing resource with the following command.

```
pcs resource op add resource_id operation_action [operation_properties]
```

Use the following command to delete a configured resource operation.

```
pcs resource op remove resource_id operation_name operation_properties
```



### Note

You must specify the exact operation properties to properly remove an existing operation.

To change the values of a monitoring option, you can update the resource. For example, you can create a **VirtualIP** with the following command.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99
cidr_netmask=24 nic=eth2
```

By default, this command creates these operations.

```
Operations: start interval=0s timeout=20s (VirtualIP-start-timeout-20s)
stop interval=0s timeout=20s (VirtualIP-stop-timeout-20s)
monitor interval=10s timeout=20s (VirtualIP-monitor-interval-10s)
```

To change the stop timeout operation, execute the following command.

```
# pcs resource update VirtualIP stop interval=0s timeout=40s
# pcs resource show VirtualIP
Resource: VirtualIP (class=ocf provider=heartbeat type=IPaddr2)
Attributes: ip=192.168.0.99 cidr_netmask=24 nic=eth2
Operations: start interval=0s timeout=20s (VirtualIP-start-timeout-20s)
monitor interval=10s timeout=20s (VirtualIP-monitor-interval-10s)
stop interval=0s timeout=40s (VirtualIP-name-stop-interval-0s-timeout-40s)
```

To set global default values for monitoring operations, use the following command.

```
pcs resource op defaults [options]
```

For example, the following command sets a global default of a **timeout** value of 240 seconds for all monitoring operations.

```
# pcs resource op defaults timeout=240s
```

To display the currently configured default values for monitoring operations, do not specify any options when you execute the **pcs resource op defaults** command.

For example, following command displays the default monitoring operation values for a cluster which has been configured with a **timeout** value of 240 seconds.

```
# pcs resource op defaults
timeout: 240s
```

## 6.7. Displaying Configured Resources

To display a list of all configured resources, use the following command.

```
pcs resource show
```

For example, if your system is configured with a resource named **VirtualIP** and a resource named **WebSite**, the **pcs resource show** command yields the following output.

```
# pcs resource show
VirtualIP (ocf::heartbeat:IPaddr2): Started
WebSite (ocf::heartbeat:apache): Started
```

To display a list of all configured resources and the parameters configured for those resources, use the **--full** option of the **pcs resource show** command, as in the following example.

```
# pcs resource show --full
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
  Attributes: ip=192.168.0.120 cidr_netmask=24
  Operations: monitor interval=30s
Resource: WebSite (type=apache class=ocf provider=heartbeat)
  Attributes: statusurl=http://localhost/server-status configfile=/etc/httpd/conf/httpd.conf
  Operations: monitor interval=1min
```

To display the configured parameters for a resource, use the following command.

```
pcs resource show resource_id
```

For example, the following command displays the currently configured parameters for resource **VirtualIP**.

```
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
  Attributes: ip=192.168.0.120 cidr_netmask=24
  Operations: monitor interval=30s
```

## 6.8. Modifying Resource Parameters

To modify the parameters of a configured resource, use the following command.

```
pcs resource update resource_id [resource_options]
```

The following sequence of commands show the initial values of the configured parameters for resource **VirtualIP**, the command to change the value of the **ip** parameter, and the values following the update command.

```
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.168.0.120 cidr_netmask=24
Operations: monitor interval=30s
# pcs resource update VirtualIP ip=192.169.0.120
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.169.0.120 cidr_netmask=24
Operations: monitor interval=30s
```

## 6.9. Multiple Monitoring Operations

You can configure a single resource with as many monitor operations as a resource agent supports. In this way you can do a superficial health check every minute and progressively more intense ones at higher intervals.

### Note

When configuring multiple monitor operations, you must ensure that no two operations are performed at the same interval.

To configure additional monitoring operations for a resource that supports more in-depth checks at different levels, you add an **OCF\_CHECK\_LEVEL=n** option.

For example, if you configure the following **IPaddr2** resource, by default this creates a monitoring operation with an interval of 10 seconds and a timeout value of 20 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99
cidr_netmask=24 nic=eth2
```

If the Virtual IP supports a different check with a depth of 10, the following command causes Packemaker to perform the more advanced monitoring check every 60 seconds in addition to the normal Virtual IP check every 10 seconds. (As noted, you should not configure the additional monitoring operation with a 10-second interval as well.)

```
# pcs resource op add VirtualIP monitor interval=60s OCF_CHECK_LEVEL=10
```

## 6.10. Enabling and Disabling Cluster Resources

The following command enables the resource specified by **resource\_id**.

```
pcs resource enable resource_id
```

The following command disables the resource specified by **resource\_id**.

```
pcs resource disable resource_id
```

## 6.11. Cluster Resources Cleanup

If a resource has failed, a failure message appears when you display the cluster status. If you resolve that resource, you can clear that failure status with the **pcs resource cleanup** command. This command resets the resource status and **failcount**, telling the cluster to forget the operation history of a resource and re-detect its current state.

The following command cleans up the resource specified by *resource\_id*.

```
pcs resource cleanup resource_id
```

If you do not specify a *resource\_id*, this command resets the resource status and **failcount** for all resources.

## Chapter 7. Resource Constraints

You can determine the behavior of a resource in a cluster by configuring constraints for that resource. You can configure the following categories of constraints:

- **location** constraints — A location constraint determines which nodes a resource can run on. Location constraints are described in [Section 7.1, “Location Constraints”](#).
- **order** constraints — An order constraint determines the order in which the resources run. Order constraints are described in [Section 7.2, “Order Constraints”](#).
- **colocation** constraints — A colocation constraint determines where resources will be placed relative to other resources. Colocation constraints are described in [Section 7.3, “Colocation of Resources”](#).

As a shorthand for configuring a set of constraints that will locate a set of resources together and ensure that the resources start sequentially and stop in reverse order, Pacemaker supports the concept of resource groups. For information on resource groups, see [Section 6.5, “Resource Groups”](#).

### 7.1. Location Constraints

Location constraints determine which nodes a resource can run on. You can configure location constraints to determine whether a resource will prefer or avoid a specified node.

[Table 7.1, “Location Constraint Options”](#). summarizes the options for configuring location constraints.

**Table 7.1. Location Constraint Options**

Field	Description
<b>rsc</b>	A resource name
<b>node</b>	A node's name
<b>score</b>	Value to indicate the preference for whether a resource should run on or avoid a node.  A Value of <b>INFINITY</b> changes "should" to "must"; <b>INFINITY</b> is the default <b>score</b> value for a resource location constraint.

Field	Description
<b>resource-discovery</b>	<p>Value to indicate the preference for whether Pacemaker should perform resource discovery on this node for the specified resource. Limiting resource discovery to a subset of nodes the resource is physically capable of running on can significantly boost performance when a large set of nodes is present. When <code>pacemaker_remote</code> is in use to expand the node count into the hundreds of nodes range, this option should be considered. Possible values include:</p> <p><b>always</b>: Always perform resource discovery for the specified resource on this node.</p> <p><b>never</b>: Never perform resource discovery for the specified resource on this node.</p> <p><b>exclusive</b>: Perform resource discovery for the specified resource only on this node (and other nodes similarly marked as <b>exclusive</b>). Multiple location constraints using <b>exclusive</b> discovery for the same resource across different nodes creates a subset of nodes <b>resource-discovery</b> is exclusive to. If a resource is marked for <b>exclusive</b> discovery on one or more nodes, that resource is only allowed to be placed within that subset of nodes.</p> <p>Note that setting this option to <b>never</b> or <b>exclusive</b> allows the possibility for the resource to be active in those locations without the cluster's knowledge. This can lead to the resource being active in more than one location if the service is started outside the cluster's control (for example, by <code>systemd</code> or by an administrator). This can also occur if the <b>resource-discovery</b> property is changed while part of the cluster is down or suffering split-brain, or if the <b>resource-discovery</b> property is changed for a resource and node while the resource is active on that node. For this reason, using this option is appropriate only when you have more than eight nodes and there is a way to guarantee that the resource can run only in a particular location (for example, when the required software is not installed anywhere else).</p> <p><b>always</b> is the default <b>resource-discovery</b> value for a resource location constraint.</p>

The following command creates a location constraint for a resource to prefer the specified node or nodes.

```
pcs constraint location rsc prefers node[=score] ...
```

The following command creates a location constraint for a resource to avoid the specified node or nodes.

```
pcs constraint location rsc avoids node[=score] ...
```

There are two alternative strategies for specifying which nodes a resources can run on:

- » Opt-In Clusters — Configure a cluster in which, by default, no resource can run anywhere and then selectively enable allowed nodes for specific resources. The procedure for configuring an opt-in cluster is described in [Section 7.1.1, “Configuring an “Opt-In” Cluster”](#).
- » Opt-Out Clusters — Configure a cluster in which, by default, all resources can run anywhere and then create location constraints for resources that are not allowed to run on specific nodes. The procedure for configuring an opt-out cluster is described in [Section 7.1.2, “Configuring an “Opt-Out” Cluster”](#).

Whether you should choose to configure an opt-in or opt-out cluster depends both on your personal preference and the make-up of your cluster. If most of your resources can run on most of the nodes, then an opt-out arrangement is likely to result in a simpler configuration. On the other hand, if most resources can only run on a small subset of nodes an opt-in configuration might be simpler.

### 7.1.1. Configuring an “Opt-In” Cluster

To create an opt-in cluster, set the **symmetric-cluster** cluster property to **false** to prevent resources from running anywhere by default.

```
# pcs property set symmetric-cluster=false
```

Enable nodes for individual resources. The following commands configure location constraints so that the resource **Webserver** prefers node **example-1**, the resource **Database** prefers node **example-2**, and both resources can fail over to node **example-3** if their preferred node fails.

```
# pcs constraint location Webserver prefers example-1=200
# pcs constraint location Webserver prefers example-3=0
# pcs constraint location Database prefers example-2=200
# pcs constraint location Database prefers example-3=0
```

### 7.1.2. Configuring an “Opt-Out” Cluster

To create an opt-out cluster, set the **symmetric-cluster** cluster property to **true** to allow resources to run everywhere by default.

```
# pcs property set symmetric-cluster=true
```

The following commands will then yield a configuration that is equivalent to the example in [Section 7.1.1, “Configuring an “Opt-In” Cluster”](#). Both resources can fail over to node **example-3** if their preferred node fails, since every node has an implicit score of 0.

```
# pcs constraint location Webserver prefers example-1=200
# pcs constraint location Webserver avoids example-2=INFINITY
# pcs constraint location Database avoids example-1=INFINITY
# pcs constraint location Database prefers example-2=200
```

Note that it is not necessary to specify a score of INFINITY in these commands, since that is the default value for the score.

## 7.2. Order Constraints

Order constraints determine the order in which the resources run. You can configure an order constraint to determine the order in which resources start and stop.

Use the following command to configure an order constraint.

```
pcs constraint order [action] resource_id then [action] resource_id [options]
```

[Table 7.2, “Properties of an Order Constraint”](#). summarizes the properties and options for configuring order constraints.

**Table 7.2. Properties of an Order Constraint**

Field	Description
resource_id	The name of a resource on which an action is performed.
action	<p>The action to perform on a resource. Possible values of the <i>action</i> property are as follows:</p> <ul style="list-style-type: none"> <li>* <b>start</b> - Start the resource.</li> <li>* <b>stop</b> - Stop the resource.</li> <li>* <b>promote</b> - Promote the resource from a slave resource to a master resource.</li> <li>* <b>demote</b> - Demote the resource from a master resource to a slave resource.</li> </ul> <p>If no action is specified, the default action is <b>start</b>. For information on master and slave resources, refer to <a href="#">Section 9.2, “MultiState Resources: Resources That Have Multiple Modes”</a>.</p>
<b>kind</b> option	<p>How to enforce the constraint. The possible values of the <b>kind</b> option are as follows:</p> <ul style="list-style-type: none"> <li>* <b>Optional</b> - Only applies if both resources are executing the specified action. For information on optional ordering, refer to <a href="#">Section 7.2.2, “Advisory Ordering”</a>.</li> <li>* <b>Mandatory</b> - Always (default value). If the first resource you specified is stopping or cannot be started, the second resource you specified must be stopped. For information on mandatory ordering, refer to <a href="#">Section 7.2.1, “Mandatory Ordering”</a>.</li> <li>* <b>Serialize</b> - Ensure that no two stop/start actions occur concurrently for a set of resources.</li> </ul>
<b>symmetrical</b> option	If true, which is the default, stop the resources in the reverse order. Default value: <b>true</b>

## 7.2.1. Mandatory Ordering

A mandatory constraints indicates that the second resource you specify cannot run without the first resource you specify being active. This is the default value of the **kind** option. Leaving the default value ensures that the second resource you specify will react when the first resource you specify changes state.

- If the first resource you specified resource was running and is stopped, the second resource you specified will also be stopped (if it is running).

- » If the first resource you specified resource was not running and cannot be started, the second resource you specified will be stopped (if it is running).
- » If the first resource you specified is (re)started while the second resource you specified is running, the second resource you specified will be stopped and restarted.

## 7.2.2. Advisory Ordering

When the **kind=Optional** option is specified for an order constraint, the constraint is considered optional and only applies if both resources are executing the specified actions. Any change in state by the first resource you specify will have no effect on the second resource you specify.

The following command configures an advisory ordering constraint for the resources named **VirtualIP** and **dummy\_resource**.

```
# pcs constraint order VirtualIP then dummy_resource kind=Optional
```

## 7.2.3. Ordered Resource Sets

A common situation is for an administrator to create a chain of ordered resources, where, for example, resource A starts before resource B which starts before resource C. If your configuration requires that you create a set of resources that is colocated and started in order, you can configure a resource group that contains those resources, as described in [Section 6.5, “Resource Groups”](#). If, however, you need to configure resources to start in order and the resources are not necessarily colocated, you can create an order constraint on a set or sets of resources with the **pcs constraint order set** command.

You can set the following options for a set of resources with the **pcs constraint order set** command.

- » **sequential**, which can be set to **true** or **false** to indicate whether the set of resources must be ordered relative to each other.
- Setting **sequential** to **false** allows a set to be ordered relative to other sets in the ordering constraint, without its members being ordered relative to each other. Therefore, this option makes sense only if multiple sets are listed in the constraint; otherwise, the constraint has no effect.
- » **require-all**, which can be set to **true** or **false** to indicate whether all of the resources in the set must be active before continuing. Setting **require-all** to **false** means that only one resource in the set needs to be started before continuing on to the next set. Setting **require-all** to **false** has no effect unless used in conjunction with unordered sets, which are sets for which **sequential** is set to **false**.
  - » **action**, which can be set to **start**, **promote**, **demote** or **stop**, as described in [Table 7.2, “Properties of an Order Constraint”](#).

You can set the following constraint options for a set of resources following the **setoptions** parameter of the **pcs constraint order set** command.

- » **id**, to provide a name for the constraint you are defining.
- » **score**, to indicate the degree of preference for this constraint. For information on this option, see [Table 7.3, “Properties of a Colocation Constraint”](#).

```
pcs constraint order set resource1 resource2 [resourceN]... [options] [set resourceX resourceY ... [options]] [setoptions [constraint_options]]
```

If you have three resources named **D1**, **D2**, and **D3**, the following command configures them as an ordered resource set.

```
# pcs constraint order set D1 D2 D3
```

### 7.2.4. Removing Resources From Ordering Constraints

Use the following command to remove resources from any ordering constraint.

```
pcs constraint order remove resource1 [resourceN]...
```

## 7.3. Colocation of Resources

A colocation constraint determines that the location of one resource depends on the location of another resource.

There is an important side effect of creating a colocation constraint between two resources: it affects the order in which resources are assigned to a node. This is because you cannot place resource A relative to resource B unless you know where resource B is. So when you are creating colocation constraints, it is important to consider whether you should colocate resource A with resource B or resource B with resource A.

Another thing to keep in mind when creating colocation constraints is that, assuming resource A is collocated with resource B, the cluster will also take into account resource A's preferences when deciding which node to choose for resource B.

The following command creates a colocation constraint.

```
pcs constraint colocation add [master|slave] source_resource with [master|slave] target_resource  
[score] [options]
```

For information on master and slave resources, see [Section 9.2, “MultiState Resources: Resources That Have Multiple Modes”](#).

[Table 7.3, “Properties of a Colocation Constraint”](#). summarizes the properties and options for configuring colocation constraints.

**Table 7.3. Properties of a Colocation Constraint**

Field	Description
source_resource	The colocation source. If the constraint cannot be satisfied, the cluster may decide not to allow the resource to run at all.
target_resource	The colocation target. The cluster will decide where to put this resource first and then decide where to put the source resource.
score	Positive values indicate the resource should run on the same node. Negative values indicate the resources should not run on the same node. A value of + <b>INFINITY</b> , the default value, indicates that the <i>source_resource</i> must run on the same node as the <i>target_resource</i> . A value of - <b>INFINITY</b> indicates that the <i>source_resource</i> must not run on the same node as the <i>target_resource</i> .

### 7.3.1. Mandatory Placement

Mandatory placement occurs any time the constraint's score is **+INFINITY** or **-INFINITY**. In such cases, if the constraint cannot be satisfied, then the *source\_resource* is not permitted to run. For **score=INFINITY**, this includes cases where the *target\_resource* is not active.

If you need **myresource1** to always run on the same machine as **myresource2**, you would add the following constraint:

```
# pcs constraint colocation add myresource1 with myresource2 score=INFINITY
```

Because **INFINITY** was used, if **myresource2** cannot run on any of the cluster nodes (for whatever reason) then **myresource1** will not be allowed to run.

Alternatively, you may want to configure the opposite, a cluster in which **myresource1** cannot run on the same machine as **myresource2**. In this case use **score=-INFINITY**

```
# pcs constraint colocation add myresource1 with myresource2 score=-INFINITY
```

Again, by specifying **-INFINITY**, the constraint is binding. So if the only place left to run is where **myresource2** already is, then **myresource1** may not run anywhere.

### 7.3.2. Advisory Placement

If mandatory placement is about "must" and "must not", then advisory placement is the "I would prefer if" alternative. For constraints with scores greater than **-INFINITY** and less than **INFINITY**, the cluster will try to accommodate your wishes but may ignore them if the alternative is to stop some of the cluster resources. Advisory colocation constraints can combine with other elements of the configuration to behave as if they were mandatory.

### 7.3.3. Colocating Sets of Resources

If your configuration requires that you create a set of resources that is colocated and started in order, you can configure a resource group that contains those resources, as described in [Section 6.5, “Resource Groups”](#). If, however, you need to colocate a set of resources but the resources do not necessarily need to start in order, you can create a colocation constraint on a set or sets of resources with the **pcs constraint colocation set** command.

You can set the following options for a set of resources with the **pcs constraint colocation set** command.

- **sequential**, which can be set to **true** or **false** to indicate whether the members of the set must be colocated with each other.

Setting **sequential** to **false** allows the members of this set to be colocated with another set listed later in the constraint, regardless of which members of this set are active. Therefore, this option makes sense only if another set is listed after this one in the constraint; otherwise, the constraint has no effect.

- **role**, which can be set to **Stopped**, **Started**, **Master**, or **Slave**. For information on multi-state resources, see [Section 9.2, “MultiState Resources: Resources That Have Multiple Modes”](#).

You can set the following constraint options for a set of resources following the **setoptions** parameter of the **pcs constraint colocation set** command.

- **kind**, to indicate how to enforce the constraint. For information on this option, see [Table 7.2, “Properties of an Order Constraint”](#).
- **symmetrical**, to indicate the order in which to stop the resources. If true, which is the default, stop

the resources in the reverse order. Default value: **true**

- » **id**, to provide a name for the constraint you are defining.

When listing members of a set, each member is colocated with the one before it. For example, "set A B" means "B is colocated with A". However, when listing multiple sets, each set is colocated with the one after it. For example, "set C D sequential=false set A B" means "set C D (where C and D have no relation between each other) is colocated with set A B (where B is colocated with A)".

The following command creates a colocation constraint on a set or sets of resources.

```
pcs constraint colocation set resource1 resource2 [resourceN]... [options] [set resourceX resourceY
... [options]] [setoptions [constraint_options]]
```

### 7.3.4. Removing Colocation Constraints

Use the following command to remove colocation constraints with *source\_resource*.

```
pcs constraint colocation remove source_resource target_resource
```

## 7.4. Displaying Constraints

There are several commands you can use to display constraints that have been configured.

The following command lists all current location, order, and colocation constraints.

```
pcs constraint list|show
```

The following command lists all current location constraints.

- » If **resources** is specified, location constraints are displayed per resource. This is the default behavior.
- » If **nodes** is specified, location constraints are displayed per node.
- » If specific resources or nodes are specified, then only information about those resources or nodes is displayed.

```
pcs constraint location [show resources|nodes [specific nodes|resources]] [--full]
```

The following command lists all current ordering constraints. If the **--full** option is specified, show the internal constraint IDs.

```
pcs constraint order show [--full]
```

The following command lists all current colocation constraints. If the **--full** option is specified, show the internal constraint IDs.

```
pcs constraint colocation show [--full]
```

The following command lists the constraints that reference specific resources.

```
pcs constraint ref resource ...
```

## Chapter 8. Managing Cluster Resources

This chapter describes various commands you can use to manage cluster resources. It provides information on the following procedures.

- » [Section 8.1, “Manually Moving Resources Around the Cluster”](#)
- » [Section 8.2, “Moving Resources Due to Failure”](#)
- » [Section 8.4, “Enabling, Disabling, and Banning Cluster Resources”](#)
- » [Section 8.5, “Disabling a Monitor Operations”](#)

### 8.1. Manually Moving Resources Around the Cluster

You can override the cluster and force resources to move from their current location. There are two occasions when you would want to do this:

- » When a node is under maintenance, and you need to move all resources running on that node to a different node
- » When individually specified resources needs to be moved

To move all resources running on a node to a different node, you put the node in standby mode. For information on putting a cluster node in standby mode, refer to [Section 4.4.5, “Standby Mode”](#).

You can move individually specified resources in either of the following ways.

- » You can use the **pcs resource move** command to move a resource off a node on which it is currently running, as described in [Section 8.1.1, “Moving a Resource from its Current Node”](#).
- » You can use the **pcs resource relocate run** command to move a resource to its preferred node, as determined by current cluster status, constraints, location of resources and other settings. For information on this command, see [Section 8.1.2, “Moving a Resource to its Preferred Node”](#).

#### 8.1.1. Moving a Resource from its Current Node

To move a resource off the node on which it is currently running, use the following command, specifying the *resource\_id* of the node as defined. Specify the **destination\_node** if you want to indicate on which node to run the resource that you are moving.

```
pcs resource move resource_id [destination_node] [--master] [lifetime=lifetime]
```

#### Note

When you execute the **pcs resource move** command, this adds a constraint to the resource to prevent it from running on the node on which it is currently running. You can execute the **pcs resource clear** or the **pcs constraint delete** command to remove the constraint. This does not necessarily move the resources back to the original node; where the resources can run at that point depends on how you have configured your resources initially.

If you specify the **--master** parameter of the **pcs resource ban** command, the scope of the constraint is limited to the master role and you must specify *master\_id* rather than *resource\_id*.

You can optionally configure a **lifetime** parameter for the **pcs resource move** command to indicate a period of time the constraint should remain. You specify the units of a **lifetime** parameter according to the format defined in ISO 8601, which requires that you specify the unit as a capital letter such as Y (for years), M (for months), W (for weeks), D (for days), H (for hours), M (for minutes), and S (for seconds).

To distinguish a unit of minutes(M) from a unit of months(M), you must specify PT before indicating the value in minutes. For example, a **lifetime** parameter of 5M indicates an interval of five months, while a **lifetime** parameter of PT5M indicates an interval of five minutes.

The **lifetime** parameter is checked at intervals defined by the **cluster-recheck-interval** cluster property. By default this value is 15 minutes. If your configuration requires that you check this parameter more frequently, you can reset this value with the following command.

```
pcs property set cluster-recheck-interval=value
```

You can optionally configure a **--wait[=n]** parameter for the **pcs resource ban** command to indicate the number of seconds to wait for the resource to start on the destination node before returning 0 if the resource is started or 1 if the resource has not yet started. If you do not specify n, the default resource timeout will be used.

The following command moves the resource **resource1** to node **example-node2** and prevents it from moving back to the node on which it was originally running for one hour and thirty minutes.

```
pcs resource move resource1 example-node2 lifetime=PT1H30M
```

The following command moves the resource **resource1** to node **example-node2** and prevents it from moving back to the node on which it was originally running for thirty minutes.

```
pcs resource move resource1 example-node2 lifetime=PT30M
```

For information on resource constraints, refer to [Chapter 7, Resource Constraints](#).

### 8.1.2. Moving a Resource to its Preferred Node

After a resource has moved, either due to a failover or to an administrator manually moving the node, it will not necessarily move back to its original node even after the circumstances that caused the failover have been corrected. To relocate resources to their preferred node, use the following command. A preferred node is determined by the current cluster status, constraints, resource location, and other settings and may change over time.

```
pcs resource relocate run [resource1] [resource2] ...
```

If you do not specify any resources, all resource are relocated to their preferred nodes.

This command calculates the preferred node for each resource while ignoring resource stickiness. After calculating the preferred node, it creates location constraints which will cause the resources to move to their preferred nodes. Once the resources have been moved, the constraints are deleted automatically. To remove all constraints created by the **pcs resource relocate run** command, you can enter the **pcs resource relocate clear** command. To display the current status of resources and their optimal node ignoring resource stickiness, enter the **pcs resource relocate show** command.

## 8.2. Moving Resources Due to Failure

When you create a resource, you can configure the resource so that it will move to a new node after a defined number of failures by setting the **migration-threshold** option for that resource. Once the threshold has been reached, this node will no longer be allowed to run the failed resource until:

- » The administrator manually resets the resource's **failcount** using the **pcs resource failcount** command.
- » The resource's **failure-timeout** value is reached.

There is no threshold defined by default.

### Note

Setting a **migration-threshold** for a resource is not the same as configuring a resource for migration, in which the resource moves to another location without loss of state.

The following example adds a migration threshold of 10 to the resource named **dummy\_resource**, which indicates that the resource will move to a new node after 10 failures.

```
# pcs resource meta dummy_resource migration-threshold=10
```

You can add a migration threshold to the defaults for the whole cluster with the following command.

```
# pcs resource defaults migration-threshold=10
```

To determine the resource's current failure status and limits, use the **pcs resource failcount** command.

There are two exceptions to the migration threshold concept; they occur when a resource either fails to start or fails to stop. If the cluster property **start-failure-is-fatal** is set to **true** (which is the default), start failures cause the **failcount** to be set to **INFINITY** and thus always cause the resource to move immediately.

Stop failures are slightly different and crucial. If a resource fails to stop and STONITH is enabled, then the cluster will fence the node in order to be able to start the resource elsewhere. If STONITH is not enabled, then the cluster has no way to continue and will not try to start the resource elsewhere, but will try to stop it again after the failure timeout.

## 8.3. Moving Resources Due to Connectivity Changes

Setting up the cluster to move resources when external connectivity is lost is a two step process.

1. Add a **ping** resource to the cluster. The **ping** resource uses the system utility of the same name to test if a list of machines (specified by DNS host name or IPv4/IPv6 address) are reachable and uses the results to maintain a node attribute called **pingd**.
2. Configure a location constraint for the resource that will move the resource to a different node when connectivity is lost.

[Table 6.1, “Resource Properties”](#) describes the properties you can set for a **ping** resource.

### Table 8.1. Properties of a ping resources

Field	Description
<b>dampen</b>	The time to wait (dampening) for further changes to occur. This prevents a resource from bouncing around the cluster when cluster nodes notice the loss of connectivity at slightly different times.
<b>multiplier</b>	The number of connected ping nodes gets multiplied by this value to get a score. Useful when there are multiple ping nodes configured.
<b>host_list</b>	The machines to contact in order to determine the current connectivity status. Allowed values include resolvable DNS host names, IPv4 and IPv6 addresses. The entries in the host list are space separated.

The following example command creates a **ping** resource that verifies connectivity to **gateway.example.com**. In practice, you would verify connectivity to your network gateway/router. You configure the **ping** resource as a clone so that the resource will run on all cluster nodes.

```
# pcs resource create ping ocf:pacemaker:ping dampen=5s multiplier=1000
host_list=gateway.example.com --clone
```

The following example configures a location constraint rule for the existing resource named **Webserver**. This will cause the **Webserver** resource to move to a host that is able to ping **www.example.com** if the host that it is currently running on cannot ping **www.example.com**

```
# pcs constraint location Webserver rule score=-INFINITY pingd lt 1 or not_defined
pingd
```

## 8.4. Enabling, Disabling, and Banning Cluster Resources

In addition to the **pcs resource move** and **pcs resource relocate** commands described in [Section 8.1, “Manually Moving Resources Around the Cluster”](#), there are a variety of other commands you can use to control the behavior of cluster resources.

You can manually stop a running resource and prevent the cluster from starting it again with the following command. Depending on the rest of the configuration (constraints, options, failures, and so on), the resource may remain started. If you specify the **--wait** option, **pcs** will wait up to 30 seconds (or 'n' seconds, as specified) for the resource to stop and then return 0 if the resource is stopped or 1 if the resource has not stopped.

```
pcs resource disable resource_id [--wait[=n]]
```

You can use the following command to allow the cluster to start a resource. Depending on the rest of the configuration, the resource may remain stopped. If you specify the **--wait** option, **pcs** will wait up to 30 seconds (or 'n' seconds, as specified) for the resource to start and then return 0 if the resource is started or 1 if the resource has not started.

```
pcs resource enable resource_id [--wait[=n]]
```

Use the following command to prevent a resource from running on a specified node, or on the current node if no node is specified.

```
pcs resource ban resource_id [node] [--master] [lifetime=lifetime] [--wait[=n]]
```

Note that when you execute the **pcs resource ban** command, this adds a -INFINITY location constraint to the resource to prevent it from running on the indicated node. You can execute the **pcs**

**resource clear** or the **pcs constraint delete** command to remove the constraint. This does not necessarily move the resources back to the indicated node; where the resources can run at that point depends on how you have configured your resources initially. For information on resource constraints, refer to [Chapter 7, Resource Constraints](#).

If you specify the **--master** parameter of the **pcs resource ban** command, the scope of the constraint is limited to the master role and you must specify *master\_id* rather than *resource\_id*.

You can optionally configure a **lifetime** parameter for the **pcs resource ban** command to indicate a period of time the constraint should remain. For information on specifying units for the **lifetime** parameter and on specifying the intervals at which the **lifetime** parameter should be checked, see [Section 8.1, “Manually Moving Resources Around the Cluster”](#).

You can optionally configure a **--wait[=n]** parameter for the **pcs resource ban** command to indicate the number of seconds to wait for the resource to start on the destination node before returning 0 if the resource is started or 1 if the resource has not yet started. If you do not specify n, the default resource timeout will be used.

You can use the **debug-start** parameter of the **pcs resource** command to force a specified resource to start on the current node, ignoring the cluster recommendations and printing the output from starting the resource. This is mainly used for debugging resources; starting resources on a cluster is (almost) always done by Pacemaker and not directly with a **pcs** command. If your resource is not starting, it is usually due to either a misconfiguration of the resource (which you debug in the system log), constraints that the resource from starting, or the resource being disabled. You can use this command to test resource configuration, but it should not normally be used to start resources in a cluster.

The format of the **debug-start** command is as follows.

```
pcs resource debug-start resource_id
```

## 8.5. Disabling a Monitor Operations

The easiest way to stop a recurring monitor is to delete it. However, there can be times when you only want to disable it temporarily. In such cases, add **enabled="false"** to the operation's definition. When you want to reinstate the monitoring operation, set **enabled="true"** to the operation's definition.

## 8.6. Managed Resources

You can set a resource to **unmanaged** mode, which indicates that the resource is still in the configuration but Pacemaker does not manage the resource.

The following command sets the indicated resources to **unmanaged** mode.

```
pcs resource unmanage resource1 [resource2] ...
```

The following command sets resources to **managed** mode, which is the default state.

```
pcs resource manage resource1 [resource2] ...
```

You can specify the name of a resource group with the **pcs resource manage** or **pcs resource unmanage** command. The command will act on all of the resources in the group, so that you can set all of the resources in a group to **managed** or **unmanaged** mode with a single command and then manage the contained resources individually.

# Chapter 9. Advanced Resource types

This chapter describes advanced resource types that Pacemaker supports.

## 9.1. Resource Clones

You can clone a resource so that the resource can be active on multiple nodes. For example, you can use cloned resources to configure multiple instances of an IP resource to distribute throughout a cluster for node balancing. You can clone any resource provided the resource agent supports it. A clone consists of one resource or one resource group.



### Note

Only resources that can be active on multiple nodes at the same time are suitable for cloning. For example, a **Filesystem** resource mounting a non-clustered file system such as **ext4** from a shared memory device should not be cloned. Since the **ext4** partition is not cluster aware, this file system is not suitable for read/write operations occurring from multiple nodes at the same time.

### 9.1.1. Creating and Removing a Cloned Resource

You can create a resource and a clone of that resource at the same time with the following command.

```
pcs resource create resource_id standard:provider:type|type [resource options] \
clone [meta clone_options]
```

The name of the clone will be **resource\_id-clone**.

You cannot create a resource group and a clone of that resource group in a single command.

Alternately, you can create a clone of a previously-created resource or resource group with the following command.

```
pcs resource clone resource_id | group_name [clone_options]...
```

The name of the clone will be **resource\_id-clone** or **group\_name-clone**.



### Note

You need to configure resource configuration changes on one node only.



### Note

When configuring constraints, always use the name of the group or clone.

When you create a clone of a resource, the clone takes on the name of the resource with **-clone** appended to the name. The following commands creates a resource of type **apache** named **webfarm** and a clone of that resource named **webfarm-clone**.

```
# pcs resource create webfarm apache clone
```

Use the following command to remove a clone of a resource or a resource group. This does not remove the resource or resource group itself.

```
pcs resource unclone resource_id | group_name
```

For information on resource options, refer to [Section 6.1, “Resource Creation”](#).

[Table 9.1, “Resource Clone Options”](#) describes the options you can specify for a cloned resource.

**Table 9.1. Resource Clone Options**

Field	Description
<b>priority, target-role, is-managed</b>	Options inherited from resource that is being cloned, as described in <a href="#">Table 6.3, “Resource Meta Options”</a> .
<b>clone-max</b>	How many copies of the resource to start. Defaults to the number of nodes in the cluster.
<b>clone-node-max</b>	How many copies of the resource can be started on a single node; the default value is <b>1</b> .
<b>notify</b>	When stopping or starting a copy of the clone, tell all the other copies beforehand and when the action was successful. Allowed values: <b>false, true</b> . The default value is <b>false</b> .
<b>globally-unique</b>	Does each copy of the clone perform a different function? Allowed values: <b>false, true</b>  If the value of this option is <b>false</b> , these resources behave identically everywhere they are running and thus there can be only one copy of the clone active per machine.  If the value of this option is <b>true</b> , a copy of the clone running on one machine is not equivalent to another instance, whether that instance is running on another node or on the same node. The default value is <b>true</b> if the value of <b>clone-node-max</b> is greater than one; otherwise the default value is <b>false</b> .
<b>ordered</b>	Should the copies be started in series (instead of in parallel). Allowed values: <b>false, true</b> . The default value is <b>false</b> .
<b>interleave</b>	Changes the behavior of ordering constraints (between clones/masters) so that copies of the first clone can start or stop as soon as the copy on the same node of the second clone has started or stopped (rather than waiting until every instance of the second clone has started or stopped). Allowed values: <b>false, true</b> . The default value is <b>false</b> .

### 9.1.2. Clone Constraints

In most cases, a clone will have a single copy on each active cluster node. You can, however, set **clone-max** for the resource clone to a value that is less than the total number of nodes in the cluster. If this is the case, you can indicate which nodes the cluster should preferentially assign copies to with resource location constraints. These constraints are written no differently to those for regular

resources except that the clone's id must be used.

The following command creates a location constraint for the cluster to preferentially assign resource clone **webfarm-clone** to **node1**.

```
# pcs constraint location webfarm-clone prefers node1
```

Ordering constraints behave slightly differently for clones. In the example below, **webfarm-stats** will wait until all copies of **webfarm-clone** that need to be started have done so before being started itself. Only if no copies of **webfarm-clone** can be started then **webfarm-stats** will be prevented from being active. Additionally, **webfarm-clone** will wait for **webfarm-stats** to be stopped before stopping itself.

```
# pcs constraint order start webfarm-clone then webfarm-stats
```

Colocation of a regular (or group) resource with a clone means that the resource can run on any machine with an active copy of the clone. The cluster will choose a copy based on where the clone is running and the resource's own location preferences.

Colocation between clones is also possible. In such cases, the set of allowed locations for the clone is limited to nodes on which the clone is (or will be) active. Allocation is then performed as normally.

The following command creates a colocation constraint to ensure that the resource **webfarm-stats** runs on the same node as an active copy of **webfarm-clone**.

```
# pcs constraint colocation add webfarm-stats with webfarm-clone
```

### 9.1.3. Clone Stickiness

To achieve a stable allocation pattern, clones are slightly sticky by default. If no value for **resource-stickiness** is provided, the clone will use a value of 1. Being a small value, it causes minimal disturbance to the score calculations of other resources but is enough to prevent Pacemaker from needlessly moving copies around the cluster.

## 9.2. MultiState Resources: Resources That Have Multiple Modes

Multistate resources are a specialization of Clone resources. They allow the instances to be in one of two operating modes; these are called **Master** and **Slave**. The names of the modes do not have specific meanings, except for the limitation that when an instance is started, it must come up in the **Slave** state.

You can create a resource as a master/slave clone with the following single command.

```
pcs resource create resource_id standard:provider:type|type [resource options] \
--master [meta master_options]
```

The name of the master/slave clone will be **resource\_id-master**.

Alternately, you can create a master/slave resource from a previously-created resource or resource group with the following command: When you use this command, you can specify a name for the master/slave clone. If you do not specify a name, the name of the master/slave clone will be **resource\_id-master** or **group\_name-master**.

```
pcs resource master master/slave_name resource_id|group_name [master_options]
```

For information on resource options, refer to [Section 6.1, “Resource Creation”](#).

[Table 9.2, “Properties of a Multistate Resource”](#) describes the options you can specify for a multistate resource.

**Table 9.2. Properties of a Multistate Resource**

Field	Description
<b>id</b>	Your name for the multistate resource
<b>priority</b> , <b>target-role</b> , <b>is-managed</b>	See <a href="#">Table 6.3, “Resource Meta Options”</a> .
<b>clone-max</b> , <b>clone-node-max</b> , <b>notify</b> , <b>globally-unique</b> , <b>ordered</b> , <b>interleave</b>	See <a href="#">Table 9.1, “Resource Clone Options”</a> .
<b>master-max</b>	How many copies of the resource can be promoted to <b>master</b> status; default 1.
<b>master-node-max</b>	How many copies of the resource can be promoted to <b>master</b> status on a single node; default 1.

### 9.2.1. Monitoring Multi-State Resources

To add a monitoring operation for the master resource only, you can add an additional monitor operation to the resource. Note, however, that every monitor operation on a resource must have a different interval.

The following example configures a monitor operation with an interval of 11 seconds on the master resource for **ms\_resource**. This monitor operation is in addition to the default monitor operation with the default monitor interval of 10 seconds.

```
# pcs resource op add ms_resource interval=11s role=Master
```

### 9.2.2. Multistate Constraints

In most cases, a multistate resources will have a single copy on each active cluster node. If this is not the case, you can indicate which nodes the cluster should preferentially assign copies to with resource location constraints. These constraints are written no differently than those for regular resources.

For information on resource location constraints, see [Section 7.1, “Location Constraints”](#).

You can create a colocation constraint which specifies whether the resources are master or slave resources. The following command creates a resource colocation constraint.

```
pcs constraint colocation add [master|slave] source_resource with [master|slave] target_resource [score] [options]
```

For information on colocation constraints, see [Section 7.3, “Colocation of Resources”](#).

When configuring an ordering constraint that includes multistate resources, one of the actions that you can specify for the resources is **promote**, indicating that the resource be promoted from slave to master. Additionally, you can specify an action of **demote**, indicated that the resource be demoted from master to slave.

The command for configuring an order constraint is as follows.

```
pcs constraint order [action] resource_id then [action] resource_id [options]
```

For information on resource order constraints, see [Section 7.2, “Order Constraints”](#).

### 9.2.3. Multistate Stickiness

To achieve a stable allocation pattern, multistate resources are slightly sticky by default. If no value for **resource-stickiness** is provided, the multistate resource will use a value of 1. Being a small value, it causes minimal disturbance to the score calculations of other resources but is enough to prevent Pacemaker from needlessly moving copies around the cluster.

## 9.3. The pacemaker\_remote Service

The **pacemaker\_remote** service allows nodes not running **corosync** to integrate into the cluster and have the cluster manage their resources just as if they were real cluster nodes.

Among the capabilities that the **pacemaker\_remote** service provides are the following:

- » The **pacemaker\_remote** service allows you to scale beyond the **corosync** 16-node limit.
- » The **pacemaker\_remote** service allows you to manage a virtual environment as a cluster resource and also to manage individual services within the virtual environment as cluster resources.

The following terms are used to describe the **pacemaker\_remote** service.

- » *cluster node* — A node running the High Availability services (**pacemaker** and **corosync**).
- » *remote node* — A node running **pacemaker\_remote** to remotely integrate into the cluster without requiring **corosync** cluster membership. A remote node is configured as a cluster resource that uses the **ocf:pacemaker:remote** resource agent.
- » *guest node* — A virtual guest node running the **pacemaker\_remote** service. A guest node is configured using the **remote-node** metadata option of a resource agent such as **ocf:pacemaker:VirtualDomain**. The virtual guest resource is managed by the cluster; it is both started by the cluster and integrated into the cluster as a remote node.
- » *pacemaker\_remote* — A service daemon capable of performing remote application management within remote nodes and guest nodes (KVM and LXC) in a Pacemaker cluster environment. This service is an enhanced version of Pacemaker’s local resource management daemon (LRMD) that is capable of managing resources remotely on a node not running corosync.
- » *LXC* — A Linux Container defined by the **libvirt-lxc** Linux container driver.

A Pacemaker cluster running the **pacemaker\_remote** service has the following characteristics.

- » Remote nodes and guest nodes run the **pacemaker\_remote** service (with very little configuration required on the virtual machine side).
- » The cluster stack (**pacemaker** and **corosync**), running on the cluster nodes, connects to the **pacemaker\_remote** service on the remote nodes, allowing them to integrate into the cluster.
- » The cluster stack (**pacemaker** and **corosync**), running on the cluster nodes, launches the guest nodes and immediately connects to the **pacemaker\_remote** service on the guest nodes, allowing them to integrate into the cluster.

The key difference between the cluster nodes and the remote and guest nodes that the cluster nodes manage is that the remote and guest nodes are not running the cluster stack. This means the remote and guest nodes have the following limitations:

- » they do not take place in quorum
- » they do not execute fencing device actions
- » they are not eligible to be the cluster's Designated Controller (DC)
- » they do not themselves run the full range of **pcs** commands

On the other hand, remote nodes and guest nodes are not bound to the scalability limits associated with the cluster stack.

Other than these noted limitations, the remote nodes behave just like cluster nodes in respect to resource management, and the remote and guest nodes can themselves be fenced. The cluster is fully capable of managing and monitoring resources on each remote and guest node: You can build constraints against them, put them in standby, or perform any other action you perform on cluster nodes with the **pcs** commands. Remote and guest nodes appear in cluster status output just as cluster nodes do.

### 9.3.1. Host and Guest Authentication

The connection between cluster nodes and `pacemaker_remote` is secured using Transport Layer Security (TLS) with pre-shared key (PSK) encryption and authentication over TCP (using port 3121 by default). This means both the cluster node and the node running `pacemaker_remote` must share the same private key. By default this key must be placed at `/etc/pacemaker/authkey` on both cluster nodes and remote nodes.

### 9.3.2. Guest Node Resource Options

When configuring a virtual machine or LXC resource to act as a guest node, you create a **VirtualDomain** resource, which manages the virtual machine. For descriptions of the options you can set for a **VirtualDomain** resource, use the following command.

```
# pcs resource describe VirtualDomain
```

In addition to the **VirtualDomain** resource options, you can configure metadata options to both enable the resource as a guest node and define the connection parameters. [Table 9.3, “Metadata Options for Configuring KVM/LXC Resources as Remote Nodes”](#) describes these metadata options.

**Table 9.3. Metadata Options for Configuring KVM/LXC Resources as Remote Nodes**

Field	Default	Description
<code>remote-node</code>	<none>	The name of the guest node this resource defines. This both enables the resource as a guest node and defines the unique name used to identify the guest node. <b>WARNING:</b> This value cannot overlap with any resource or node IDs.
<code>remote-port</code>	3121	Configures a custom port to use for the guest connection to <code>pacemaker_remote</code> .
<code>remote-addr</code>	<code>remote-node</code> value used as host name	The IP address or host name to connect to if remote node's name is not the host name of the guest

Field	Default	Description
remote-connect-timeout	60s	Amount of time before a pending guest connection will time out

### 9.3.3. Remote Node Resource Options

You configure a remote node as a cluster resource with the `pcs resource create`, specifying `ocf:pacemaker:remote` as the resource type. [Table 9.4, “Resource Options for Remote Nodes”](#) describes the resource options you can configure for a `remote` resource.

**Table 9.4. Resource Options for Remote Nodes**

Field	Default	Description
<code>reconnect_interval</code>	0	Time in seconds to wait before attempting to reconnect to a remote node after an active connection to the remote node has been severed. This wait is recurring. If reconnect fails after the wait period, a new reconnect attempt will be made after observing the wait time. When this option is in use, Pacemaker will keep attempting to reach out and connect to the remote node indefinitely after each wait interval.
<code>server</code>		Server location to connect to. This can be an IP address or host name.
<code>port</code>		TCP port to connect to.

### 9.3.4. Changing Default pacemaker\_remote Options

If you need to change the default port or `authkey` location for either Pacemaker or `pacemaker_remote`, there are environment variables you can set that affect both of those daemons. These environment variables can be enabled by placing them in the `/etc/sysconfig/pacemaker` file as follows.

```
##### Pacemaker Remote
# Use a custom directory for finding the authkey.
PCMK_authkey_location=/etc/pacemaker/authkey
#
# Specify a custom port for Pacemaker Remote connections
PCMK_remote_port=3121
```

Note that when you change the default key location on a particular node (cluster node, guest node or remote node), it is sufficient to set `PCMK_authkey_location` on that node (and put the key in that location). It is not necessary that the location be the same on every node, although doing so makes administration easier.

When changing the default port used by a particular guest node or remote node, the `PCMK_remote_port` variable must be set in that node's `/etc/sysconfig/pacemaker` file, and the cluster resource creating the guest node or remote node connection must also be configured with the same port number (using the `remote-port` metadata option for guest nodes, or the `port` option for remote nodes).

### 9.3.5. Configuration Overview: KVM Guest Node

This section provides a high-level summary overview of the steps to perform to have Pacemaker launch a virtual machine and to integrate that machine as a guest node, using **libvirt** and KVM virtual guests.

1. After installing the virtualization software and enabling the **libvirtd** service on the cluster nodes, put the same encryption key with the path **/etc/pacemaker/authkey** on every cluster node and virtual machine. This secures remote communication and authentication.

Run the following set of commands on every node to create the **authkey** directory with secure permissions.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
```

The following command shows one method to create an encryption key. You should create the key only once and then copy it to all of the nodes.

```
# dd if=/dev/urandom of=/etc/pacemaker/authkey bs=4096 count=1
```

2. On every virtual machine, install **pacemaker\_remote** packages, start the **pacemaker\_remote** service and enable it to run on startup, and allow TCP port 3121 through the firewall.

```
# yum install pacemaker-remote resource-agents
# systemctl start pacemaker_remote.service
# systemctl enable pacemaker_remote.service
# firewall-cmd --add-port 3121/tcp --permanent
# firewall-cmd --reload
```

3. Give each virtual machine a static network address and unique host name, which should be known to all nodes. For information on setting a static IP address for the guest virtual machine, see the *Virtualization Deployment and Administration Guide*.
4. To create the **VirtualDomain** resource agent for the management of the virtual machine, Pacemaker requires the virtual machine's xml config file to be dumped to a file on disk. For example, if you created a virtual machine named **guest1**, dump the xml to a file somewhere on the host. You can use a file name of your choosing; this example uses **/etc/pacemaker/guest1.xml**.

```
# virsh dumpxml guest1 > /etc/pacemaker/guest1.xml
```

5. If it is running, shut down the guest node. Pacemaker will start the node when it is configured in the cluster.
6. Create the **VirtualDomain** resource, configuring the **remote-node** resource meta option to indicate that the virtual machine is a guest node capable of running resources.

In the example below, the meta-attribute **remote-node=guest1** tells pacemaker that this resource is a guest node with the host name **guest1** that is capable of being integrated into the cluster. The cluster will attempt to contact the virtual machine's **pacemaker\_remote** service at the host name **guest1** after it launches.

From a cluster node, enter the following command.

```
# pcs resource create vm-guest1 VirtualDomain hypervisor="qemu:///system"
config="/virtual_machines/vm-guest1.xml" meta remote-node=guest1
```

- After creating the **VirtualDomain** resource, you can treat the guest node just as you would treat any other node in the cluster. For example, you can create a resource and place a resource constraint on the resource to run on the guest node as in the following commands, which are run from a cluster node. As of Red Hat Enterprise Linux 7.3, you can include guest nodes in groups, which allows you to group a storage device, file system, and VM.

```
# pcs resource create webserver apache params
  configfile=/etc/httpd/conf/httpd.conf op monitor interval=30s
# pcs constraint location webserver prefers guest1
```

### 9.3.6. Configuration Overview: Remote Node

This section provides a high-level summary overview of the steps to perform to configure a Pacemaker remote node and to integrate that node into an existing Pacemaker cluster environment.

- On the node that you will be configuring as a remote node, allow cluster-related services through the local firewall.

```
# firewall-cmd --permanent --add-service=high-availability
success
# firewall-cmd --reload
success
```

#### Note

If you are using **iptables** directly, or some other firewall solution besides **firewalld**, simply open the following ports, which can be used by various clustering components: TCP ports 2224, 3121, and 21064, and UDP port 5405.

- Install the **pacemaker\_remote** daemon on the remote node.

```
# yum install -y pacemaker-remote resource-agents pcs
```

- All nodes (both cluster nodes and remote nodes) must have the same authentication key installed for the communication to work correctly. If you already have a key on an existing node, use that key and copy it to the remote node. Otherwise, create a new key on the remote node.

Run the following set of commands on the remote node to create a directory for the authentication key with secure permissions.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
```

The following command shows one method to create an encryption key on the remote node.

```
# dd if=/dev/urandom of=/etc/pacemaker/authkey bs=4096 count=1
```

- Start and enable the **pacemaker\_remote** daemon on the remote node.

```
# systemctl enable pacemaker_remote.service
# systemctl start pacemaker_remote.service
```

5. On the cluster node, create a location for the shared authentication key with the same path as the authentication key on the remote node and copy the key into that directory. In this example, the key is copied from the remote node where the key was created.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
# scp remote1:/etc/pacemaker/authkey /etc/pacemaker/authkey
```

6. Run the following command from a cluster node to create a **remote** resource. In this case the remote node is **remote1**.

```
# pcs resource create remote1 ocf:pacemaker:remote
```

7. After creating the **remote** resource, you can treat the remote node just as you would treat any other node in the cluster. For example, you can create a resource and place a resource constraint on the resource to run on the remote node as in the following commands, which are run from a cluster node.

```
# pcs resource create webserver apache params
configfile=/etc/httpd/conf/httpd.conf op monitor interval=30s
# pcs constraint location webserver prefers remote1
```



### Warning

Never involve a remote node connection resource in a resource group, colocation constraint, or order constraint.

8. Configure fencing resources for the remote node. Remote nodes are fenced the same way as cluster nodes. Configure fencing resources for use with remote nodes the same as you would with cluster nodes. Note, however, that remote nodes can never initiate a fencing action. Only cluster nodes are capable of actually executing a fencing operation against another node.

#### 9.3.7. System Upgrades and pacemaker\_remote

As of Red Hat Enterprise Linux 7.3, if the **pacemaker\_remote** service is stopped on an active Pacemaker Remote node, the cluster will gracefully migrate resources off the node before stopping the node. This allows you to perform software upgrades and other routine maintenance procedures without removing the node from the cluster. Once **pacemaker\_remote** is shut down, however, the cluster will immediately try to reconnect. If **pacemaker\_remote** is not restarted within the resource's monitor timeout, the cluster will consider the monitor operation as failed.

If you wish to avoid monitor failures when the **pacemaker\_remote** service is stopped on an active Pacemaker Remote node, you can use the following procedure to take the node out of the cluster before performing any system administration that might stop **pacemaker\_remote**



## Warning

For Red Hat Enterprise Linux release 7.2 and earlier, if **pacemaker\_remote** stops on a node that is currently integrated into a cluster, the cluster will fence that node. If the stop happens automatically as part of a **yum update** process, the system could be left in an unusable state (particularly if the kernel is also being upgraded at the same time as **pacemaker\_remote**). For Red Hat Enterprise Linux release 7.2 and earlier you must use the following procedure to take the node out of the cluster before performing any system administration that might stop **pacemaker\_remote**.

1. Stop the node's connection resource with the **pcs resource disable resource\_name**, which will move all services off the node. For guest nodes, this will also stop the VM, so the VM must be started outside the cluster (for example, using **virsh**) to perform any maintenance.
2. Perform the desired maintenance.
3. When ready to return the node to the cluster, re-enable the resource with the **pcs resource enable**.

### 9.3.8. Converting a VM Resource to a Guest Node

Use the following command to convert an existing **VirtualDomain** resource into a guest node. You do not need to run this command if the resource was originally created as a guest node.

```
pcs cluster remote-node add hostname resource_id [options]
```

Use the following command to disable a resource configured as a guest node on the specified host.

```
pcs cluster remote-node remove hostname
```

# Chapter 10. Cluster Quorum

A Red Hat Enterprise Linux High Availability Add-On cluster uses the **votequorum** service, in conjunction with fencing, to avoid split brain situations. A number of votes is assigned to each system in the cluster, and cluster operations are allowed to proceed only when a majority of votes is present. The service must be loaded into all nodes or none; if it is loaded into a subset of cluster nodes, the results will be unpredictable. For information on the configuration and operation of the **votequorum** service, see the **votequorum(5)** man page.

## 10.1. Configuring Quorum Options

There are some special features of quorum configuration that you can set when you create a cluster with the **pcs cluster setup** command. [Table 10.1, “Quorum Options”](#) summarizes these options.

**Table 10.1. Quorum Options**

Option	Description
<b>--auto_tie_breaker</b>	<p>When enabled, the cluster can suffer up to 50% of the nodes failing at the same time, in a deterministic fashion. The cluster partition, or the set of nodes that are still in contact with the <b>nodeid</b> configured in <b>auto_tie_breaker_node</b> (or lowest <b>nodeid</b> if not set), will remain quorate. The other nodes will be inquorate.</p> <p>The <b>auto_tie_breaker</b> option is principally used for clusters with an even number of nodes, as it allows the cluster to continue operation with an even split. For more complex failures, such as multiple, uneven splits, it is recommended that you use a quorum device, as described in <a href="#">Section 10.5, “Quorum Devices (Technical Preview)”</a>. The <b>auto_tie_breaker</b> option is incompatible with quorum devices.</p>
<b>--wait_for_all</b>	<p>When enabled, the cluster will be quorate for the first time only after all nodes have been visible at least once at the same time.</p> <p>The <b>wait_for_all</b> option is primarily used for two-node clusters and for even-node clusters using the quorum device <b>Ims</b> (last man standing) algorithm.</p> <p>The <b>wait_for_all</b> option is automatically enabled when a cluster has two nodes, does not use a quorum device, and <b>auto_tie_breaker</b> is disabled. You can override this by explicitly setting <b>wait_for_all</b> to 0.</p>
<b>--last_man_standing</b>	<p>When enabled, the cluster can dynamically recalculate <b>expected_votes</b> and quorum under specific circumstances. You must enable <b>wait_for_all</b> when you enable this option. The <b>last_man_standing</b> option is incompatible with quorum devices.</p>
<b>--last_man_standing_window</b> <b>w</b>	<p>The time, in milliseconds, to wait before recalculating <b>expected_votes</b> and quorum after a cluster loses nodes.</p>

For further information about configuring and using these options, see the **votequorum(5)** man page.

## 10.2. Quorum Administration Commands (Red Hat Enterprise Linux 7.3 and Later)

Once a cluster is running, you can enter the following cluster quorum commands.

The following command shows the quorum configuration.

```
pcs quorum [config]
```

The following command shows the quorum runtime status.

```
pcs quorum status
```

If you take nodes out of a cluster for a long period of time and the loss of those nodes would cause quorum loss, you can change the value of the **expected\_votes** parameter for the live cluster with the **pcs quorum expected-votes** command. This allows the cluster to continue operation when it does not have quorum.



### Warning

Changing the expected votes in a live cluster should be done with extreme caution. If less than 50% of the cluster is running because you have manually changed the expected votes, then the other nodes in the cluster could be started separately and run cluster services, causing data corruption and other unexpected results. If you change this value, you should ensure that the **wait\_for\_all** parameter is enabled.

The following command sets the expected votes in the live cluster to the specified value. This affects the live cluster only and does not change the configuration file; the value of **expected\_votes** is reset to the value in the configuration file in the event of a reload.

```
pcs quorum expected-votes votes
```

## 10.3. Modifying Quorum Options (Red Hat Enterprise Linux 7.3 and later)

As of Red Hat Enterprise Linux 7.3, you can modify general quorum options for your cluster with the **pcs quorum update** command. Executing this command requires that the cluster be stopped. For information on the quorum options, see the **votequorum(5)** man page.

The format of the **pcs quorum update** command is as follows.

```
pcs quorum update [auto_tie_breaker=[0|1]] [[last_man_standing=[0|1]]  
[last_man_standing_window=[time-in-ms]] [wait_for_all=[0|1]]]
```

The following series of commands modifies the **wait\_for\_all** quorum option and displays the updated status of the option. Note that the system does not allow you to execute this command while the cluster is running.

```
[root@node1:~]# pcs quorum update wait_for_all=1  
Checking corosync is not running on nodes...
```

```
Error: node1: corosync is running
Error: node2: corosync is running

[root@node1:~]# pcs cluster stop --all
node2: Stopping Cluster (pacemaker)...
node1: Stopping Cluster (pacemaker)...
node1: Stopping Cluster (corosync)...
node2: Stopping Cluster (corosync)...

[root@node1:~]# pcs quorum update wait_for_all=1
Checking corosync is not running on nodes...
node2: corosync is not running
node1: corosync is not running
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded

[root@node1:~]# pcs quorum config
Options:
  wait_for_all: 1
```

## 10.4. The quorum unblock Command

In a situation in which you know that the cluster is inquorate but you want the cluster to proceed with resource management, you can use the following command to prevent the cluster from waiting for all nodes when establishing quorum.



### Note

This command should be used with extreme caution. Before issuing this command, it is imperative that you ensure that nodes that are not currently in the cluster are switched off and have no access to shared resources.

```
# pcs cluster quorum unblock
```

## 10.5. Quorum Devices (Technical Preview)



### Important

The quorum device feature is provided for technical preview only. For details on what "technical preview" means, see [Technology Preview Features Support Scope](#).

As of Red Hat Enterprise Linux 7.3, you can configure a separate quorum device which acts as a third-party arbitration device for the cluster. Its primary use is to allow a cluster to sustain more node failures than standard quorum rules allow. A quorum device is recommended for clusters with an even number of nodes and highly recommended for two-node clusters.

You must take the following into account when configuring a quorum device.

- » It is recommended that a quorum device be run on a different physical network at the same site as the cluster that uses the quorum device. Ideally, the quorum device host should be in a separate rack than the main cluster, or at least on a separate PSU and not on the same network segment as the corosync ring or rings.
- » You cannot use more than one quorum device in a cluster at the same time.
- » Although you cannot use more than one quorum device in a cluster at the same time, a single quorum device may be used by several clusters at the same time. Each cluster using that quorum device can use different algorithms and quorum options, as those are stored on the cluster nodes themselves. For example, a single quorum device can be used by one cluster with an **ffsplit** (fifty/fifty split) algorithm and by a second cluster with an **lms** (last man standing) algorithm.
- » A quorum device should not be run on an existing cluster node.

### 10.5.1. Installing Quorum Device Packages

Configuring a quorum device for a cluster requires that you install the following packages:

- » Install **corosync-qdevice** on the cluster nodes.

```
[root@node1:~]# yum install corosync-qdevice
[root@node2:~]# yum install corosync-qdevice
```

- » Install **corosync-qnetd** on the quorum device host.

```
[root@qdevice:~]# yum install corosync-qnetd
```

### 10.5.2. Configuring a Quorum Device

This section provides a sample procedure to configure a quorum device in a Red Hat high availability cluster. The following procedure configures a quorum device and adds it to the cluster. In this example:

- » The node used for a quorum device is **qdevice**.
- » The quorum device model is **net**, which is currently the only supported model. The **net** model supports the following algorithms:
  - **ffsplit**: fifty-fifty split. This provides exactly one vote to the partition with the highest number of active nodes.
  - **lms**: last-man-standing. If the node is the only one left in the cluster that can see the **qnetd** server, then it returns a vote.



#### Warning

The LMS algorithm allows the cluster to remain quorate even with only one remaining node, but it also means that the voting power of the quorum device is great since it is the same as `number_of_nodes - 1`. Losing connection with the quorum device means losing `number_of_nodes - 1` votes, which means that only a cluster with all nodes active can remain quorate (by overvoting the quorum device); any other cluster becomes inquorate.

For more detailed information on the implementation of these algorithms, see the **corosync-qdevice(8)** man page.

- » The cluster nodes are **node1** and **node2**.

The following procedure configures a quorum device and adds that quorum device to a cluster.

1. On the node that you will use to host your quorum device, configure the quorum device with the following command. This command configures and starts the quorum device model **net** and configures the device to start on boot.

```
[root@qdevice:~]# pcs qdevice setup model net --enable --start
Quorum device 'net' initialized
quorum device enabled
Starting quorum device...
quorum device started
```

After configuring the quorum device, you can check its status. This should show that the **corosync-qnetd** daemon is running and, at this point, there are no clients connected to it. The **--full** command option provides detailed output.

```
[root@qdevice:~]# pcs qdevice status net --full
QNetd address:          *:5403
TLS:                  Supported (client certificate required)
Connected clients:      0
Connected clusters:     0
Maximum send/receive size: 32768/32768 bytes
```

2. Add the quorum device to the cluster.

Before adding the quorum device, you can check the current configuration and status for the quorum device for later comparison. The output for these commands indicates that the cluster is not yet using a quorum device.

```
[root@node1:~]# pcs quorum config
Options:
```

```
[root@node1:~]# pcs quorum status
Quorum information
-----
Date:      Wed Jun 29 13:15:36 2016
Quorum provider: corosync_votequorum
Nodes:      2
Node ID:    1
Ring ID:    1/8272
Quorate:   Yes

Votequorum information
-----
Expected votes: 2
Highest expected: 2
Total votes: 2
Quorum: 1
Flags: 2Node Quorate
```

## Membership information

```
-----  
Nodeid  Votes  Qdevice Name  
 1      1      NR node1 (local)  
 2      1      NR node2
```

The following command adds the quorum device that you have previously created to the cluster. You cannot use more than one quorum device in a cluster at the same time. However, one quorum device can be used by several clusters at the same time. This example command configures the quorum device to use the **ffsplit** algorithm. For information on the configuration options for the quorum device, see the **corosync-qdevice(8)** man page.

```
[root@node1:~]# pcs quorum device add model net host=qdevice  
algorithm=ffsplit  
Setting up qdevice certificates on nodes...  
node2: Succeeded  
node1: Succeeded  
Enabling corosync-qdevice...  
node1: corosync-qdevice enabled  
node2: corosync-qdevice enabled  
Sending updated corosync.conf to nodes...  
node1: Succeeded  
node2: Succeeded  
Corosync configuration reloaded  
Starting corosync-qdevice...  
node1: corosync-qdevice started  
node2: corosync-qdevice started
```

### 3. Check the configuration status of the quorum device.

From the cluster side, you can execute the following commands to see how the configuration has changed.

The **pcs quorum config** shows the quorum device that has been configured.

```
[root@node1:~]# pcs quorum config  
Options:  
Device:  
  Model: net  
    algorithm: ffsplit  
    host: qdevice
```

The **pcs quorum status** command shows the quorum runtime status, indicating that the quorum device is in use.

```
[root@node1:~]# pcs quorum status  
Quorum information  
-----  
Date:       Wed Jun 29 13:17:02 2016  
Quorum provider: corosync_votequorum  
Nodes:      2  
Node ID:    1  
Ring ID:    1/8272  
Quorate:   Yes
```

```

Votequorum information
-----
Expected votes: 3
Highest expected: 3
Total votes: 3
Quorum: 2
Flags: Quorate Qdevice

Membership information
-----
  Nodeid   Votes   Qdevice Name
    1        1   A,V,NMW node1 (local)
    2        1   A,V,NMW node2
    0        1       Qdevice

```

The **pcs quorum device status** shows the quorum device runtime status.

```

[root@node1:~]# pcs quorum device status
Qdevice information
-----
Model: Net
Node ID: 1
Configured node list:
  0 Node ID = 1
  1 Node ID = 2
Membership node list: 1, 2

Qdevice-net information
-----
Cluster name: mycluster
QNetd host: qdevice:5403
Algorithm: ffsplit
Tie-breaker: Node with lowest node ID
State: Connected

```

From the quorum device side, you can execute the following status command, which shows the status of the **corosync-qnetd** daemon.

```

[root@qdevice:~]# pcs qdevice status net --full
QNetd address: *:5403
TLS: Supported (client certificate required)
Connected clients: 2
Connected clusters: 1
Maximum send/receive size: 32768/32768 bytes
Cluster "mycluster":
  Algorithm: ffsplit
  Tie-breaker: Node with lowest node ID
  Node ID 2:
    Client address: ::ffff:192.168.122.122:50028
    HB interval: 8000ms
    Configured node list: 1, 2
    Ring ID: 1.2050
    Membership node list: 1, 2
    TLS active: Yes (client certificate verified)

```

Vote:	ACK (ACK)
Node ID 1:	
Client address:	::ffff:192.168.122.121:48786
HB interval:	8000ms
Configured node list:	1, 2
Ring ID:	1.2050
Membership node list:	1, 2
TLS active:	Yes (client certificate verified)
Vote:	ACK (ACK)

### 10.5.3. Managing the Quorum Device Service

PCS provides the ability to manage the quorum device service on the local host (**corosync-qnetd**), as shown in the following example commands. Note that these commands affect only the **corosync-qnetd** service.

```
[root@qdevice:~]# pcs qdevice start net
[root@qdevice:~]# pcs qdevice stop net
[root@qdevice:~]# pcs qdevice enable net
[root@qdevice:~]# pcs qdevice disable net
[root@qdevice:~]# pcs qdevice kill net
```

### 10.5.4. Managing the Quorum Device Settings in a Cluster

The following sections describe the PCS commands that you can use to manage the quorum device settings in a cluster, showing examples that are based on the quorum device configuration in [Section 10.5.2, “Configuring a Quorum Device”](#).

#### 10.5.4.1. Changing Quorum Device Settings

You can change the setting of a quorum device with the **pcs quorum device update** command.



#### Warning

To change the **host** option of quorum device model **net**, use the **pcs quorum device remove** and the **pcs quorum device add** commands to set up the configuration properly, unless the old and the new host are the same machine.

The following command changes the quorum device algorithm to **Ims**.

```
[root@node1:~]# pcs quorum device update model algorithm=Ims
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
Corosync configuration reloaded
Reloading qdevice configuration on nodes...
node1: corosync-qdevice stopped
node2: corosync-qdevice stopped
node1: corosync-qdevice started
node2: corosync-qdevice started
```

#### 10.5.4.2. Removing a Quorum Device

Use the following command to remove a quorum device configured on a cluster node.

```
[root@node1:~]# pcs quorum device remove
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
Corosync configuration reloaded
Disabling corosync-qdevice...
node1: corosync-qdevice disabled
node2: corosync-qdevice disabled
Stopping corosync-qdevice...
node1: corosync-qdevice stopped
node2: corosync-qdevice stopped
Removing qdevice certificates from nodes...
node1: Succeeded
node2: Succeeded
```

After you have removed a quorum device, you should see the following error message when displaying the quorum device status.

```
[root@node1:~]# pcs quorum device status
Error: Unable to get quorum status: corosync-qdevice-tool: Can't connect to QDevice socket (is
QDevice running?): No such file or directory
```

#### 10.5.4.3. Destroying a Quorum Device

To disable and stop a quorum device on the quorum device host and delete all of its configuration files, use the following command.

```
[root@qdevice:~]# pcs qdevice destroy net
Stopping quorum device...
quorum device stopped
quorum device disabled
Quorum device 'net' configuration files removed
```

## Chapter 11. Pacemaker Rules

Rules can be used to make your configuration more dynamic. One common example is to set one value for **resource-stickiness** during working hours, to prevent resources from being moved back to their most preferred location, and another on weekends when no one is around to notice an outage.

Another use of rules might be to assign machines to different processing groups (using a node attribute) based on time and to then use that attribute when creating location constraints.

Each rule can contain a number of expressions, date-expressions and even other rules. The results of the expressions are combined based on the rule's **boolean-op** field to determine if the rule ultimately evaluates to **true** or **false**. What happens next depends on the context in which the rule is being used.

**Table 11.1. Properties of a Rule**

Field	Description
<b>role</b>	Limits the rule to apply only when the resource is in that role. Allowed values: <b>Started</b> , <b>Slave</b> , and <b>Master</b> . NOTE: A rule with <b>role="Master"</b> cannot determine the initial location of a clone instance. It will only affect which of the active instances will be promoted.
<b>score</b>	The score to apply if the rule evaluates to <b>true</b> . Limited to use in rules that are part of location constraints.
<b>score-attribute</b>	The node attribute to look up and use as a score if the rule evaluates to <b>true</b> . Limited to use in rules that are part of location constraints.
<b>boolean-op</b>	How to combine the result of multiple expression objects. Allowed values: <b>and</b> and <b>or</b> . The default value is <b>and</b> .

### 11.1. Node Attribute Expressions

Node attribute expressions are used to control a resource based on the attributes defined by a node or nodes.

**Table 11.2. Properties of an Expression**

Field	Description
<b>value</b>	User supplied value for comparison
<b>attribute</b>	The node attribute to test
<b>type</b>	Determines how the value(s) should be tested. Allowed values: <b>string</b> , <b>integer</b> , <b>version</b>

Field	Description
<b>operation</b>	The comparison to perform. Allowed values: <ul style="list-style-type: none"><li>* <b>lt</b> - True if the node attribute's value is less than <b>value</b></li><li>* <b>gt</b> - True if the node attribute's value is greater than <b>value</b></li><li>* <b>lte</b> - True if the node attribute's value is less than or equal to <b>value</b></li><li>* <b>gte</b> - True if the node attribute's value is greater than or equal to <b>value</b></li><li>* <b>eq</b> - True if the node attribute's value is equal to <b>value</b></li><li>* <b>ne</b> - True if the node attribute's value is not equal to <b>value</b></li><li>* <b>defined</b> - True if the node has the named attribute</li><li>* <b>not_defined</b> - True if the node does not have the named attribute</li></ul>

## 11.2. Time/Date Based Expressions

Date expressions are used to control a resource or cluster option based on the current date/time. They can contain an optional date specification.

**Table 11.3. Properties of a Date Expression**

Field	Description
<b>start</b>	A date/time conforming to the ISO8601 specification.
<b>end</b>	A date/time conforming to the ISO8601 specification.
<b>operation</b>	Compares the current date/time with the start or the end date or both the start and end date, depending on the context. Allowed values: <ul style="list-style-type: none"><li>* <b>gt</b> - True if the current date/time is after <b>start</b></li><li>* <b>lt</b> - True if the current date/time is before <b>end</b></li><li>* <b>in-range</b> - True if the current date/time is after <b>start</b> and before <b>end</b></li><li>* <b>date-spec</b> - performs a cron-like comparison to the current date/time</li></ul>

## 11.3. Date Specifications

Date specifications are used to create cron-like expressions relating to time. Each field can contain a single number or a single range. Instead of defaulting to zero, any field not supplied is ignored.

For example, **monthdays="1"** matches the first day of every month and **hours="09-17"** matches the hours between 9 am and 5 pm (inclusive). However, you cannot specify **weekdays="1,2"** or **weekdays="1-2,5-6"** since they contain multiple ranges.

**Table 11.4. Properties of a Date Specification**

Field	Description
<b>id</b>	A unique name for the date

Field	Description
<code>hours</code>	Allowed values: 0-23
<code>monthdays</code>	Allowed values: 0-31 (depending on month and year)
<code>weekdays</code>	Allowed values: 1-7 (1=Monday, 7=Sunday)
<code>yeardays</code>	Allowed values: 1-366 (depending on the year)
<code>months</code>	Allowed values: 1-12
<code>weeks</code>	Allowed values: 1-53 (depending on <code>weekyear</code> )
<code>years</code>	Year according the Gregorian calendar
<code>weekyears</code>	May differ from Gregorian years; for example, <b>2005-001</b> <b>Ordinal</b> is also <b>2005-01-01</b> <b>Gregorian</b> is also <b>2004-W53-6 Weekly</b>
<code>moon</code>	Allowed values: 0-7 (0 is new, 4 is full moon).

## 11.4. Durations

Durations are used to calculate a value for `end` when one is not supplied to `in_range` operations. They contain the same fields as `date_spec` objects but without the limitations (ie. you can have a duration of 19 months). Like `date_specs`, any field not supplied is ignored.

## 11.5. Configuring Rules with pcs

To configure a rule, use the following command. If `score` is omitted, it defaults to INFINITY. If `id` is omitted, one is generated from the `constraint_id`. The `rule_type` should be `expression` or `date_expression`.

```
pcs constraint rule add constraint_id [rule_type] [score=score [id=rule_id]  
expression|date_expression|date_spec options]
```

To remove a rule, use the following. If the rule that you are removing is the last rule in its constraint, the constraint will be removed.

```
pcs constraint rule remove rule_id
```

## 11.6. Sample Time Based Expressions

The following command configures an expression that is true if now is any time in the year 2005.

```
# pcs constraint location Webserver rule score=INFINITY date-spec years=2005
```

The following command configures an expression that is true from 9 am to 5 pm, Monday through Friday. Note that the hours value of 16 matches up to 16:59:59, as the numeric value (hour) still matches.

```
# pcs constraint location Webserver rule score=INFINITY date-spec hours="9-16"  
weekdays="1-5"
```

The following command configures an expression that is true when there is a full moon on Friday the thirteenth.

```
# pcs constraint location Webserver rule date-spec weekdays=5 monthdays=13  
moon=4
```

## 11.7. Using Rules to Determine Resource Location

You can use a rule to determine a resource's location with the following command.

```
pcs constraint location resource_id rule [rule_id] [role=master|slave] [score=score expression]
```

The *expression* can be one of the following:

- » **defined|not\_defined *attribute***
- » ***attribute* lt|gt|lte|gte|eq|ne *value***
- » **date [*start=start*] [*end=end*] operation=gt|lt|in-range**
- » **date-spec *date\_spec\_options***

## Chapter 12. Pacemaker Cluster Properties

Cluster properties control how the cluster behaves when confronted with situations that may occur during cluster operation.

- [Table 12.1, “Cluster Properties”](#) describes the cluster properties options.
- [Section 12.2, “Setting and Removing Cluster Properties”](#) describes how to set cluster properties.
- [Section 12.3, “Querying Cluster Property Settings”](#) describes how to list the currently set cluster properties.

### 12.1. Summary of Cluster Properties and Options

[Table 12.1, “Cluster Properties”](#) summarizes the Pacemaker cluster properties, showing the default values of the properties and the possible values you can set for those properties.

#### Note

In addition to the properties described in this table, there are additional cluster properties that are exposed by the cluster software. For these properties, it is recommended that you not change their values from their defaults.

**Table 12.1. Cluster Properties**

Option	Default	Description
<b>batch-limit</b>	30	The number of jobs that the transition engine (TE) is allowed to execute in parallel. The "correct" value will depend on the speed and load of your network and cluster nodes.
<b>migration-limit</b>	-1 (unlimited)	The number of migration jobs that the TE is allowed to execute in parallel on a node.
<b>no-quorum-policy</b>	stop	What to do when the cluster does not have quorum. Allowed values:  * ignore - continue all resource management  * freeze - continue resource management, but do not recover resources from nodes not in the affected partition  * stop - stop all resources in the affected cluster partition  * suicide - fence all nodes in the affected cluster partition
<b>symmetric-cluster</b>	true	Indicates whether resources can run on any node by default.

Option	Default	Description
<b>stonith-enabled</b>	true	Indicates that failed nodes and nodes with resources that cannot be stopped should be fenced. Protecting your data requires that you set this <b>true</b> .  If <b>true</b> , or unset, the cluster will refuse to start resources unless one or more STONITH resources have been configured also.
<b>stonith-action</b>	reboot	Action to send to STONITH device. Allowed values: <b>reboot</b> , <b>off</b> . The value <b>poweroff</b> is also allowed, but is only used for legacy devices.
<b>cluster-delay</b>	60s	Round trip delay over the network (excluding action execution). The "correct" value will depend on the speed and load of your network and cluster nodes.
<b>stop-orphan-resources</b>	true	Indicates whether deleted resources should be stopped.
<b>stop-orphan-actions</b>	true	Indicates whether deleted actions should be canceled.
<b>start-failure-is-fatal</b>	true	Indicates whether a failure to start a resource on a particular node prevents further start attempts on that node. When set to <b>false</b> , the cluster will decide whether to try starting on the same node again based on the resource's current failure count and migration threshold. For information on setting the <b>migration-threshold</b> option for a resource, see <a href="#">Section 8.2, “Moving Resources Due to Failure”</a> .
<b>pe-error-series-max</b>	-1 (all)	The number of PE inputs resulting in ERRORS to save. Used when reporting problems.
<b>pe-warn-series-max</b>	-1 (all)	The number of PE inputs resulting in WARNINGS to save. Used when reporting problems.
<b>pe-input-series-max</b>	-1 (all)	The number of "normal" PE inputs to save. Used when reporting problems.
<b>cluster-infrastructure</b>		The messaging stack on which Pacemaker is currently running. Used for informational and diagnostic purposes; not user-configurable.
<b>dc-version</b>		Version of Pacemaker on the cluster's Designated Controller (DC). Used for diagnostic purposes; not user-configurable.
<b>last-lrm-refresh</b>		Last refresh of the Local Resource Manager, given in units of seconds since epoch. Used for diagnostic purposes; not user-configurable.
<b>cluster-recheck-interval</b>	15 minutes	Polling interval for time-based changes to options, resource parameters and constraints. Allowed values: Zero disables polling, positive values are an interval in seconds (unless other SI units are specified, such as 5min).
<b>default-action-timeout</b>	20s	Timeout value for a Pacemaker action. The setting for an operation in a resource itself always takes precedence over the default value set as a cluster option.

Option	Default	Description
<b>maintenance-mode</b>	false	Maintenance Mode tells the cluster to go to a "hands off" mode, and not start or stop any services until told otherwise. When maintenance mode is completed, the cluster does a sanity check of the current state of any services, and then stops or starts any that need it.
<b>shutdown-escalation</b>	20min	The time after which to give up trying to shut down gracefully and just exit. Advanced use only.
<b>stonith-timeout</b>	60s	How long to wait for a STONITH action to complete.
<b>stop-all-resources</b>	false	Should the cluster stop all resources.
<b>default-resource-stickiness</b>	5000	Indicates how much a resource prefers to stay where it is. It is recommended that you set this value as a resource/operation default rather than as a cluster option.
<b>is-managed-default</b>	true	Indicates whether the cluster is allowed to start and stop a resource. It is recommended that you set this value as a resource/operation default rather than as a cluster option.
<b>enable-acl</b>	false	(Red Hat Enterprise Linux 7.1 and later) Indicates whether the cluster can use access control lists, as set with the <b>pcs acl</b> command.

## 12.2. Setting and Removing Cluster Properties

To set the value of a cluster property, use the following **pcs** command.

```
pcs property set property=value
```

For example, to set the value of **symmetric-cluster** to **false**, use the following command.

```
# pcs property set symmetric-cluster=false
```

You can remove a cluster property from the configuration with the following command.

```
pcs property unset property
```

Alternately, you can remove a cluster property from a configuration by leaving the value field of the **pcs property set** command blank. This restores that property to its default value. For example, if you have previously set the **symmetric-cluster** property to **false**, the following command removes the value you have set from the configuration and restores the value of **symmetric-cluster** to **true**, which is its default value.

```
# pcs property set symmetric-cluster=
```

## 12.3. Querying Cluster Property Settings

In most cases, when you use the **pcs** command to display values of the various cluster components, you can use **pcs list** or **pcs show** interchangeably. In the following examples, **pcs list** is the format used to display an entire list of all settings for more than one property, while **pcs show** is the format used to display the values of a specific property.

To display the values of the property settings that have been set for the cluster, use the following **pcs** command.

```
pcs property list
```

To display all of the values of the property settings for the cluster, including the default values of the property settings that have not been explicitly set, use the following command.

```
pcs property list --all
```

To display the current value of a specific cluster property, use the following command.

```
pcs property show property
```

For example, to display the current value of the **cluster-infrastructure** property, execute the following command:

```
# pcs property show cluster-infrastructure
Cluster Properties:
  cluster-infrastructure: cman
```

For informational purposes, you can display a list of all of the default values for the properties, whether they have been set to a value other than the default or not, by using the following command.

```
pcs property [list|show] --defaults
```

# Chapter 13. Triggering Scripts for Cluster Events

A Pacemaker cluster is an event-driven system, where an event might be a resource or node failure, a configuration change, or a resource starting or stopping. You can configure Pacemaker cluster alerts to take some external action when a cluster event occurs. You can configure cluster alerts in one of two ways:

- As of Red Hat Enterprise Linux 7.3, you can configure Pacemaker alerts by means of alert agents, which are external programs that the cluster calls in the same manner as the cluster calls resource agents to handle resource configuration and operation. This is the preferred, simpler method of configuring cluster alerts. Pacemaker alert agents are described in [Section 13.1, “Pacemaker Alert Agents \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- The **ocf:pacemaker:ClusterMon** resource can monitor the cluster status and trigger alerts on each cluster event. This resource runs the **crm\_mon** command in the background at regular intervals. For information on the **ClusterMon** resource see [Section 13.2, “Event Notification with Monitoring Resources”](#).

## 13.1. Pacemaker Alert Agents (Red Hat Enterprise Linux 7.3 and later)

You can create Pacemaker alert agents to take some external action when a cluster event occurs. The cluster passes information about the event to the agent by means of environment variables. Agents can do anything desired with this information, such as send an email message or log to a file or update a monitoring system.

- Pacemaker provides several sample alert agents, which are installed in **/usr/share/pacemaker/alerts** by default. These sample scripts may be copied and used as is, or they may be used as templates to be edited to suit your purposes. Refer to the source code of the sample agents for the full set of attributes they support. See [Section 13.1.1, “Using the Sample Alert Agents”](#) for an example of a basic procedure for configuring an alert that uses a sample alert agent.
- General information on configuring and administering alert agents is provided in [Section 13.1.2, “Alert Creation”](#), [Section 13.1.3, “Displaying, Modifying, and Removing Alerts”](#), [Section 13.1.4, “Alert Recipients”](#), [Section 13.1.5, “Alert Meta Options”](#), and [Section 13.1.6, “Alert Configuration Command Examples”](#).
- You can write your own alert agents for a Pacemaker alert to call. For information on writing alert agents, see [Section 13.1.7, “Writing an Alert Agent”](#).

### 13.1.1. Using the Sample Alert Agents

When you use one of sample alert agents, you should review the script to ensure that it suits your needs. These sample agents are provided as a starting point for custom scripts for specific cluster environments.

To use one of the sample alert agents, you must install the agent on each node in the cluster. For example, the following command installs the **alert\_snmp.sh.sample** script as **alert\_snmp.sh**.

```
# install --mode=0755 /usr/share/pacemaker/alerts/alert_snmp.sh.sample
/var/lib/pacemaker/alert_snmp.sh
```

After you have installed the script, you can create an alert that uses the script. The following example configures an alert that uses the installed **alert\_snmp.sh** alert agent to send cluster events as SNMP traps. By default, the script will send all events except successful monitor calls to the SNMP

server. This example configures the timestamp format as a meta option. For information about meta options, see [Section 13.1.5, “Alert Meta Options”](#). After configuring the alert, this example configures a recipient for the alert and displays the alert configuration.

```
# pcs alert create id=snmp_alert path=/var/lib/pacemaker/alert_snmp.sh meta
timestamp_format="%Y-%m-%d,%H:%M:%S.%01N".
# pcs alert recipient add snmp_alert 192.168.1.2
# pcs alert
Alerts:
Alert: snmp_alert (path=/var/lib/pacemaker/alert_snmp.sh)
Meta options: timestamp_format=%Y-%m-%d,%H:%M:%S.%01N.
Recipients:
Recipient: snmp_alert-recipient (value=192.168.1.2)
```

The following example installs the **alert\_smtp.sh** agent and then configures an alert that uses the installed alert agent to send cluster events as email messages. After configuring the alert, this example configures a recipient and displays the alert configuration.

```
# install --mode=0755 /usr/share/pacemaker/alerts/alert_smtp.sh.sample
/var/lib/pacemaker/alert_smtp.sh
# pcs alert create id=smtp_alert path=/var/lib/pacemaker/alert_smtp.sh options
email_sender=donotreply@example.com
# pcs alert recipient add smtp_alert admin@example.com
# pcs alert
Alerts:
Alert: smtp_alert (path=/var/lib/pacemaker/alert_smtp.sh)
Options: email_sender=donotreply@example.com
Recipients:
Recipient: smtp_alert-recipient (value=admin@example.com)
```

For more information on the format of the **pcs alert create** and **pcs alert recipient add** commands, see [Section 13.1.2, “Alert Creation”](#) and [Section 13.1.4, “Alert Recipients”](#).

### 13.1.2. Alert Creation

The following command creates a cluster alert. The options that you configure are agent-specific configuration values that are passed to the alert agent script at the path you specify as additional environment variables. If you do not specify a value for **id**, one will be generated. For information on alert meta options, [Section 13.1.5, “Alert Meta Options”](#).

```
pcs alert create path=path [id=alert-id] [description=description] [options [option=value]...] [meta
[meta-option=value]...]
```

Multiple alert agents may be configured; the cluster will call all of them for each event. Alert agents will be called only on cluster nodes. They will be called for events involving Pacemaker Remote nodes, but they will never be called on those nodes.

The following example creates a simple alert that will call **my-script.sh** for each event.

```
# pcs alert create id=my_alert path=/path/to/myscript.sh
```

For an example that shows how to create a cluster alert that uses one of the sample alert agents, see [Section 13.1.1, “Using the Sample Alert Agents”](#).

### 13.1.3. Displaying, Modifying, and Removing Alerts

The following command shows all configured alerts along with the values of the configured options.

```
pcs alert [config|show]
```

The following command updates an existing alert with the specified *alert-id* value.

```
pcs alert update alert-id [path=path] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

The following command removes an alert with the specified *alert-id* value.

```
pcs alert remove alert-id
```

### 13.1.4. Alert Recipients

Usually alerts are directed towards a recipient. Thus each alert may be additionally configured with one or more recipients. The cluster will call the agent separately for each recipient.

The recipient may be anything the alert agent can recognize: an IP address, an email address, a file name, or whatever the particular agent supports.

The following command adds a new recipient to the specified alert.

```
pcs alert recipient add alert-id recipient-value [id=recipient-id] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

The following command updates an existing alert recipient.

```
pcs alert recipient update recipient-id [value=recipient-value] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

The following command removes the specified alert recipient.

```
pcs alert recipient remove recipient-id
```

The following example command adds the alert recipient **my-alert-recipient** with a recipient ID of **my-recipient-id** to the alert **my-alert**. This will configure the cluster to call the alert script that has been configured for **my-alert** for each event, passing the recipient **some-address** as an environment variable.

```
# pcs alert recipient add my-alert my-alert-recipient id=my-recipient-id options value=some-address
```

### 13.1.5. Alert Meta Options

As with resource agents, meta options can be configured for alert agents to affect how Pacemaker calls them. [Table 13.1, “Alert Meta Options”](#) describes the alert meta options. Meta options can be configured per alert agent as well as per recipient.

**Table 13.1. Alert Meta Options**

Meta-Attribute	Default	Description
<b>timestamp-format</b>	%H:%M:%S.%06N	Format the cluster will use when sending the event's timestamp to the agent. This is a string as used with the <b>date(1)</b> command.
<b>timeout</b>	30s	If the alert agent does not complete within this amount of time, it will be terminated.

The following example configures an alert that calls the script **my-script.sh** and then adds two recipients to the alert. The first recipient has an ID of **my-alert-recipient1** and the second recipient has an ID of **my-alert-recipient2**. The script will get called twice for each event, with each call using a 15-second timeout. One call will be passed to the recipient **someuser@example.com** with a timestamp in the format **%D %H:%M**, while the other call will be passed to the recipient **otheruser@example.com** with a timestamp in the format **%c**.

```
# pcs alert create id=my-alert path=/path/to/my-script.sh meta timeout=15s
# pcs alert recipient add my-alert someuser@example.com id=my-alert-recipient1
meta timestamp-format=%D %H:%M
# pcs alert recipient add my-alert otheruser@example.com id=my-alert-recipient2
meta timestamp-format=%c
```

### 13.1.6. Alert Configuration Command Examples

The following sequential examples show some basic alert configuration commands to show the format to use to create alerts, add recipients, and display the configured alerts.

The following commands create a simple alert, add two recipients to the alert, and display the configured values.

- » Since no alert ID value is specified, the system creates an alert ID value of **alert**.
- » The first recipient creation command specifies a recipient of **rec\_value**. Since this command does not specify a recipient ID, the value of **alert-recipient** is used as the recipient ID.
- » The second recipient creation command specifies a recipient of **rec\_value2**. This command specifies a recipient ID of **my-recipient** for the recipient.

```
# pcs alert create path=/my/path
# pcs alert recipient add alert rec_value
# pcs alert recipient add alert rec_value2 id=my-recipient
# pcs alert config
Alerts:
Alert: alert (path=/my/path)
Recipients:
Recipient: alert-recipient (value=rec_value)
Recipient: my-recipient (value=rec_value2)
```

This following commands add a second alert and a recipient for that alert. The alert ID for the second alert is **my-alert** and the recipient value is **my-other-recipient**. Since no recipient ID is specified, the system provides a recipient id of **my-alert-recipient**.

```
# pcs alert create id=my-alert path=/path/to/script description=alert_description
options option1=value1 opt=val meta meta-option1=2 m=val
# pcs alert recipient add my-alert my-other-recipient
# pcs alert
Alerts:
```

```

Alert: alert (path=/my/path)
Recipients:
  Recipient: alert-recipient (value=rec_value)
  Recipient: my-recipient (value=rec_value2)
Alert: my-alert (path=/path/to/script)
Description: alert_description
Options: opt=val option1=value1
Meta options: m=val meta-option1=2
Recipients:
  Recipient: my-alert-recipient (value=my-other-recipient)

```

The following commands modify the alert values for the alert **my-alert** and for the recipient **my-alert-recipient**.

```

# pcs alert update my-alert options option1=newvalue1 meta m=newval
# pcs alert recipient update my-alert-recipient options option1=new meta
metaopt1=newopt
# pcs alert
Alerts:
Alert: alert (path=/my/path)
Recipients:
  Recipient: alert-recipient (value=rec_value)
  Recipient: my-recipient (value=rec_value2)
Alert: my-alert (path=/path/to/script)
Description: alert_description
Options: opt=val option1=newvalue1
Meta options: m=newval meta-option1=2
Recipients:
  Recipient: my-alert-recipient (value=my-other-recipient)
  Options: option1=new
  Meta options: metaopt1=newopt

```

The following command removes the recipient **my-alert-recipient** from alert.

```

# pcs alert recipient remove my-recipient
# pcs alert
Alerts:
Alert: alert (path=/my/path)
Recipients:
  Recipient: alert-recipient (value=rec_value)
Alert: my-alert (path=/path/to/script)
Description: alert_description
Options: opt=val option1=newvalue1
Meta options: m=newval meta-option1=2
Recipients:
  Recipient: my-alert-recipient (value=my-other-recipient)
  Options: option1=new
  Meta options: metaopt1=newopt

```

The following command removes **myalert** from the configuration.

```

# pcs alert remove my-alert
# pcs alert
Alerts:
Alert: alert (path=/my/path)

```

Recipients:  
Recipient: alert-recipient (value=rec\_value)

### 13.1.7. Writing an Alert Agent

There are three types of Pacemaker alerts: node alerts, fencing alerts, and resource alerts. The environment variables that are passed to the alert agents can differ, depending on the type of alert.

[Table 13.2, “Environment Variables Passed to Alert Agents”](#) describes the environment variables that are passed to alert agents and specifies when the environment variable is associated with a specific alert type.

**Table 13.2. Environment Variables Passed to Alert Agents**

Environment Variable	Description
<b>CRM_alert_kind</b>	The type of alert (node, fencing, or resource)
<b>CRM_alert_version</b>	The version of Pacemaker sending the alert
<b>CRM_alert_recipient</b>	The configured recipient
<b>CRM_alert_node_sequence</b>	A sequence number increased whenever an alert is being issued on the local node, which can be used to reference the order in which alerts have been issued by Pacemaker. An alert for an event that happened later in time reliably has a higher sequence number than alerts for earlier events. Be aware that this number has no cluster-wide meaning.
<b>CRM_alert_timestamp</b>	A timestamp created prior to executing the agent, in the format specified by the <b>timestamp-format</b> meta option. This allows the agent to have a reliable, high-precision time of when the event occurred, regardless of when the agent itself was invoked (which could potentially be delayed due to system load or other circumstances).
<b>CRM_alert_node</b>	Name of affected node
<b>CRM_alert_desc</b>	Detail about event. For node alerts, this is the node’s current state (member or lost). For fencing alerts, this is a summary of the requested fencing operation, including origin, target, and fencing operation error code, if any. For resource alerts, this is a readable string equivalent of <b>CRM_alert_status</b> .
<b>CRM_alert_nodeid</b>	ID of node whose status changed (provided with node alerts only)
<b>CRM_alert_task</b>	The requested fencing or resource operation (provided with fencing and resource alerts only)
<b>CRM_alert_rc</b>	The numerical return code of the fencing or resource operation (provided with fencing and resource alerts only)
<b>CRM_alert_rsc</b>	The name of the affected resource (resource alerts only)
<b>CRM_alert_interval</b>	The interval of the resource operation (resource alerts only)
<b>CRM_alert_target_rc</b>	The expected numerical return code of the operation (resource alerts only)
<b>CRM_alert_status</b>	A numerical code used by Pacemaker to represent the operation result (resource alerts only)

When writing an alert agent, you must take the following concerns into account.

- » Alert agents may be called with no recipient (if none is configured), so the agent must be able to handle this situation, even if it only exists in that case. Users may modify the configuration in stages, and add a recipient later.

- » If more than one recipient is configured for an alert, the alert agent will be called once per recipient. If an agent is not able to run concurrently, it should be configured with only a single recipient. The agent is free, however, to interpret the recipient as a list.
- » When a cluster event occurs, all alerts are fired off at the same time as separate processes. Depending on how many alerts and recipients are configured and on what is done within the alert agents, a significant load burst may occur. The agent could be written to take this into consideration, for example by queueing resource-intensive actions into some other instance, instead of directly executing them.
- » Alert agents are run as the **hacluster** user, which has a minimal set of permissions. If an agent requires additional privileges, it is recommended to configure **sudo** to allow the agent to run the necessary commands as another user with the appropriate privileges.
- » Take care to validate and sanitize user-configured parameters, such as **CRM\_alert\_timestamp** (whose content is specified by the user-configured **timestamp-format**), **CRM\_alert\_recipient**, and all alert options. This is necessary to protect against configuration errors. In addition, if some user can modify the CIB without having **hacluster**-level access to the cluster nodes, this is a potential security concern as well, and you should avoid the possibility of code injection.
- » If a cluster contains resources for which the **onfail** parameter is set to **fence**, there will be multiple fence notifications on failure, one for each resource for which this parameter is set plus one additional notification. Both the STONITH daemon and the **crmd** daemon will send notifications. Pacemaker performs only one actual fence operation in this case, however, no matter how many notifications are sent.



### Note

The alerts interface is designed to be backward compatible with the external scripts interface used by the **ocf:pacemaker:ClusterMon** resource. To preserve this compatibility, the environment variables passed to alert agents are available prepended with **CRM\_notify\_** as well as **CRM\_alert\_**. One break in compatibility is that the **ClusterMon** resource ran external scripts as the root user, while alert agents are run as the **hacluster** user. For information on configuring scripts that are triggered by the **ClusterMon**, see [Section 13.2, “Event Notification with Monitoring Resources”](#).

## 13.2. Event Notification with Monitoring Resources

The **ocf:pacemaker:ClusterMon** resource can monitor the cluster status and trigger alerts on each cluster event. This resource runs the **crm\_mon** command in the background at regular intervals.

By default, the **crm\_mon** command listens for resource events only; to enable listing for fencing events you can provide the **--watch-fencing** option to the command when you configure the **ClusterMon** resource. The **crm\_mon** command does not monitor for membership issues but will print a message when fencing is started and when monitoring is started for that node, which would imply that a member just joined the cluster.

The **ClusterMon** resource can execute an external program to determine what to do with cluster notifications by means of the **extra\_options** parameter. [Table 13.3, “Environment Variables Passed to the External Monitor Program”](#) lists the environment variables that are passed to that program, which describe the type of cluster event that occurred.

**Table 13.3. Environment Variables Passed to the External Monitor Program**

Environment Variable	Description
<b>CRM_notify_recipient</b>	The static external-recipient from the resource definition
<b>CRM_notify_node</b>	The node on which the status change happened
<b>CRM_notify_rsc</b>	The name of the resource that changed the status
<b>CRM_notify_task</b>	The operation that caused the status change
<b>CRM_notify_desc</b>	The textual output relevant error code of the operation (if any) that caused the status change
<b>CRM_notify_rc</b>	The return code of the operation
<b>CRM_target_rc</b>	The expected return code of the operation
<b>CRM_notify_status</b>	The numerical representation of the status of the operation

The following example configures a **ClusterMon** resource that executes the external program **crm\_logger.sh** which will log the event notifications specified in the program.

The following procedure creates the **crm\_logger.sh** program that this resource will use.

1. On one node of the cluster, create the program that will log the event notifications.

```
# cat <<-END >/usr/local/bin/crm_logger.sh
#!/bin/sh
logger -t "ClusterMon-External" "${CRM_notify_node} ${CRM_notify_rsc} \
${CRM_notify_task} ${CRM_notify_desc} ${CRM_notify_rc} \
${CRM_notify_target_rc} ${CRM_notify_status} ${CRM_notify_recipient}";
exit;
END
```

2. Set the ownership and permissions for the program.

```
# chmod 700 /usr/local/bin/crm_logger.sh
# chown root.root /usr/local/bin/crm_logger.sh
```

3. Use the **scp** command to copy the **crm\_logger.sh** program to the other nodes of the cluster, putting the program in the same location on those nodes and setting the same ownership and permissions for the program.

The following example configures the **ClusterMon** resource, named **ClusterMon-External**, that runs the program **/usr/local/bin/crm\_logger.sh**. The **ClusterMon** resource outputs the cluster status to an **html** file, which is **/var/www/html/cluster\_mon.html** in this example. The **pidfile** detects whether **ClusterMon** is already running; in this example that file is **/var/run/crm\_mon-external.pid**. This resource is created as a clone so that it will run on every node in the cluster. The **watch-fencing** is specified to enable monitoring of fencing events in addition to resource events, including the start/stop/monitor, start/monitor, and stop of the fencing resource.

```
# pcs resource create ClusterMon-External ClusterMon user=root \
update=10 extra_options="-E /usr/local/bin/crm_logger.sh --watch-fencing" \
htmlfile=/var/www/html/cluster_mon.html \
pidfile=/var/run/crm_mon-external.pid clone
```



## Note

The **crm\_mon** command that this resource executes and which could be run manually is as follows:

```
# /usr/sbin/crm_mon -p /var/run/crm_mon-manual.pid -d -i 5 \
-h /var/www/html/crm_mon-manual.html -E "/usr/local/bin/crm_logger.sh" \
--watch-fencing
```

The following example shows the format of the output of the monitoring notifications that this example yields.

```
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
st_notify_fence Operation st_notify_fence requested by rh6node1pcmk.examplerh.com for peer
rh6node2pcmk.examplerh.com: OK (ref=b206b618-e532-42a5-92eb-44d363ac848e) 0 0 0 #177
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
start OK 0 0 0
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com fence_xvms
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com
ClusterMon-External start OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com fence_xvms
start OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
start OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com
ClusterMon-External monitor OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 8:
monitor ClusterMon-External:1_monitor_0 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 16: start
ClusterMon-External:1_start_0 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
stop OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 15:
monitor ClusterMon-External_monitor_10000 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com
ClusterMon-External start OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
start OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
```

# Chapter 14. Configuring Multi-Site Clusters with Pacemaker (Technical Preview)



## Important

The Booth ticket manager is provided for technical preview only. For details on what "technical preview" means, see [Technology Preview Features Support Scope](#).

When a cluster spans more than one site, issues with network connectivity between the sites can lead to split-brain situations. When connectivity drops, there is no way for a node on one site to determine whether a node on another site has failed or is still functioning with a failed site interlink. In addition, it can be problematic to provide high availability services across two sites which are too far apart to keep synchronous.

To address these issues, Red Hat Enterprise Linux release 7.3 and later provides the ability to configure high availability clusters that span multiple sites through the use of a Booth cluster ticket manager. The Booth *ticket manager* is a distributed service that is meant to be run on a different physical network than the networks that connect the cluster nodes at particular sites. It yields another, loose cluster, a *Booth formation*, that sits on top of the regular clusters at the sites. This aggregated communication layer facilitates consensus-based decision processes for individual Booth tickets.

A Booth *ticket* is a singleton in the Booth formation and represents a time-sensitive, movable unit of authorization. Resources can be configured to require a certain ticket to run. This can ensure that resources are run at only one site at a time, for which a ticket or tickets have been granted.

You can think of a Booth formation as an overlay cluster consisting of clusters running at different sites, where all the original clusters are independent of each other. It is the Booth service which communicates to the clusters whether they have been granted a ticket, and it is Pacemaker that determines whether to run resources in a cluster based on a Pacemaker ticket constraint. This means that when using the ticket manager, each of the clusters can run its own resources as well as shared resources. For example there can be resources A, B and C running only in one cluster, resources D, E, and F running only in the other cluster, and resources G and H running in either of the two clusters as determined by a ticket. It is also possible to have an additional resource J that could run in either of the two clusters as determined by a separate ticket.

The following procedure provides an outline of the steps you follow to configure a multi-site configuration that uses the Booth ticket manager.

These example commands use the following arrangement:

- » Cluster 1 consists of the nodes **cluster1-node1** and **cluster1-node2**
- » Cluster 1 has a floating IP address assigned to it of 192.168.11.100
- » Cluster 2 consists of **cluster2-node1** and **cluster2-node2**
- » Cluster 2 has a floating IP address assigned to it of 192.168.22.100
- » The arbitrator node is **arbitrator-node** with an ip address of 192.168.99.100
- » The name of the Booth ticket that this configuration uses is **apacheticket**

These example commands assume that the cluster resources for an Apache service have been configured as part of the resource group **apache\_group** for each cluster. It is not required that the resources and resource groups be the same on each cluster to configure a ticket constraint for those

resources, since the Pacemaker instance for each cluster is independent, but that is a common failover scenario.

For a full cluster configuration procedure that configures an Apache service in a cluster, see the example in *High Availability Add-On Administration*.

Note that at any time in the configuration procedure you can enter the **pcs booth config** command to display the booth configuration for the current node or cluster or the **pcs booth status** command to display the current status of booth on the local node.

1. Create a Booth configuration on one node of one cluster. The addresses you specify for each cluster and for the arbitrator must be IP addresses. For each cluster, you specify a floating IP address.

```
[cluster1-node1 ~] # pcs booth setup sites 192.168.11.100 192.168.22.100  
arbitrators 192.168.99.100
```

This command creates the configuration files **/etc/booth/booth.conf** and **/etc/booth/booth.key** on the node from which it is run.

2. Create a ticket for the Booth configuration. This is the ticket that you will use to define the resource constraint that will allow resources to run only when this ticket has been granted to the cluster.

This basic failover configuration procedure uses only one ticket, but you can create additional tickets for more complicated scenarios where each ticket is associated with a different resource or resources.

```
[cluster1-node1 ~] # pcs booth ticket add apacheticket
```

3. Synchronize the Booth configuration to all nodes in the current cluster.

```
[cluster1-node1 ~] # pcs booth sync
```

4. From the arbitrator node, pull the Booth configuration to the arbitrator. If you have not previously done so, you must first authenticate **pcs** to the node from which you are pulling the configuration.

```
[arbitrator-node ~] # pcs cluster auth cluster1-node1  
[arbitrator-node ~] # pcs booth pull cluster1-node1
```

5. Pull the Booth configuration to the other cluster and synchronize to all the nodes of that cluster. As with the arbitrator node, if you have not previously done so, you must first authenticate **pcs** to the node from which you are pulling the configuration.

```
[cluster2-node1 ~] # pcs cluster auth cluster1-node1  
[cluster2-node1 ~] # pcs booth pull cluster1-node1  
[cluster2-node1 ~] # pcs booth sync
```

6. Start and enable Booth on the arbitrator.



## Note

You must not manually start or enable Booth on any of the nodes of the clusters since Booth runs as a Pacemaker resource in those clusters.

```
[arbitrator-node ~] # pcs booth start  
[arbitrator-node ~] # pcs booth enable
```

7. Configure Booth to run as a cluster resource on both cluster sites. This creates a resource group with **booth-ip** and **booth-service** as members of that group.

```
[cluster1-node1 ~] # pcs booth create ip 192.168.11.100  
[cluster2-node1 ~] # pcs booth create ip 192.168.22.100
```

8. Add a ticket constraint to the resource group you have defined for each cluster.

```
[cluster1-node1 ~] # pcs constraint ticket add apacheticket apachegroup  
[cluster2-node1 ~] # pcs constraint ticket add apacheticket apachegroup
```

You can run the following command to display the currently configured ticket constraints.

```
pcs constraint ticket [show]
```

9. Grant the ticket you created for this setup to the first cluster.

Note that it is not necessary to have defined ticket constraints before granting a ticket. Once you have initially granted a ticket to a cluster, then Booth takes over ticket management unless you override this manually with the **pcs booth ticket revoke** command. For information on the **pcs booth** administration commands, see the PCS help screen for the **pcs booth** command.

```
[cluster1-node1 ~] # pcs booth ticket grant apacheticket
```

It is possible to add or remove tickets at any time, even after completing this procedure. After adding or removing a ticket, however, you must synchronize the configuration files to the other nodes and clusters as well as to the arbitrator and grant the ticket as is shown in this procedure.

For information on additional Booth administration commands that you can use for cleaning up and removing Booth configuration files, tickets, and resources, see the PCS help screen for the **pcs booth** command.

## Appendix A. Cluster Creation in Red Hat Enterprise Linux 6 and Red Hat Enterprise Linux 7

Configuring a Red Hat High Availability Cluster in Red Hat Enterprise Linux 7 with Pacemaker requires a different set of configuration tools with a different administrative interface than configuring a cluster in Red Hat Enterprise Linux 6 with **rgmanager**. [Section A.1, “Cluster Creation with rgmanager and with Pacemaker”](#) summarizes the configuration differences between the various cluster components.

The Red Hat Enterprise Linux 6.5 release supports cluster configuration with Pacemaker, using the **pcs** configuration tool. [Section A.2, “Cluster Creation with Pacemaker in Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7”](#) summarizes some small configuration differences between **pcs** support in Red Hat Enterprise Linux 6.5 and **pcs** support in Red Hat Enterprise Linux 7.0.

### A.1. Cluster Creation with rgmanager and with Pacemaker

[Table A.1, “Comparison of Cluster Configuration with rgmanager and with Pacemaker”](#) provides a comparative summary of how you configure the components of a cluster with **rgmanager** in Red Hat Enterprise Linux 6 and with Pacemaker in Red Hat Enterprise Linux 7.

**Table A.1. Comparison of Cluster Configuration with rgmanager and with Pacemaker**

Configuration Component	rgmanager	Pacemaker
Cluster configuration file	The cluster configuration file on each node is <b>cluster.conf</b> file, which can be edited directly if desired. Otherwise, use the <b>luci</b> or <b>ccs</b> interface to define the cluster configuration.	The cluster and Pacemaker configuration files are <b>corosync.conf</b> and <b>cib.xml</b> . Do not edit these files directly; use the <b>pcs</b> or <b>pcsd</b> interface instead.
Network setup	Configure IP addresses and SSH before configuring the cluster.	Configure IP addresses and SSH before configuring the cluster.
Cluster Configuration Tools	<b>luci</b> , <b>ccs</b> command, manual editing of <b>cluster.conf</b> file.	<b>pcs</b> or <b>pcsd</b> .
Installation	Install <b>rgmanager</b> (which pulls in all dependencies, including <b>ricci</b> , <b>luci</b> , and the resource and fencing agents). If needed, install <b>Ivm2-cluster</b> and <b>gfs2-utils</b> .	Install <b>pcs</b> , and the fencing agents you require. If needed, install <b>Ivm2-cluster</b> and <b>gfs2-utils</b> .

Configuration Component	<b>rgmanager</b>	<b>Pacemaker</b>
Starting cluster services	<p>Start and enable cluster services with the following procedure:</p> <ol style="list-style-type: none"> <li>1. Start <b>rgmanager</b>, <b>cman</b>, and, if needed, <b>clvmd</b> and <b>gfs2</b>.</li> <li>2. Start <b>ricci</b>, and start <b>luci</b> if using the <b>luci</b> interface.</li> <li>3. Run <b>chkconfig on</b> for the needed services so that they start at each runtime.</li> </ol> <p>Alternately, you can enter <b>ccs --start</b> to start and enable the cluster services.</p>	<p>Start and enable cluster services with the following procedure:</p> <ol style="list-style-type: none"> <li>1. On every node, execute <b>systemctl start pcsd.service</b>, then <b>systemctl enable pcasd.service</b> to enable <b>pcsd</b> to start at runtime.</li> <li>2. On one node in the cluster, enter <b>pcs cluster start --all</b> to start <b>corosync</b> and <b>pacemaker</b>.</li> </ol>
Controlling access to configuration tools	For <b>luci</b> , the root user or a user with <b>luci</b> permissions can access <b>luci</b> . All access requires the <b>ricci</b> password for the node.	The <b>pcsd</b> gui requires that you authenticate as user <b>hacluster</b> , which is the common system user. The root user can set the password for <b>hacluster</b> .
Cluster creation	Name the cluster and define which nodes to include in the cluster with <b>luci</b> or <b>ccs</b> , or directly edit the <b>cluster.conf</b> file.	Name the cluster and include nodes with <b>pcs cluster setup</b> command or with the <b>pcsd</b> Web UI. You can add nodes to an existing cluster with the <b>pcs cluster node add</b> command or with the <b>pcsd</b> Web UI.
Propagating cluster configuration to all nodes	When configuration a cluster with <b>luci</b> , propagation is automatic. With <b>ccs</b> , use the <b>--sync</b> option. You can also use the <b>cman_tool version -r</b> command.	Propagation of the cluster and Pacemaker configuration files, <b>corosync.conf</b> and <b>cib.xml</b> , is automatic on cluster setup or when adding a node or resource.
Global cluster properties	<p>The following feature are supported with <b>rgmanager</b> in Red Hat Enterprise Linux 6:</p> <ul style="list-style-type: none"> <li>* You can configure the system so that the system chooses which multicast address to use for IP multicasting in the cluster network.</li> <li>* If IP multicasting is not available, you can use UDP Unicast transport mechanism.</li> <li>* You can configure a cluster to use RRP protocol.</li> </ul>	<p>Pacemaker in Red Hat Enterprise Linux 7 supports the following features for a cluster:</p> <ul style="list-style-type: none"> <li>* You can set <b>no-quorum-policy</b> for the cluster to specify what the system should do when the cluster does not have quorum.</li> <li>* For additional cluster properties you can set, refer to <a href="#">Table 12.1, “Cluster Properties”</a>.</li> </ul>
Logging	You can set global and daemon-specific logging configuration.	See the file <b>/etc/sysconfig/pacemaker</b> for information on how to configure logging manually.

Configuration Component	rgmanager	Pacemaker
Validating the cluster	Cluster validation is automatic with <b>luci</b> and with <b>ccs</b> , using the cluster schema. The cluster is automatically validated on startup.	The cluster is automatically validated on startup, or you can validate the cluster with <b>pcs cluster verify</b> .
Quorum in two-node clusters	With a two-node cluster, you can configure how the system determines quorum:  * Configure a quorum disk  * Use <b>ccs</b> or edit the <b>cluster.conf</b> file to set <b>two_node=1</b> and <b>expected_votes=1</b> to allow a single node to maintain quorum.	<b>pcs</b> automatically adds the necessary options for a two-node cluster to <b>corosync</b> .
Cluster status	On <b>luci</b> , the current status of the cluster is visible in the various components of the interface, which can be refreshed. You can use the <b>-getconf</b> option of the <b>ccs</b> command to see current the configuration file. You can use the <b>clustat</b> command to display cluster status.	You can display the current cluster status with the <b>pcs status</b> command.
Resources	You add resources of defined types and configure resource-specific properties with <b>luci</b> or the <b>ccs</b> command, or by editing the <b>cluster.conf</b> configuration file.	You add resources of defined types and configure resource-specific properties with the <b>pcs resource create</b> command or with the <b>pcsd</b> Web UI. For general information on configuring cluster resources with Pacemaker refer to <a href="#">Chapter 6, Configuring Cluster Resources</a> .

Configuration Component	rgmanager	Pacemaker
Resource behavior, grouping, and start/stop order	Define cluster services to configure how resources interact.	<p>With Pacemaker, you use resource groups as a shorthand method of defining a set of resources that need to be located together and started and stopped sequentially. In addition, you define how resources behave and interact in the following ways:</p>
		<ul style="list-style-type: none"> <li>* You set some aspects of resource behavior as resource options.</li> <li>* You use location constraints to determine which nodes a resource can run on.</li> <li>* You use order constraints to determine the order in which resources run.</li> <li>* You use colocation constraints to determine that the location of one resource depends on the location of another resource.</li> </ul>
		<p>For more complete information on these topics, refer to <a href="#">Chapter 6, Configuring Cluster Resources</a> and <a href="#">Chapter 7, Resource Constraints</a>.</p>
Resource administration: Moving, starting, stopping resources	With <b>luci</b> , you can manage clusters, individual cluster nodes, and cluster services. With the <b>ccs</b> command, you can manage cluster. You can use the <b>clusvadm</b> to manage cluster services.	You can temporarily disable a node so that it cannot host resources with the <b>pcs cluster standby</b> command, which causes the resources to migrate. You can stop a resource with the <b>pcs resource disable</b> command.
Removing a cluster configuration completely	With <b>luci</b> , you can select all nodes in a cluster for deletion to delete a cluster entirely. You can also remove the <b>cluster.conf</b> from each node in the cluster.	You can remove a cluster configuration with the <b>pcs cluster destroy</b> command.
Resources active on multiple nodes, resources active on multiple nodes in multiple modes	No equivalent.	With Pacemaker, you can clone resources so that they can run in multiple nodes, and you can define cloned resources as master and slave resources so that they can run in multiple modes. For information on cloned resources and master/slave resources, refer to <a href="#">Chapter 9, Advanced Resource types</a> .

Configuration Component	rgmanager	Pacemaker
Fencing -- single fence device per node	Create fencing devices globally or locally and add them to nodes. You can define <b>post-fail delay</b> and <b>post-join delay</b> values for the cluster as a whole.	Create a fencing device for each node with the <b>pcs stonith create</b> command or with the <b>pcsd</b> Web UI. For devices that can fence multiple nodes, you need to define them only once rather than separately for each node. You can also define <b>pcmk_host_map</b> to configure fencing devices for all nodes with a single command; for information on <b>pcmk_host_map</b> refer to <a href="#">Table 5.1, “General Properties of Fencing Devices”</a> . You can define the <b>stonith-timeout</b> value for the cluster as a whole.
Multiple (backup) fencing devices per node	Define backup devices with <b>luci</b> or the <b>ccs</b> command, or by editing the <b>cluster.conf</b> file directly.	Configure fencing levels.

## A.2. Cluster Creation with Pacemaker in Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7

The Red Hat Enterprise Linux 6.5 release supports cluster configuration with Pacemaker, using the **pcs** configuration tool. There are, however, some differences in cluster installation and creation between Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7 when using Pacemaker. This section provides a brief listing of the command differences between these two releases. For further information on installation and cluster creation in Red Hat Enterprise Linux 7, see [Chapter 1, Red Hat High Availability Add-On Configuration and Management Reference Overview](#) and [Chapter 4, Cluster Creation and Administration](#).

### A.2.1. Pacemaker Installation in Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7

The following commands install the Red Hat High Availability Add-On software packages that Pacemaker requires in Red Hat Enterprise Linux 6.5 and prevent **corosync** from starting without **cman**. You must enter these commands on each node in the cluster.

```
[root@rhel6]# yum install pacemaker cman
[root@rhel6]# yum install pcs
[root@rhel6]# chkconfig corosync off
```

In Red Hat Enterprise Linux 7, in addition to installing the Red Hat High Availability Add-On software packages that Pacemaker requires, you set up a password for the **pcs** administration account named **hacluster**, and you start and enable the **pcsd** service. You also authenticate the administration account for the nodes of the cluster.

In Red Hat Enterprise Linux 7, enter the following commands on each node in the cluster.

```
[root@rhel7]# yum install pcs fence-agents-all  
[root@rhel7]# passwd hacluster  
[root@rhel7]# systemctl start pcsd.service  
[root@rhel7]# systemctl enable pcsd.service
```

In Red Hat Enterprise Linux 7, enter the following command on one node in the cluster.

```
[root@rhel7]# pcs cluster auth [node] [...] [-u username] [-p password]
```

## A.2.2. Cluster Creation with Pacemaker in Red Hat Enterprise Linux 6.5 and Red Hat Enterprise Linux 7

To create a Pacemaker cluster in Red Hat Enterprise Linux 6.5, you must create the cluster and start the cluster services on each node in the cluster. For example, to create a cluster named **my\_cluster** that consists of nodes **z1.example.com** and **z2.example.com** and start cluster services on those nodes, enter the following commands from both **z1.example.com** and **z2.example.com**.

```
[root@rhel6]# pcs cluster setup --name my_cluster z1.example.com z2.example.com  
[root@rhel6]# pcs cluster start
```

In Red Hat Enterprise Linux 7, you enter the cluster creation command from one node of the cluster. The following command, run from one node only, creates the cluster named **my\_cluster** that consists of nodes **z1.example.com** and **z2.example.com** and starts cluster services on those nodes.

```
[root@rhel7]# pcs cluster setup --start --name my_cluster z1.example.com  
z2.example.com
```

## Appendix B. Revision History

<b>Revision 3.1-6</b>	<b>Wed Nov 23 2016</b>	<b>Steven Levine</b>
Update to version for 7.3 GA publication.		
<b>Revision 3.1-4</b>	<b>Mon Oct 17 2016</b>	<b>Steven Levine</b>
Version for 7.3 GA publication.		
<b>Revision 3.1-3</b>	<b>Wed Aug 17 2016</b>	<b>Steven Levine</b>
Preparing document for 7.3 Beta publication.		
<b>Revision 2.1-8</b>	<b>Mon Nov 9 2015</b>	<b>Steven Levine</b>
Preparing document for 7.2 GA publication		
<b>Revision 2.1-5</b>	<b>Mon Aug 24 2015</b>	<b>Steven Levine</b>
Preparing document for 7.2 Beta publication.		
<b>Revision 1.1-9</b>	<b>Mon Feb 23 2015</b>	<b>Steven Levine</b>
Version for 7.1 GA release		
<b>Revision 1.1-7</b>	<b>Thu Dec 11 2014</b>	<b>Steven Levine</b>
Version for 7.1 Beta release		
<b>Revision 0.1-41</b>	<b>Mon Jun 2 2014</b>	<b>Steven Levine</b>
Version for 7.0 GA release		
<b>Revision 0.1-2</b>	<b>Thu May 16 2013</b>	<b>Steven Levine</b>
First printing of initial draft		

## Index

- , [Cluster Creation](#)

### A

#### Action

- Property
  - enabled, [Resource Operations](#)
  - id, [Resource Operations](#)
  - interval, [Resource Operations](#)
  - name, [Resource Operations](#)
  - on-fail, [Resource Operations](#)
  - timeout, [Resource Operations](#)

#### Action Property, [Resource Operations](#)

#### attribute, [Node Attribute Expressions](#)

- Constraint Expression, [Node Attribute Expressions](#)

#### Attribute Expression, [Node Attribute Expressions](#)

- attribute, [Node Attribute Expressions](#)

- operation, [Node Attribute Expressions](#)
- type, [Node Attribute Expressions](#)
- value, [Node Attribute Expressions](#)

**B****batch-limit, [Summary of Cluster Properties and Options](#)**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**boolean-op, [Pacemaker Rules](#)**

- Constraint Rule, [Pacemaker Rules](#)

**C****Clone**

- Option
  - clone-max, [Creating and Removing a Cloned Resource](#)
  - clone-node-max, [Creating and Removing a Cloned Resource](#)
  - globally-unique, [Creating and Removing a Cloned Resource](#)
  - interleave, [Creating and Removing a Cloned Resource](#)
  - notify, [Creating and Removing a Cloned Resource](#)
  - ordered, [Creating and Removing a Cloned Resource](#)

**Clone Option, [Creating and Removing a Cloned Resource](#)****Clone Resources, [Resource Clones](#)****clone-max, [Creating and Removing a Cloned Resource](#)**

- Clone Option, [Creating and Removing a Cloned Resource](#)

**clone-node-max, [Creating and Removing a Cloned Resource](#)**

- Clone Option, [Creating and Removing a Cloned Resource](#)

**Clones, [Resource Clones](#)****Cluster**

- Option
  - batch-limit, [Summary of Cluster Properties and Options](#)
  - cluster-delay, [Summary of Cluster Properties and Options](#)
  - cluster-infrastructure, [Summary of Cluster Properties and Options](#)
  - cluster-recheck-interval, [Summary of Cluster Properties and Options](#)
  - dc-version, [Summary of Cluster Properties and Options](#)
  - default-action-timeout, [Summary of Cluster Properties and Options](#)
  - default-resource-stickiness, [Summary of Cluster Properties and Options](#)
  - enable-acl, [Summary of Cluster Properties and Options](#)
  - is-managed-default, [Summary of Cluster Properties and Options](#)
  - last-lrm-refresh, [Summary of Cluster Properties and Options](#)
  - maintenance-mode, [Summary of Cluster Properties and Options](#)
  - migration-limit, [Summary of Cluster Properties and Options](#)
  - no-quorum-policy, [Summary of Cluster Properties and Options](#)
  - pe-error-series-max, [Summary of Cluster Properties and Options](#)
  - pe-input-series-max, [Summary of Cluster Properties and Options](#)
  - pe-warn-series-max, [Summary of Cluster Properties and Options](#)
  - shutdown-escalation, [Summary of Cluster Properties and Options](#)
  - start-failure-is-fatal, [Summary of Cluster Properties and Options](#)
  - stonith-action, [Summary of Cluster Properties and Options](#)
  - stonith-enabled, [Summary of Cluster Properties and Options](#)
  - stonith-timeout, [Summary of Cluster Properties and Options](#)
  - stop-all-resources, [Summary of Cluster Properties and Options](#)

- stop-orphan-actions, [Summary of Cluster Properties and Options](#)
- stop-orphan-resources, [Summary of Cluster Properties and Options](#)
- symmetric-cluster, [Summary of Cluster Properties and Options](#)
  
- Querying Properties, [Querying Cluster Property Settings](#)
- Removing Properties, [Setting and Removing Cluster Properties](#)
- Setting Properties, [Setting and Removing Cluster Properties](#)

**Cluster Option, [Summary of Cluster Properties and Options](#)****Cluster Properties, [Setting and Removing Cluster Properties](#), [Querying Cluster Property Settings](#)****cluster status**

- display, [Displaying Cluster Status](#)

**cluster-delay, [Summary of Cluster Properties and Options](#)**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**cluster-infrastructure, [Summary of Cluster Properties and Options](#)**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**cluster-recheck-interval, [Summary of Cluster Properties and Options](#)**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**Colocation, [Colocation of Resources](#)****Constraint**

- Attribute Expression, [Node Attribute Expressions](#)
  - attribute, [Node Attribute Expressions](#)
  - operation, [Node Attribute Expressions](#)
  - type, [Node Attribute Expressions](#)
  - value, [Node Attribute Expressions](#)
  
- Date Specification, [Date Specifications](#)
  - hours, [Date Specifications](#)
  - id, [Date Specifications](#)
  - monthdays, [Date Specifications](#)
  - months, [Date Specifications](#)
  - moon, [Date Specifications](#)
  - weekdays, [Date Specifications](#)
  - weeks, [Date Specifications](#)
  - weekyears, [Date Specifications](#)
  - yeardays, [Date Specifications](#)
  - years, [Date Specifications](#)
  
- Date/Time Expression, [Time/Date Based Expressions](#)
  - end, [Time/Date Based Expressions](#)
  - operation, [Time/Date Based Expressions](#)
  - start, [Time/Date Based Expressions](#)
  
- Duration, [Durations](#)
- Rule, [Pacemaker Rules](#)
  - boolean-op, [Pacemaker Rules](#)
  - role, [Pacemaker Rules](#)
  - score, [Pacemaker Rules](#)
  - score-attribute, [Pacemaker Rules](#)

**Constraint Expression, Node Attribute Expressions, Time/Date Based Expressions****Constraint Rule, Pacemaker Rules****Constraints**

- Colocation, [Colocation of Resources](#)
- Location
  - id, [Location Constraints](#)
  - score, [Location Constraints](#)
- Order, [Order Constraints](#)
  - kind, [Order Constraints](#)

**D****dampen, Moving Resources Due to Connectivity Changes**

- Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

**Date Specification, Date Specifications**

- hours, [Date Specifications](#)
- id, [Date Specifications](#)
- monthdays, [Date Specifications](#)
- months, [Date Specifications](#)
- moon, [Date Specifications](#)
- weekdays, [Date Specifications](#)
- weeks, [Date Specifications](#)
- weekyears, [Date Specifications](#)
- yeardays, [Date Specifications](#)
- years, [Date Specifications](#)

**Date/Time Expression, Time/Date Based Expressions**

- end, [Time/Date Based Expressions](#)
- operation, [Time/Date Based Expressions](#)
- start, [Time/Date Based Expressions](#)

**dc-version, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**default-action-timeout, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**default-resource-stickiness, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**Determine by Rules, Using Rules to Determine Resource Location****Determine Resource Location, Using Rules to Determine Resource Location****disabling**

- resources, [Enabling and Disabling Cluster Resources](#)

**Duration, Durations****E****enable-acl, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**enabled, Resource Operations**

- Action Property, [Resource Operations](#)

## **e**nabling

- resources, [Enabling and Disabling Cluster Resources](#)

## **e**nd, [Time/Date Based Expressions](#)

- Constraint Expression, [Time/Date Based Expressions](#)

## **F**

### **f**ailure-timeout, [Resource Meta Options](#)

- Resource Option, [Resource Meta Options](#)

### **f**eatures, new and changed, [New and Changed Features](#)

## **G**

### **g**lobally-unique, [Creating and Removing a Cloned Resource](#)

- Clone Option, [Creating and Removing a Cloned Resource](#)

### **g**roup Resources, [Resource Groups](#)

### **g**roups, [Resource Groups](#), [Group Stickiness](#)

## **H**

### **h**ost\_list, [Moving Resources Due to Connectivity Changes](#)

- Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

### **h**ours, [Date Specifications](#)

- Date Specification, [Date Specifications](#)

## **I**

### **i**d, [Resource Properties](#), [Resource Operations](#), [Date Specifications](#)

- Action Property, [Resource Operations](#)  
- Date Specification, [Date Specifications](#)  
- Location Constraints, [Location Constraints](#)  
- Multi-State Property, [MultiState Resources: Resources That Have Multiple Modes](#)  
- Resource, [Resource Properties](#)

### **i**nterleave, [Creating and Removing a Cloned Resource](#)

- Clone Option, [Creating and Removing a Cloned Resource](#)

### **i**nterval, [Resource Operations](#)

- Action Property, [Resource Operations](#)

### **i**s-managed, [Resource Meta Options](#)

- Resource Option, [Resource Meta Options](#)

### **i**s-managed-default, [Summary of Cluster Properties and Options](#)

- Cluster Option, [Summary of Cluster Properties and Options](#)

## **K**

### **k**ind, [Order Constraints](#)

- Order Constraints, [Order Constraints](#)

## **L**

**last-lrm-refresh, Summary of Cluster Properties and Options**  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

## Location

- Determine by Rules, [Using Rules to Determine Resource Location](#)
- score, [Location Constraints](#)

**Location Constraints, Location Constraints**

**Location Relative to other Resources, Colocation of Resources**

## M

**maintenance-mode, Summary of Cluster Properties and Options**  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**master-max, MultiState Resources: Resources That Have Multiple Modes**  
 - Multi-State Option, [MultiState Resources: Resources That Have Multiple Modes](#)

**master-node-max, MultiState Resources: Resources That Have Multiple Modes**  
 - Multi-State Option, [MultiState Resources: Resources That Have Multiple Modes](#)

**migration-limit, Summary of Cluster Properties and Options**  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**migration-threshold, Resource Meta Options**  
 - Resource Option, [Resource Meta Options](#)

**monthdays, Date Specifications**  
 - Date Specification, [Date Specifications](#)

**months, Date Specifications**  
 - Date Specification, [Date Specifications](#)

**moon, Date Specifications**  
 - Date Specification, [Date Specifications](#)

**Moving, Manually Moving Resources Around the Cluster**  
 - Resources, [Manually Moving Resources Around the Cluster](#)

## Multi-State

- Option
  - master-max, [MultiState Resources: Resources That Have Multiple Modes](#)
  - master-node-max, [MultiState Resources: Resources That Have Multiple Modes](#)
- Property
  - id, [MultiState Resources: Resources That Have Multiple Modes](#)

**Multi-State Option, MultiState Resources: Resources That Have Multiple Modes**

**Multi-State Property, MultiState Resources: Resources That Have Multiple Modes**

## multiple-active, Resource Meta Options

- Resource Option, [Resource Meta Options](#)

**multiplier, Moving Resources Due to Connectivity Changes**

- Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

## Multistate, [MultiState Resources: Resources That Have Multiple Modes](#), Multistate Stickiness

N

### **name**, [Resource Operations](#)

- Action Property, [Resource Operations](#)

### **no-quorum-policy**, [Summary of Cluster Properties and Options](#)

- Cluster Option, [Summary of Cluster Properties and Options](#)

### **notify**, [Creating and Removing a Cloned Resource](#)

- Clone Option, [Creating and Removing a Cloned Resource](#)

O

### **on-fail**, [Resource Operations](#)

- Action Property, [Resource Operations](#)

### **operation**, [Node Attribute Expressions, Time/Date Based Expressions](#)

- Constraint Expression, [Node Attribute Expressions, Time/Date Based Expressions](#)

### Option

- batch-limit, [Summary of Cluster Properties and Options](#)
- clone-max, [Creating and Removing a Cloned Resource](#)
- clone-node-max, [Creating and Removing a Cloned Resource](#)
- cluster-delay, [Summary of Cluster Properties and Options](#)
- cluster-infrastructure, [Summary of Cluster Properties and Options](#)
- cluster-recheck-interval, [Summary of Cluster Properties and Options](#)
- dampen, [Moving Resources Due to Connectivity Changes](#)
- dc-version, [Summary of Cluster Properties and Options](#)
- default-action-timeout, [Summary of Cluster Properties and Options](#)
- default-resource-stickiness, [Summary of Cluster Properties and Options](#)
- enable-acl, [Summary of Cluster Properties and Options](#)
- failure-timeout, [Resource Meta Options](#)
- globally-unique, [Creating and Removing a Cloned Resource](#)
- host\_list, [Moving Resources Due to Connectivity Changes](#)
- interleave, [Creating and Removing a Cloned Resource](#)
- is-managed, [Resource Meta Options](#)
- is-managed-default, [Summary of Cluster Properties and Options](#)
- last-lrm-refresh, [Summary of Cluster Properties and Options](#)
- maintenance-mode, [Summary of Cluster Properties and Options](#)
- master-max, [MultiState Resources: Resources That Have Multiple Modes](#)
- master-node-max, [MultiState Resources: Resources That Have Multiple Modes](#)
- migration-limit, [Summary of Cluster Properties and Options](#)
- migration-threshold, [Resource Meta Options](#)
- multiple-active, [Resource Meta Options](#)
- multiplier, [Moving Resources Due to Connectivity Changes](#)
- no-quorum-policy, [Summary of Cluster Properties and Options](#)
- notify, [Creating and Removing a Cloned Resource](#)
- ordered, [Creating and Removing a Cloned Resource](#)
- pe-error-series-max, [Summary of Cluster Properties and Options](#)
- pe-input-series-max, [Summary of Cluster Properties and Options](#)
- pe-warn-series-max, [Summary of Cluster Properties and Options](#)
- priority, [Resource Meta Options](#)

- requires, [Resource Meta Options](#)
- resource-stickiness, [Resource Meta Options](#)
- shutdown-escalation, [Summary of Cluster Properties and Options](#)
- start-failure-is-fatal, [Summary of Cluster Properties and Options](#)
- stonith-action, [Summary of Cluster Properties and Options](#)
- stonith-enabled, [Summary of Cluster Properties and Options](#)
- stonith-timeout, [Summary of Cluster Properties and Options](#)
- stop-all-resources, [Summary of Cluster Properties and Options](#)
- stop-orphan-actions, [Summary of Cluster Properties and Options](#)
- stop-orphan-resources, [Summary of Cluster Properties and Options](#)
- symmetric-cluster, [Summary of Cluster Properties and Options](#)
- target-role, [Resource Meta Options](#)

## Order

- kind, [Order Constraints](#)

## Order Constraints, [Order Constraints](#)

- symmetrical, [Order Constraints](#)

## ordered, [Creating and Removing a Cloned Resource](#)

- Clone Option, [Creating and Removing a Cloned Resource](#)

## Ordering, [Order Constraints](#)

### overview

- features, new and changed, [New and Changed Features](#)

## P

## pe-error-series-max, [Summary of Cluster Properties and Options](#)

- Cluster Option, [Summary of Cluster Properties and Options](#)

## pe-input-series-max, [Summary of Cluster Properties and Options](#)

- Cluster Option, [Summary of Cluster Properties and Options](#)

## pe-warn-series-max, [Summary of Cluster Properties and Options](#)

- Cluster Option, [Summary of Cluster Properties and Options](#)

## Ping Resource

- Option

- dampen, [Moving Resources Due to Connectivity Changes](#)
- host\_list, [Moving Resources Due to Connectivity Changes](#)
- multiplier, [Moving Resources Due to Connectivity Changes](#)

## Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

### priority, [Resource Meta Options](#)

- Resource Option, [Resource Meta Options](#)

## Property

- enabled, [Resource Operations](#)
- id, [Resource Properties](#), [Resource Operations](#), [MultiState Resources: Resources That Have Multiple Modes](#)
- interval, [Resource Operations](#)
- name, [Resource Operations](#)
- on-fail, [Resource Operations](#)
- provider, [Resource Properties](#)

- standard, [Resource Properties](#)
- timeout, [Resource Operations](#)
- type, [Resource Properties](#)

**provider, Resource Properties**

- Resource, [Resource Properties](#)

**Q****Querying**

- Cluster Properties, [Querying Cluster Property Settings](#)

**Querying Options, Querying Cluster Property Settings****R****Removing**

- Cluster Properties, [Setting and Removing Cluster Properties](#)

**Removing Properties, Setting and Removing Cluster Properties****requires, Resource Meta Options****Resource, Resource Properties**

- Constraint
  - Attribute Expression, [Node Attribute Expressions](#)
  - Date Specification, [Date Specifications](#)
  - Date/Time Expression, [Time/Date Based Expressions](#)
  - Duration, [Durations](#)
  - Rule, [Pacemaker Rules](#)
- Constraints
  - Colocation, [Colocation of Resources](#)
  - Order, [Order Constraints](#)
- Location
  - Determine by Rules, [Using Rules to Determine Resource Location](#)
  - Location Relative to other Resources, [Colocation of Resources](#)
  - Moving, [Manually Moving Resources Around the Cluster](#)
  - Option
    - failure-timeout, [Resource Meta Options](#)
    - is-managed, [Resource Meta Options](#)
    - migration-threshold, [Resource Meta Options](#)
    - multiple-active, [Resource Meta Options](#)
    - priority, [Resource Meta Options](#)
    - requires, [Resource Meta Options](#)
    - resource-stickiness, [Resource Meta Options](#)
    - target-role, [Resource Meta Options](#)
- Property
  - id, [Resource Properties](#)
  - provider, [Resource Properties](#)
  - standard, [Resource Properties](#)
  - type, [Resource Properties](#)
- Start Order, [Order Constraints](#)

**Resource Option, Resource Meta Options**

**resource-stickiness, Resource Meta Options**

- Groups, [Group Stickiness](#)
- Multi-State, [Multistate Stickiness](#)
- Resource Option, [Resource Meta Options](#)

**Resources, Manually Moving Resources Around the Cluster**

- Clones, [Resource Clones](#)
- Groups, [Resource Groups](#)
- Multistate, [MultiState Resources: Resources That Have Multiple Modes](#)

**resources**

- cleanup, [Cluster Resources Cleanup](#)
- disabling, [Enabling and Disabling Cluster Resources](#)
- enabling, [Enabling and Disabling Cluster Resources](#)

**role, Pacemaker Rules**

- Constraint Rule, [Pacemaker Rules](#)

**Rule, Pacemaker Rules**

- boolean-op, [Pacemaker Rules](#)
- Determine Resource Location, [Using Rules to Determine Resource Location](#)
- role, [Pacemaker Rules](#)
- score, [Pacemaker Rules](#)
- score-attribute, [Pacemaker Rules](#)

**S****score, Location Constraints, Pacemaker Rules**

- Constraint Rule, [Pacemaker Rules](#)
- Location Constraints, [Location Constraints](#)

**score-attribute, Pacemaker Rules**

- Constraint Rule, [Pacemaker Rules](#)

**Setting**

- Cluster Properties, [Setting and Removing Cluster Properties](#)

**Setting Properties, Setting and Removing Cluster Properties****shutdown-escalation, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**standard, Resource Properties**

- Resource, [Resource Properties](#)

**start, Time/Date Based Expressions**

- Constraint Expression, [Time/Date Based Expressions](#)

**Start Order, Order Constraints****start-failure-is-fatal, Summary of Cluster Properties and Options**

- Cluster Option, [Summary of Cluster Properties and Options](#)

**status**

- display, [Displaying Cluster Status](#)

**stonith-action**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**stonith-enabled**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**stonith-timeout**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**stop-all-resources**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**stop-orphan-actions**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**stop-orphan-resources**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**symmetric-cluster**, [Summary of Cluster Properties and Options](#)  
 - Cluster Option, [Summary of Cluster Properties and Options](#)

**symmetrical**, [Order Constraints](#)  
 - Order Constraints, [Order Constraints](#)

**T**

**target-role**, [Resource Meta Options](#)  
 - Resource Option, [Resource Meta Options](#)

**Time Based Expressions**, [Time/Date Based Expressions](#)

**timeout**, [Resource Operations](#)  
 - Action Property, [Resource Operations](#)

**type**, [Resource Properties](#), [Node Attribute Expressions](#)  
 - Constraint Expression, [Node Attribute Expressions](#)  
 - Resource, [Resource Properties](#)

**V**

**value**, [Node Attribute Expressions](#)  
 - Constraint Expression, [Node Attribute Expressions](#)

**W**

**weekdays**, [Date Specifications](#)  
 - Date Specification, [Date Specifications](#)

**weeks**, [Date Specifications](#)  
 - Date Specification, [Date Specifications](#)

**weekyears**, [Date Specifications](#)  
 - Date Specification, [Date Specifications](#)

**Y**

**yeardays**, [Date Specifications](#)  
 - Date Specification, [Date Specifications](#)

years, [Date Specifications](#)

- Date Specification, [Date Specifications](#)